

Computer Vision

Chapter 5: Convolution Neural Network

Why CNN for Image

- Some patterns are much smaller than the whole image

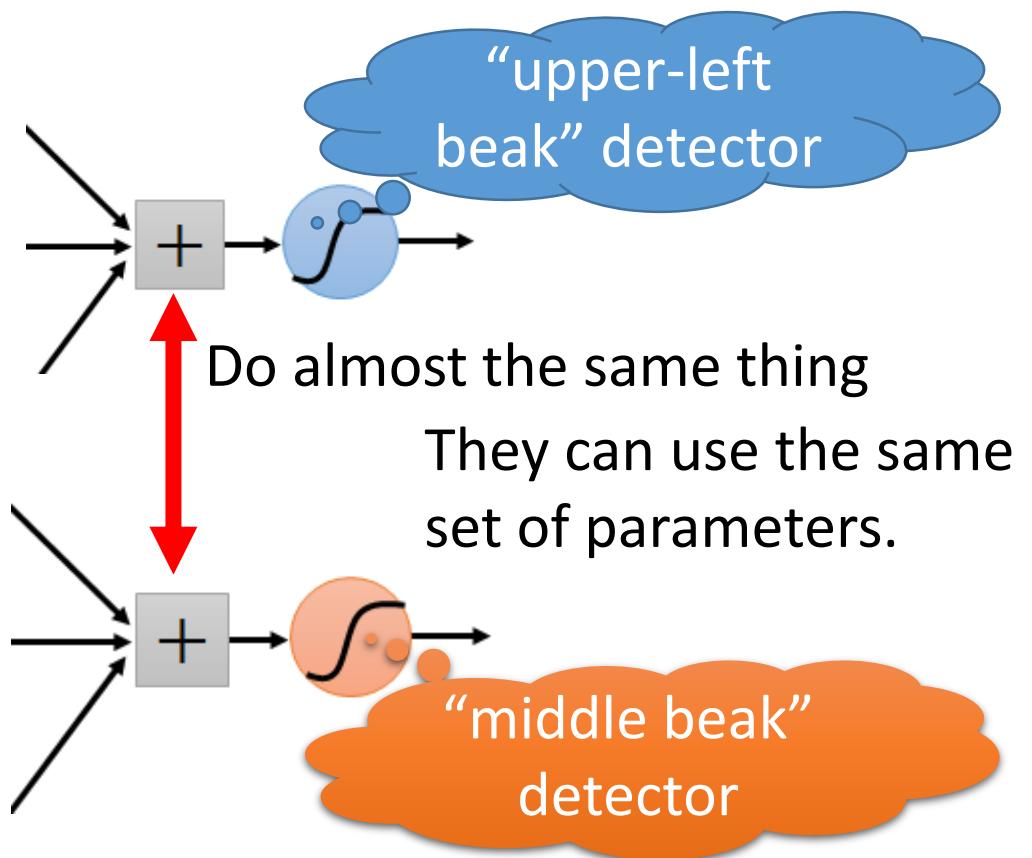
A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



Why CNN for Image

- The same patterns appear in different regions.



Why CNN for Image

- Subsampling the pixels will not change the object

bird



subsampling

bird

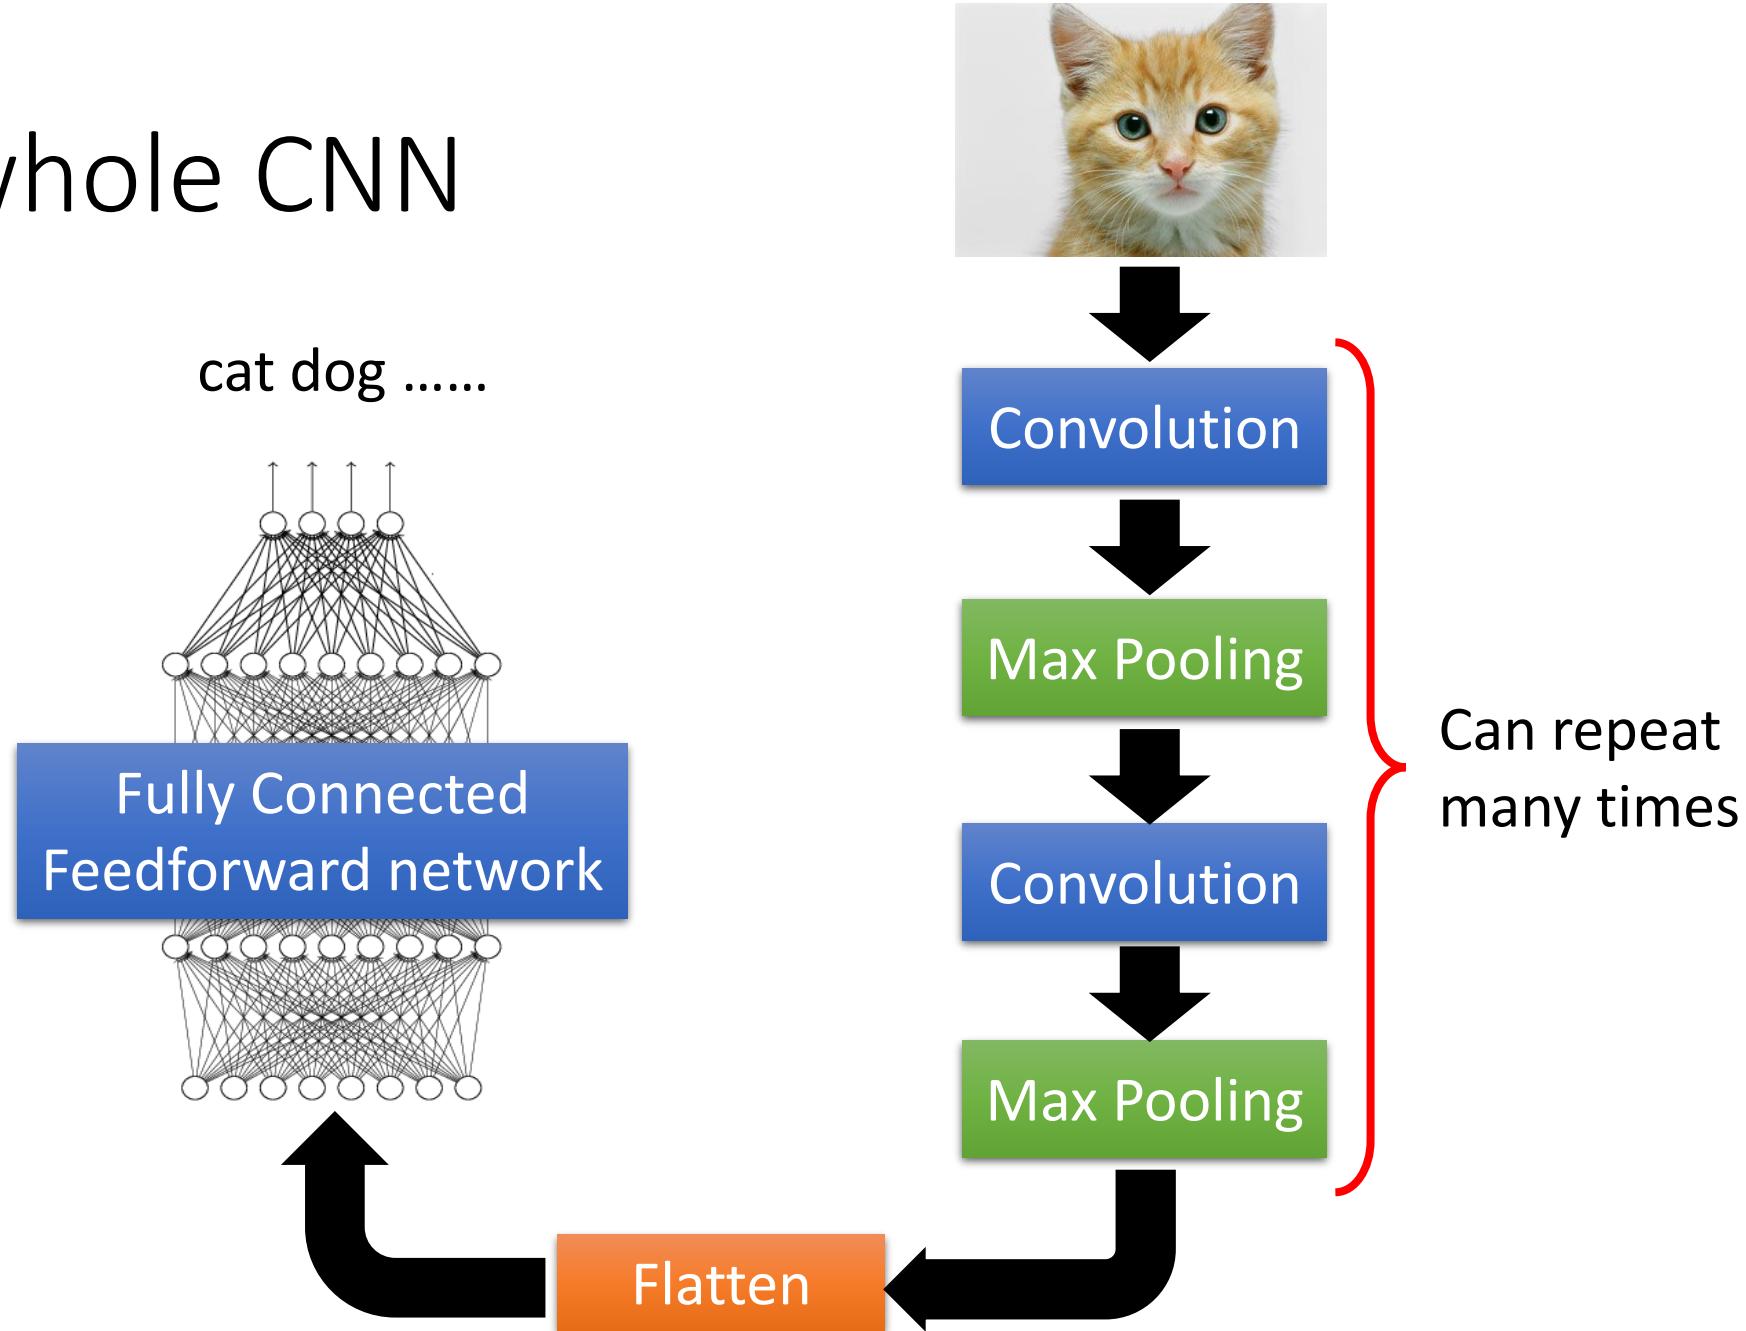


We can subsample the pixels to make image smaller



Less parameters for the network to process the image

The whole CNN



The whole CNN

Property 1

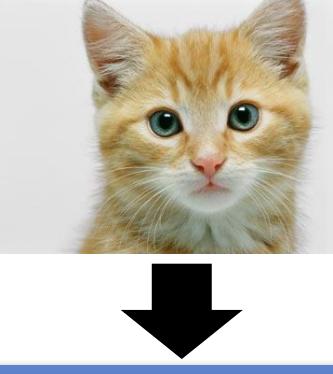
- Some patterns are much smaller than the whole image

Property 2

- The same patterns appear in different regions.

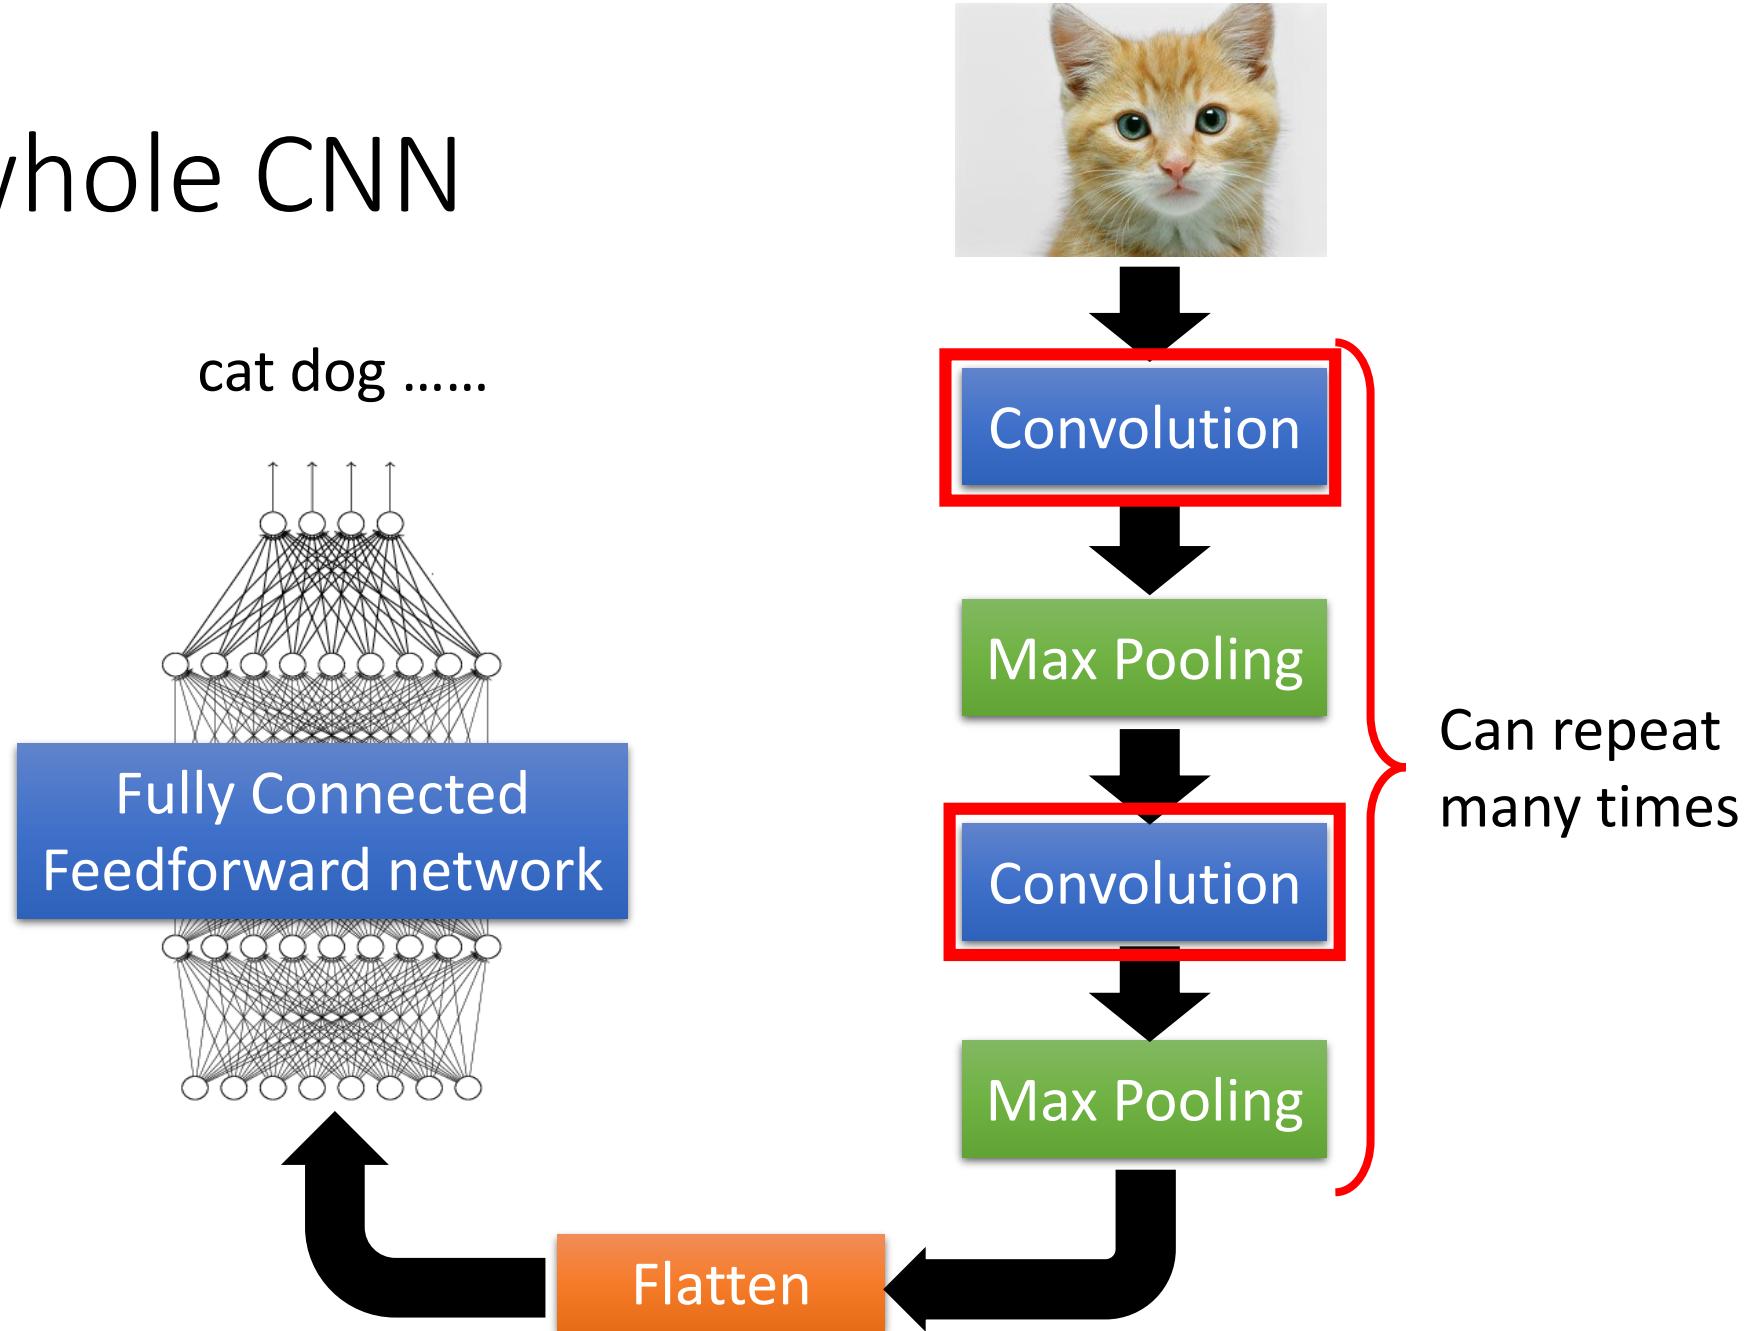
Property 3

- Subsampling the pixels will not change the object



Can repeat
many times

The whole CNN



CNN – Convolution

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

CNN – Convolution

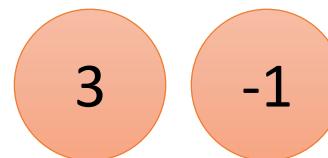
stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



CNN – Convolution

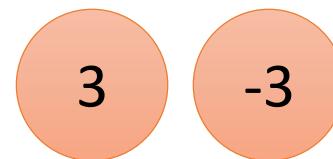
If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

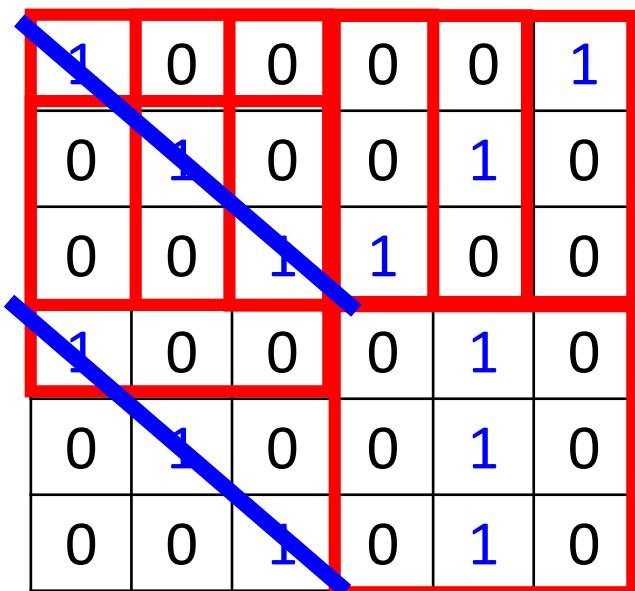
Filter 1



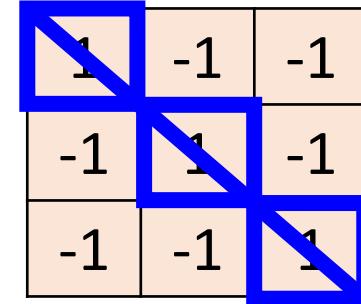
We set stride=1 below

CNN – Convolution

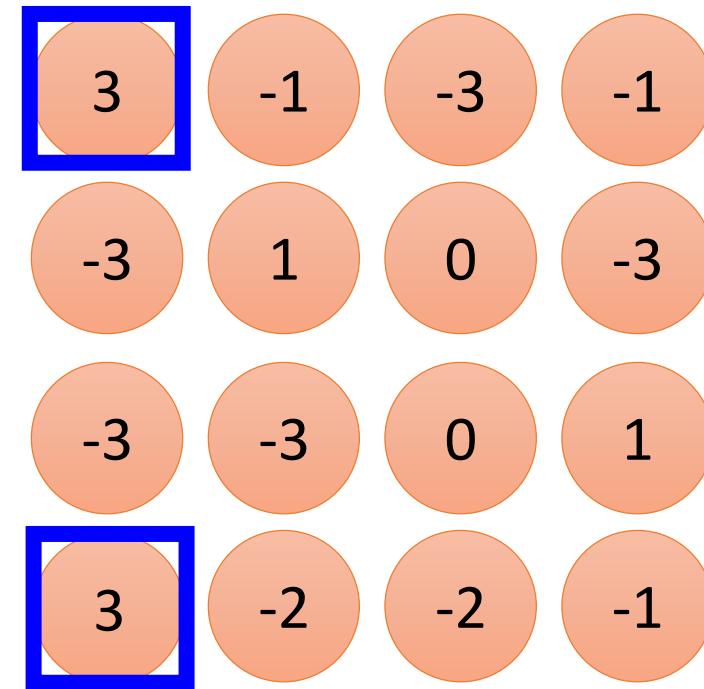
stride=1



6 x 6 image



Filter 1



Property 2

CNN – Convolution

stride=1

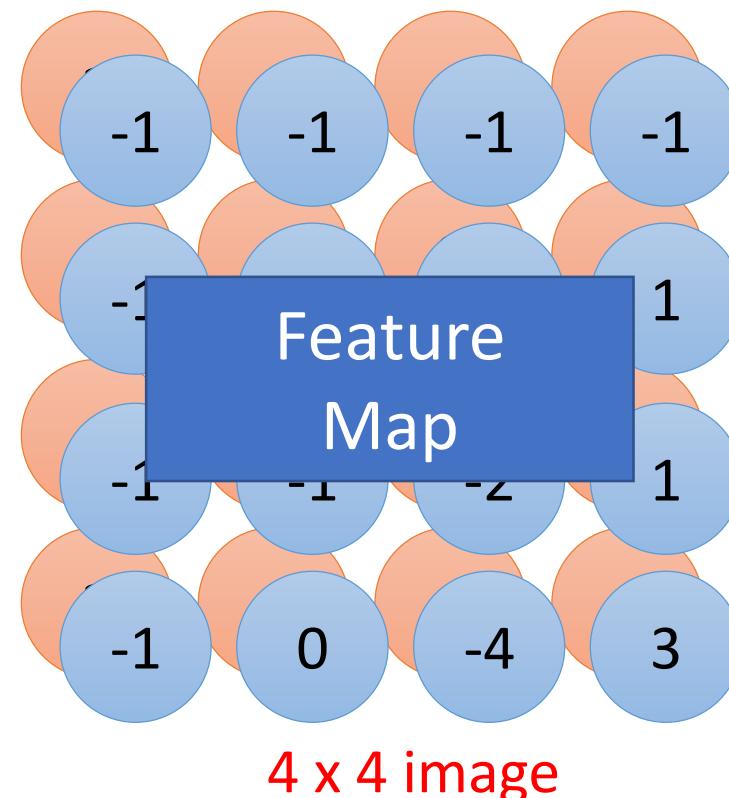
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

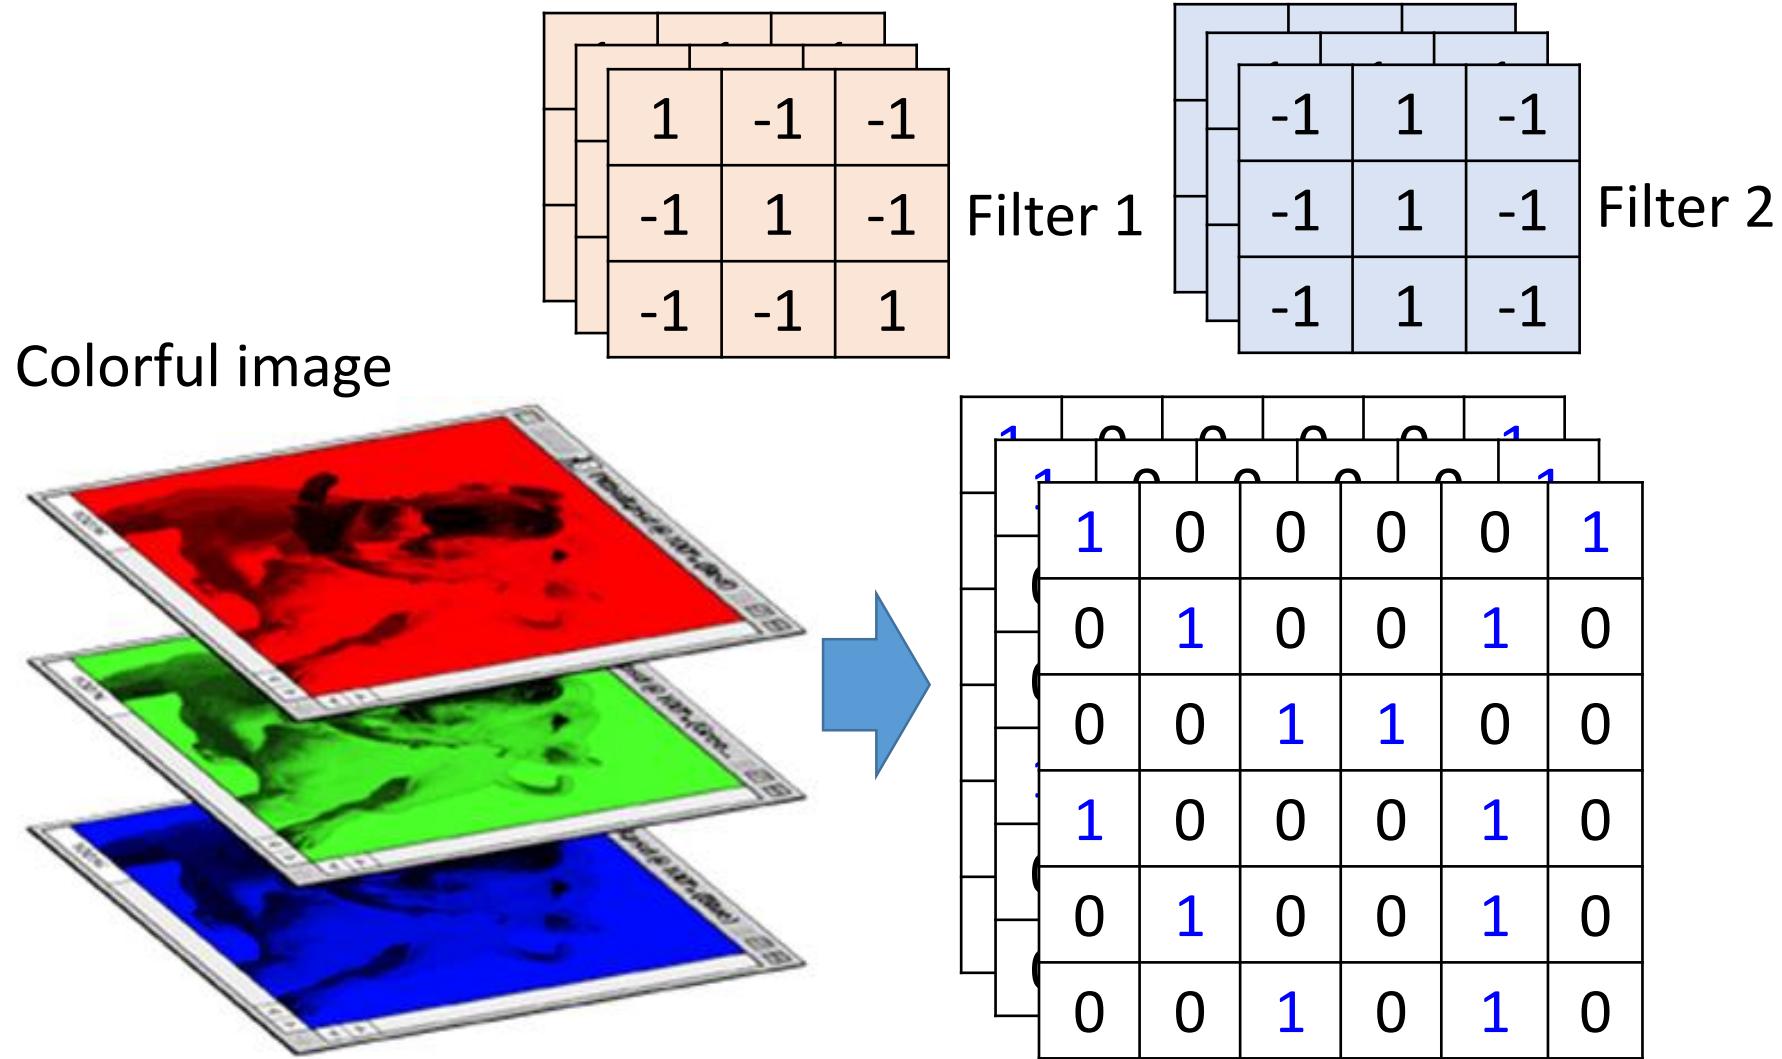
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

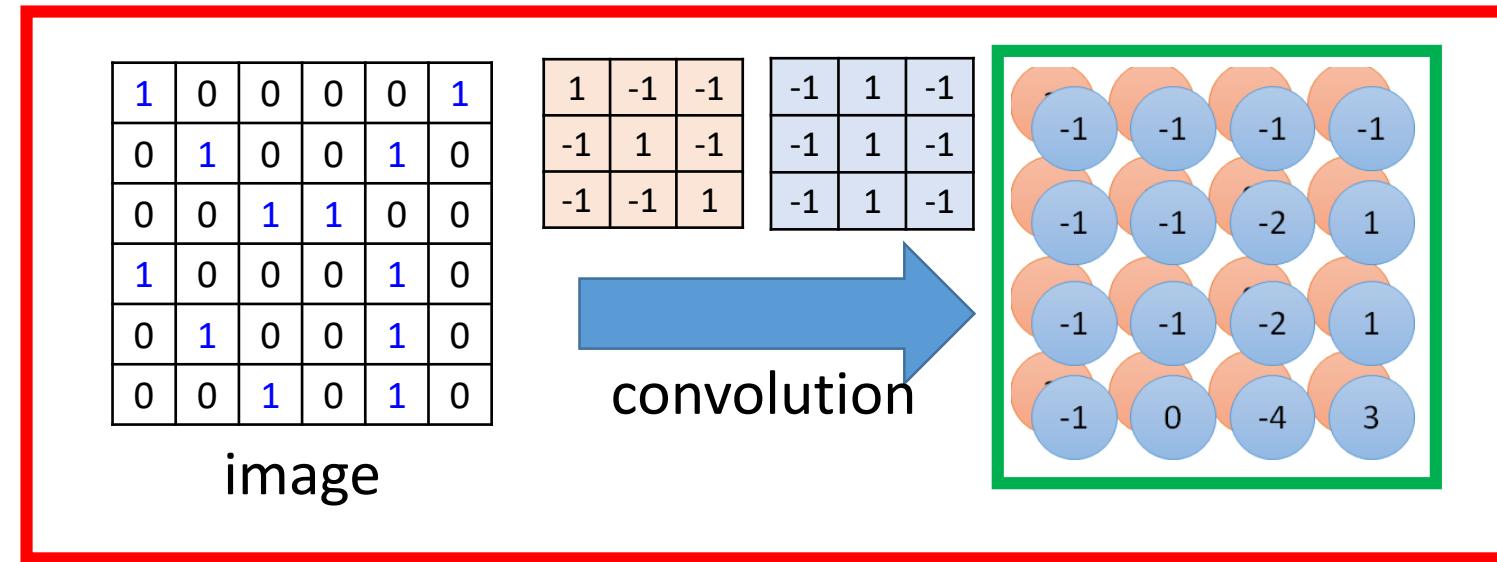
Do the same process for every filter



CNN – Colorful image

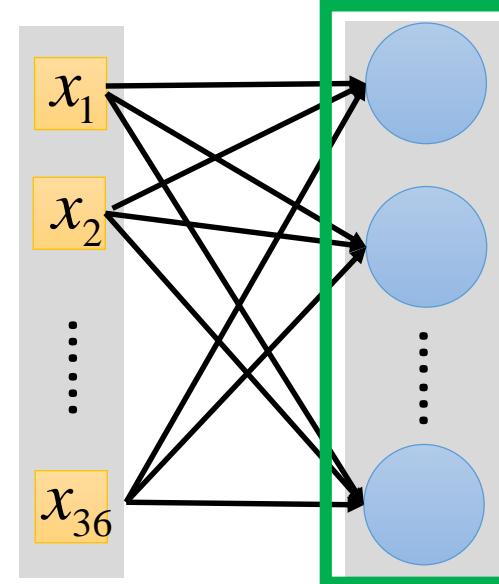


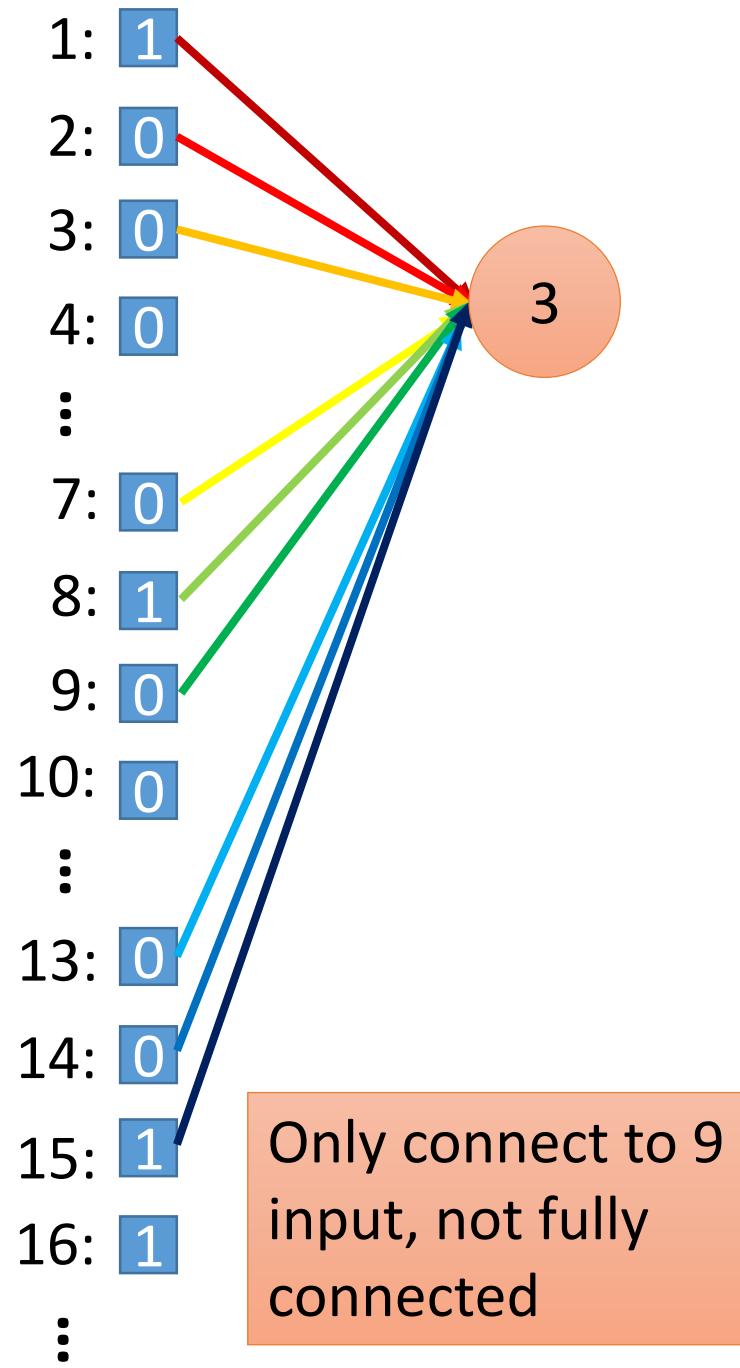
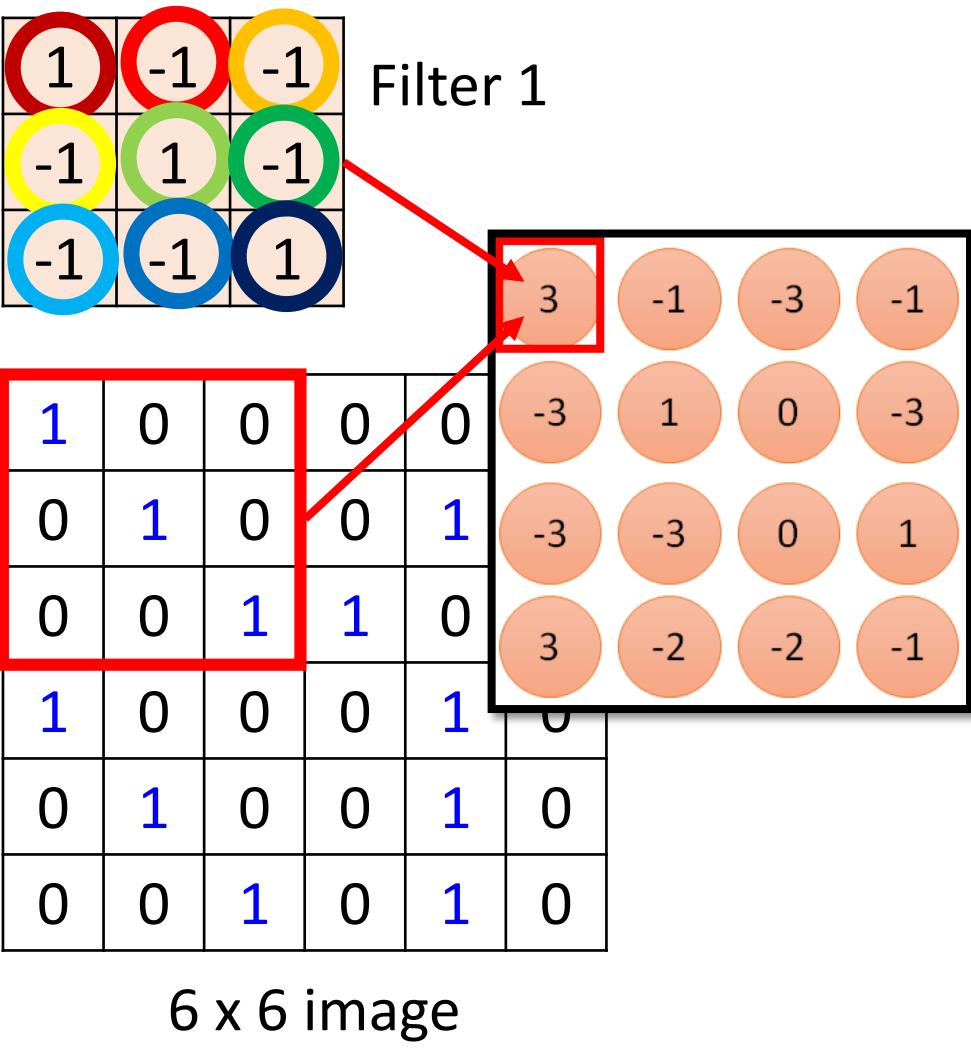
Convolution v.s. Fully Connected

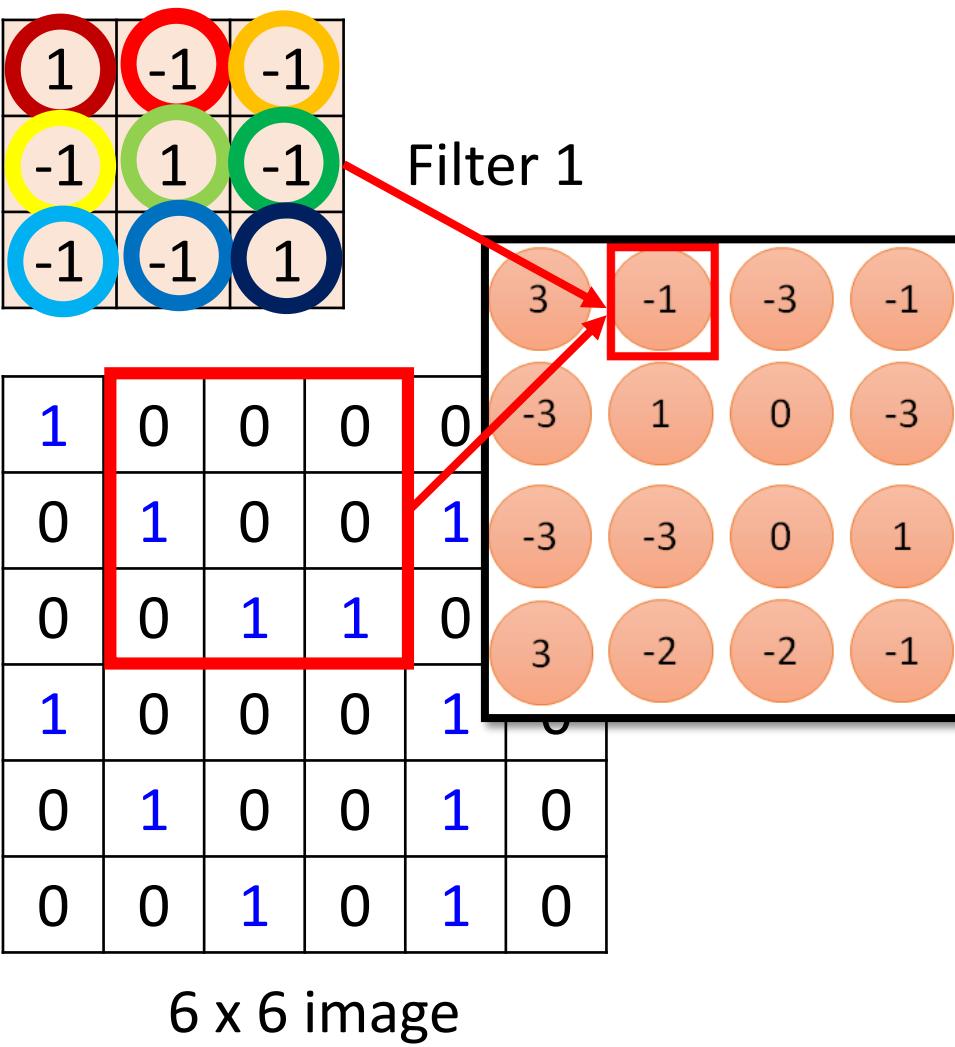


Fully-
connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

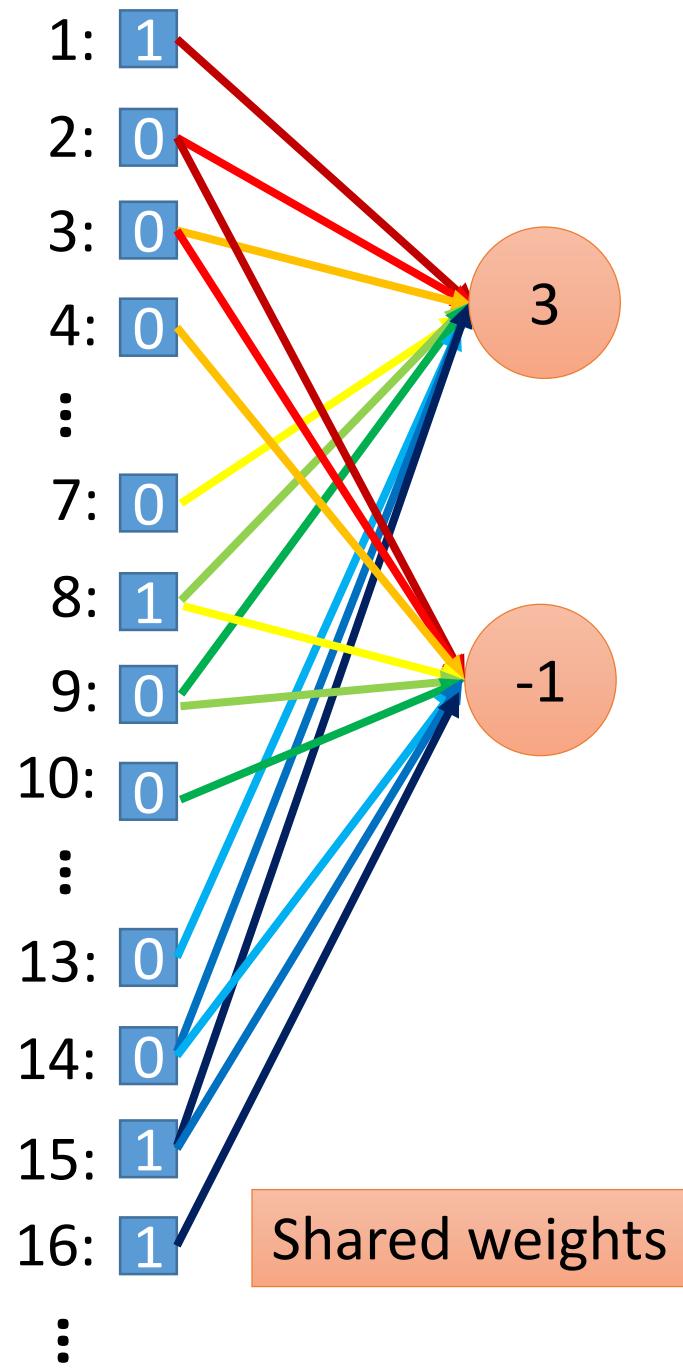




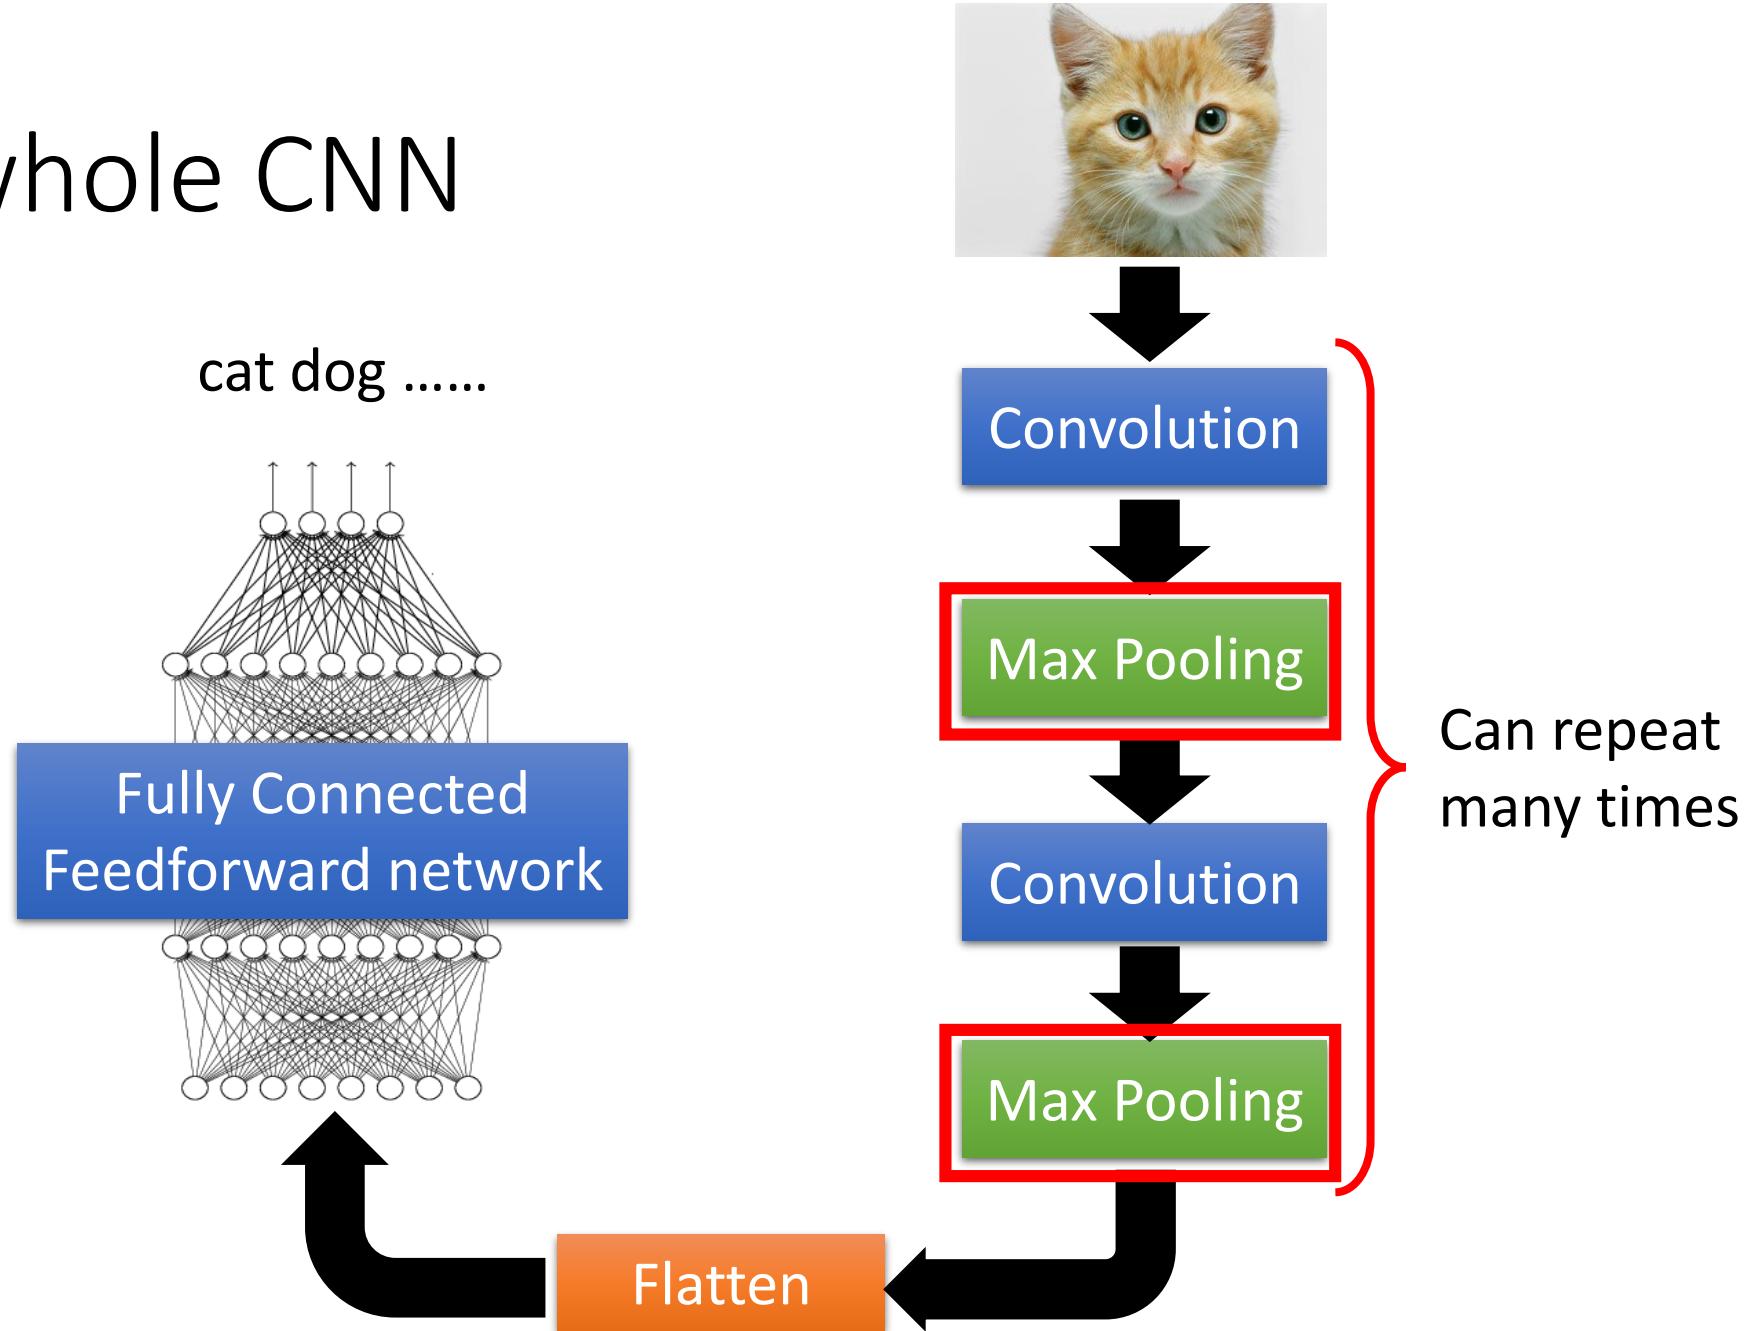


Less parameters!

Even less parameters!



The whole CNN



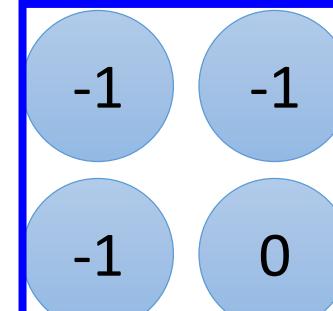
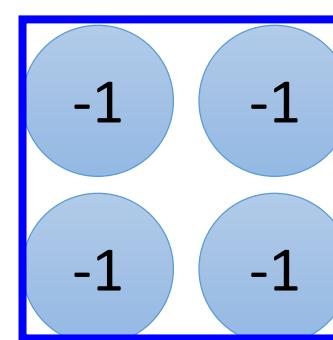
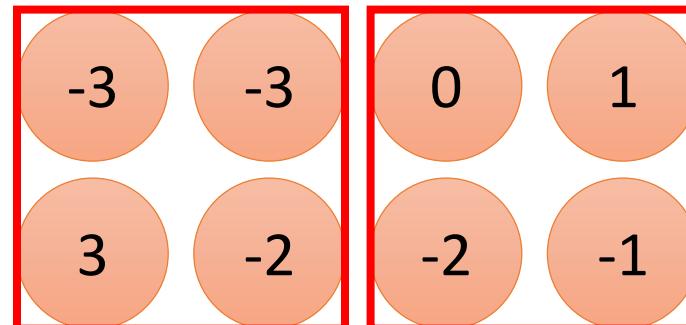
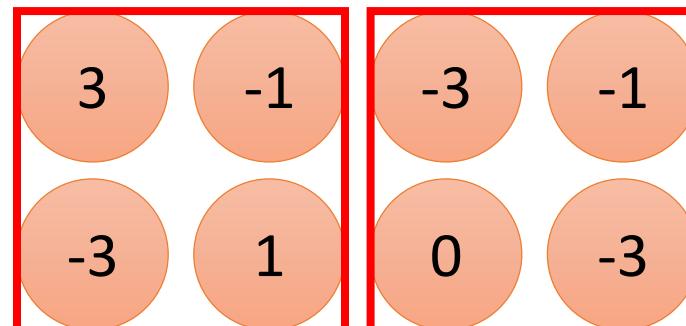
CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

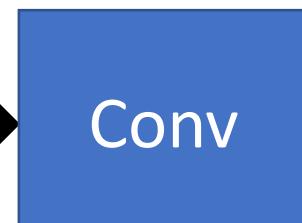
Filter 2



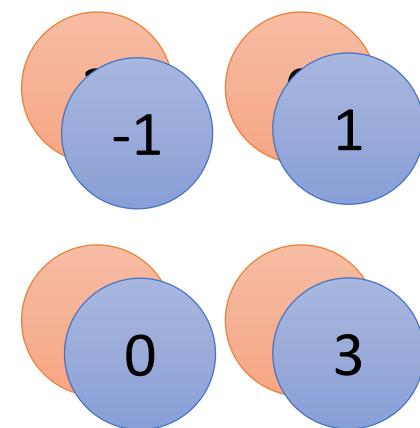
CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



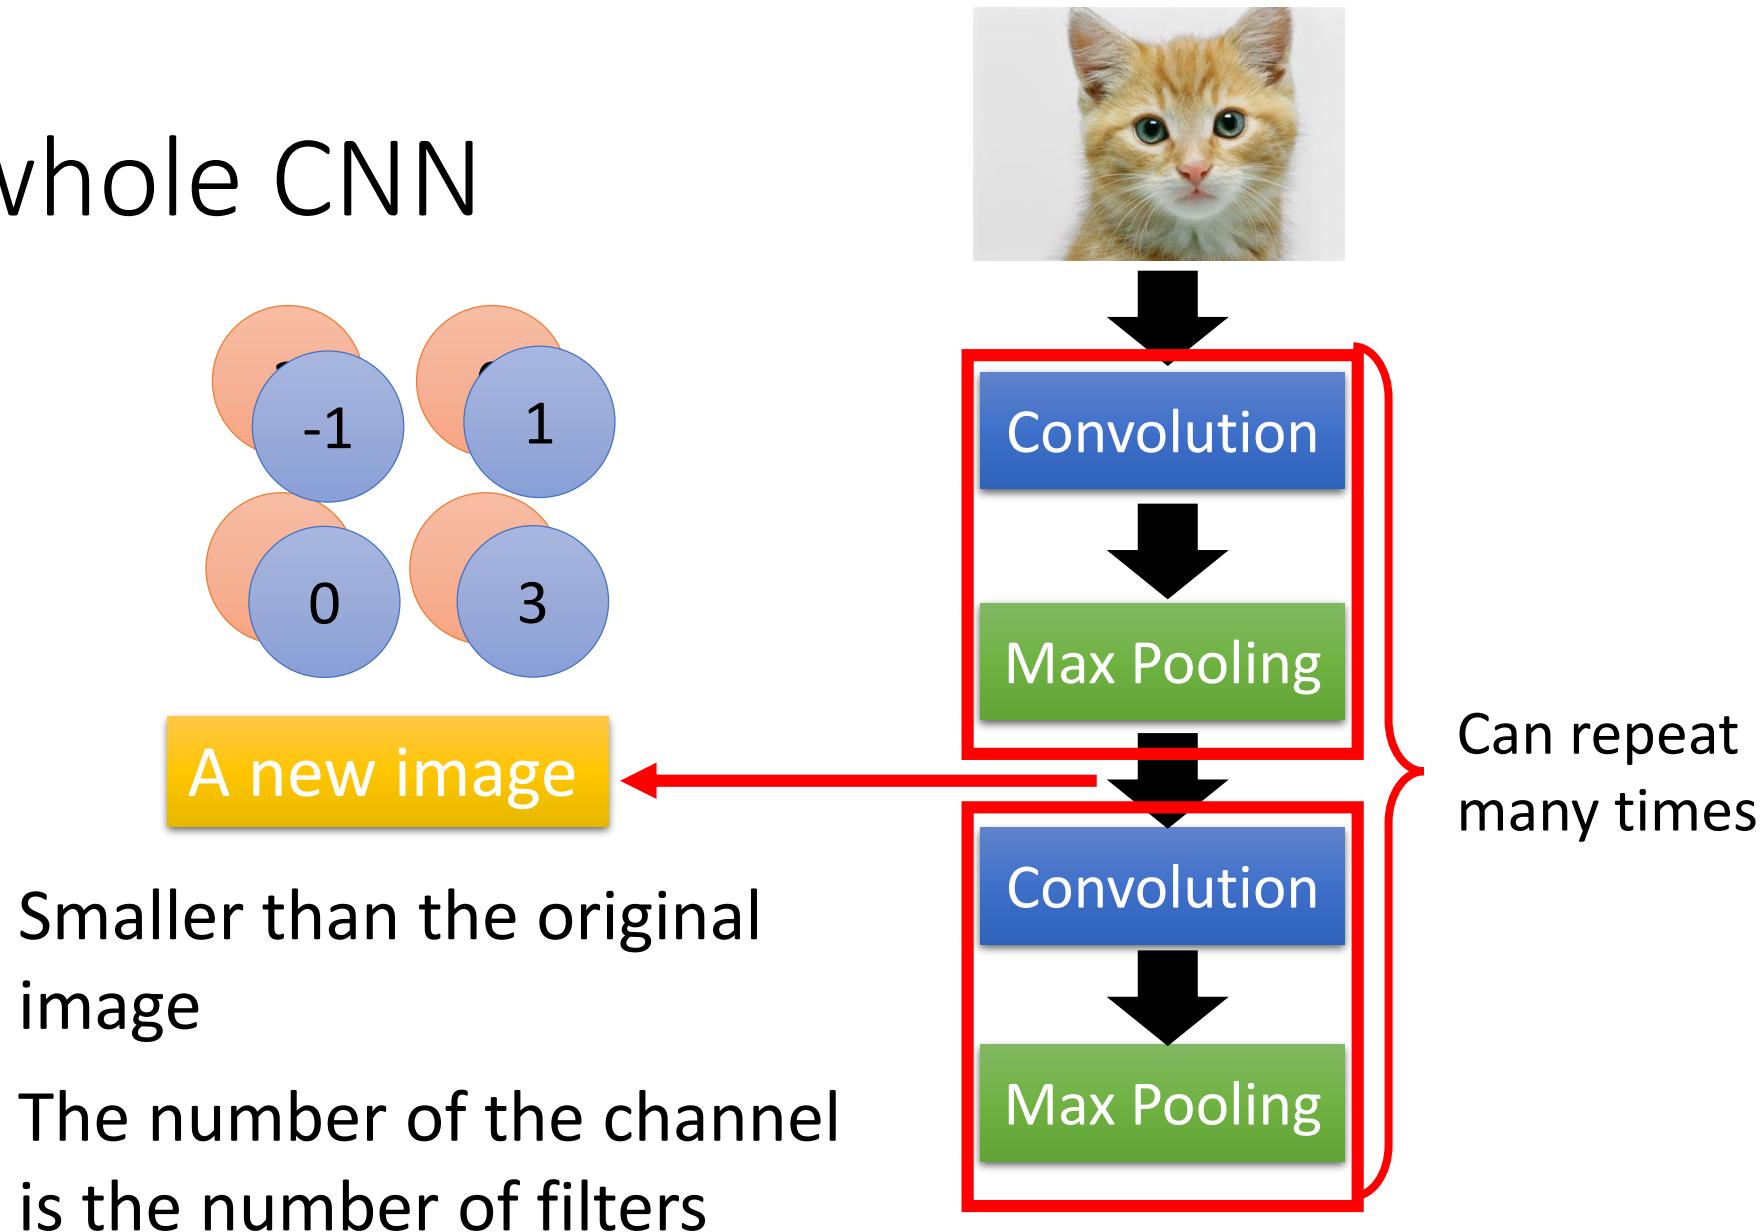
New image
but smaller



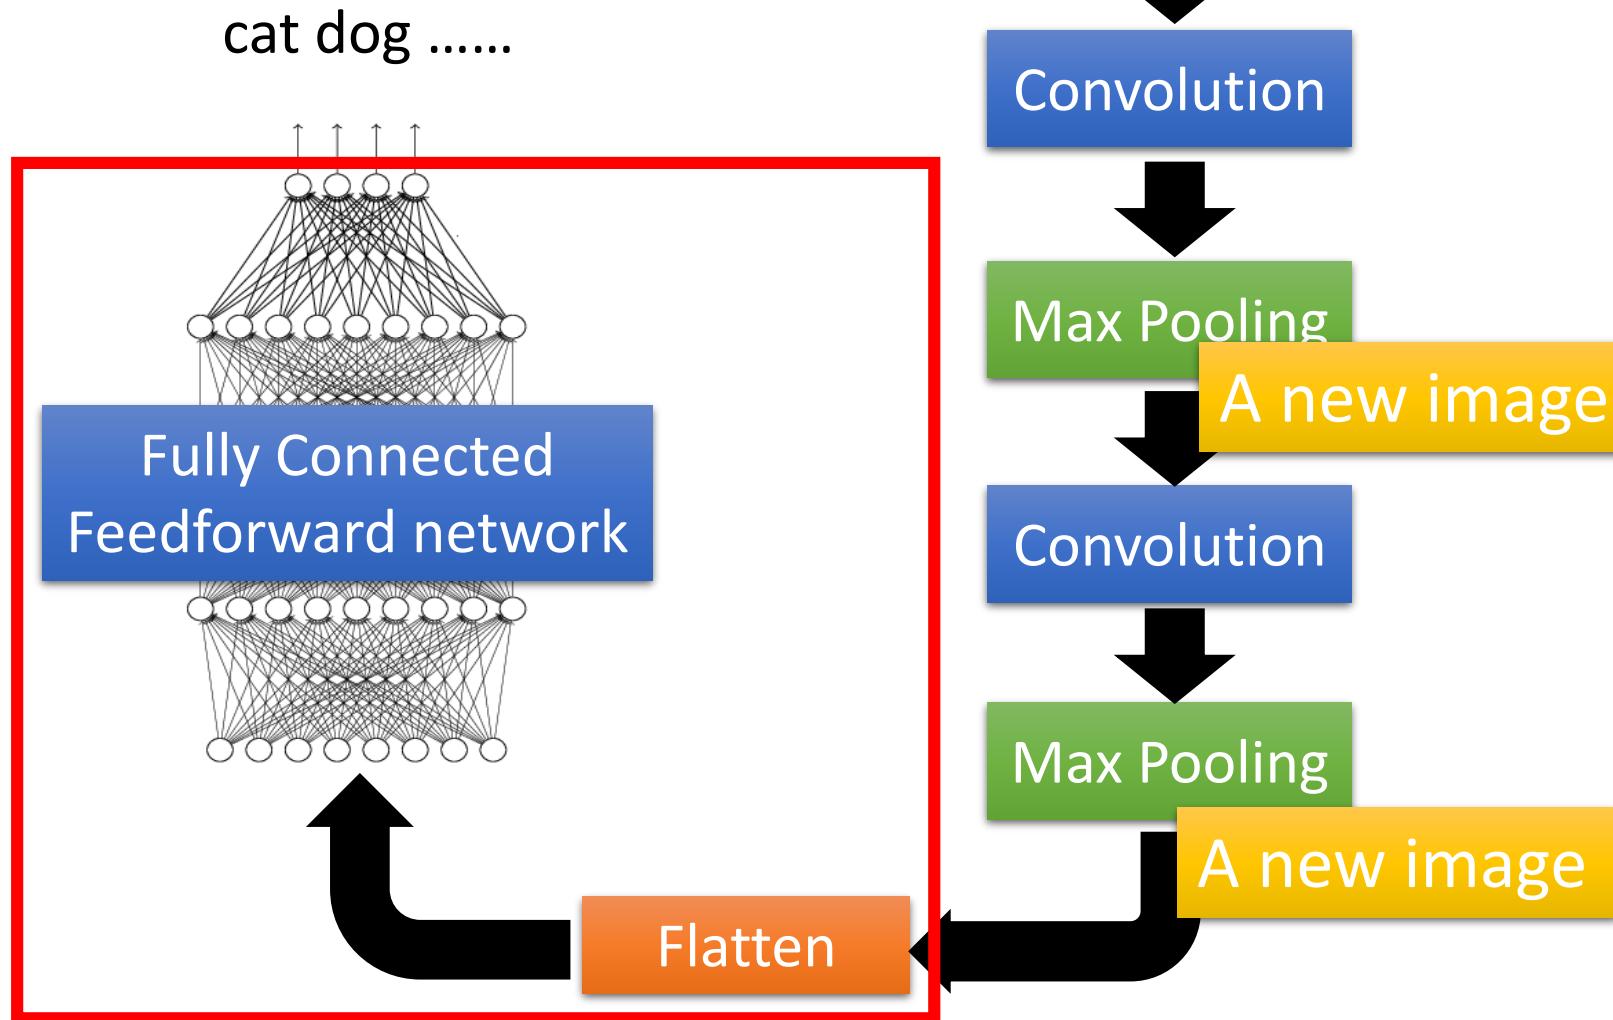
2 x 2 image

Each filter
is a channel

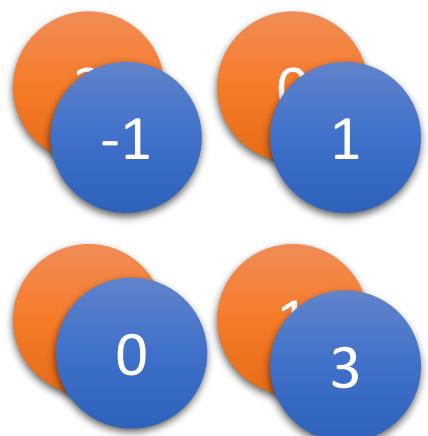
The whole CNN



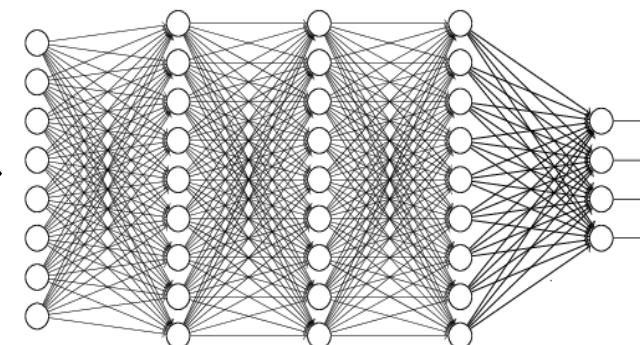
The whole CNN



Flatten



Flatten

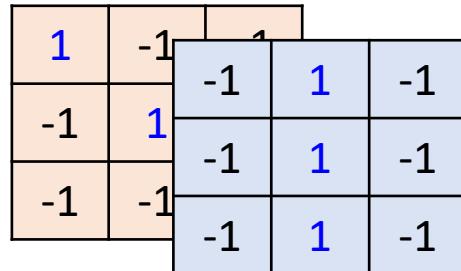


Fully Connected
Feedforward network

CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D tensor)*

```
model2.add( Convolution2D( 25, 3, 3,  
                           input_shape=(28, 28, 1)) )
```

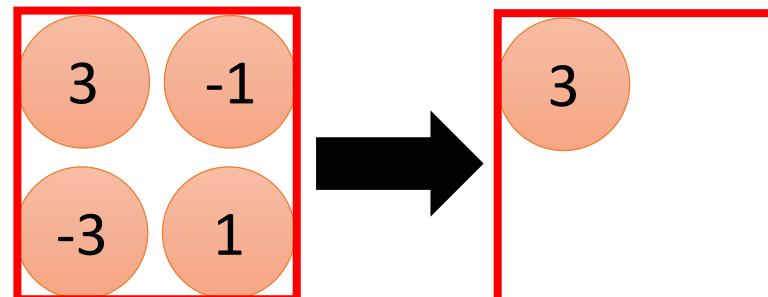


.....
There are 25
3x3 filters.

Input_shape = (28 , 28 , 1)

28 x 28 pixels 1: black/white, 3: RGB

```
model2.add(MaxPooling2D( (2, 2) ))
```



input

Convolution

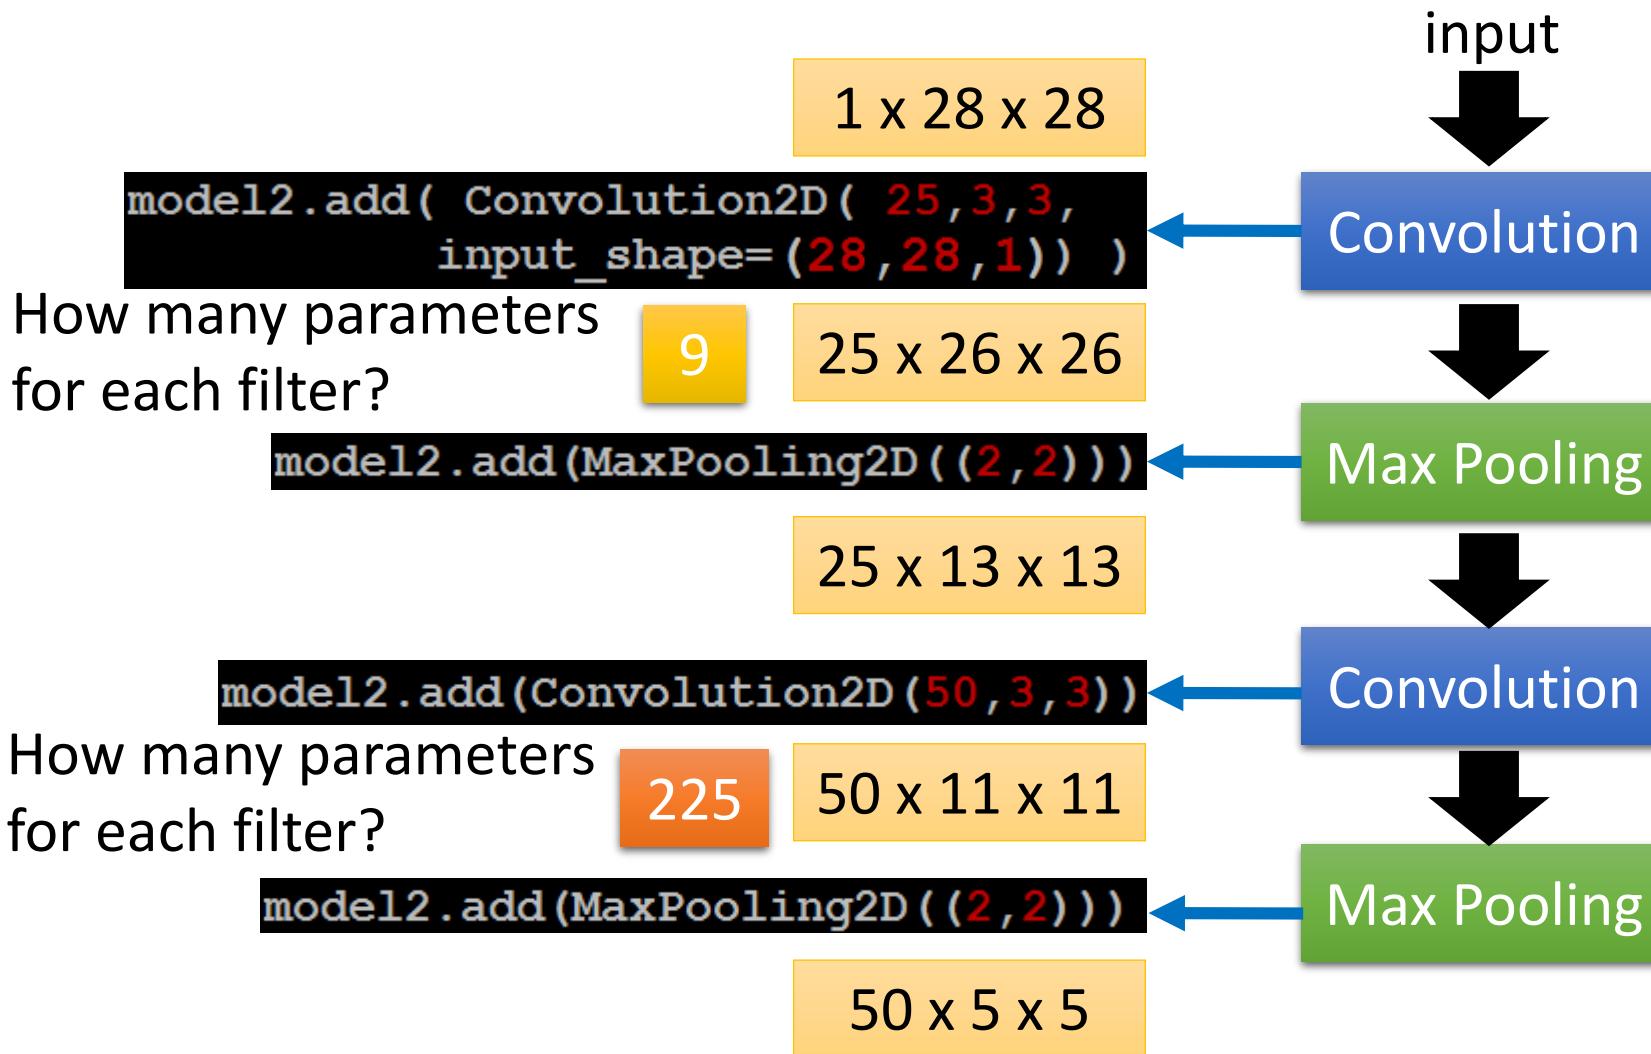
Max Pooling

Convolution

Max Pooling

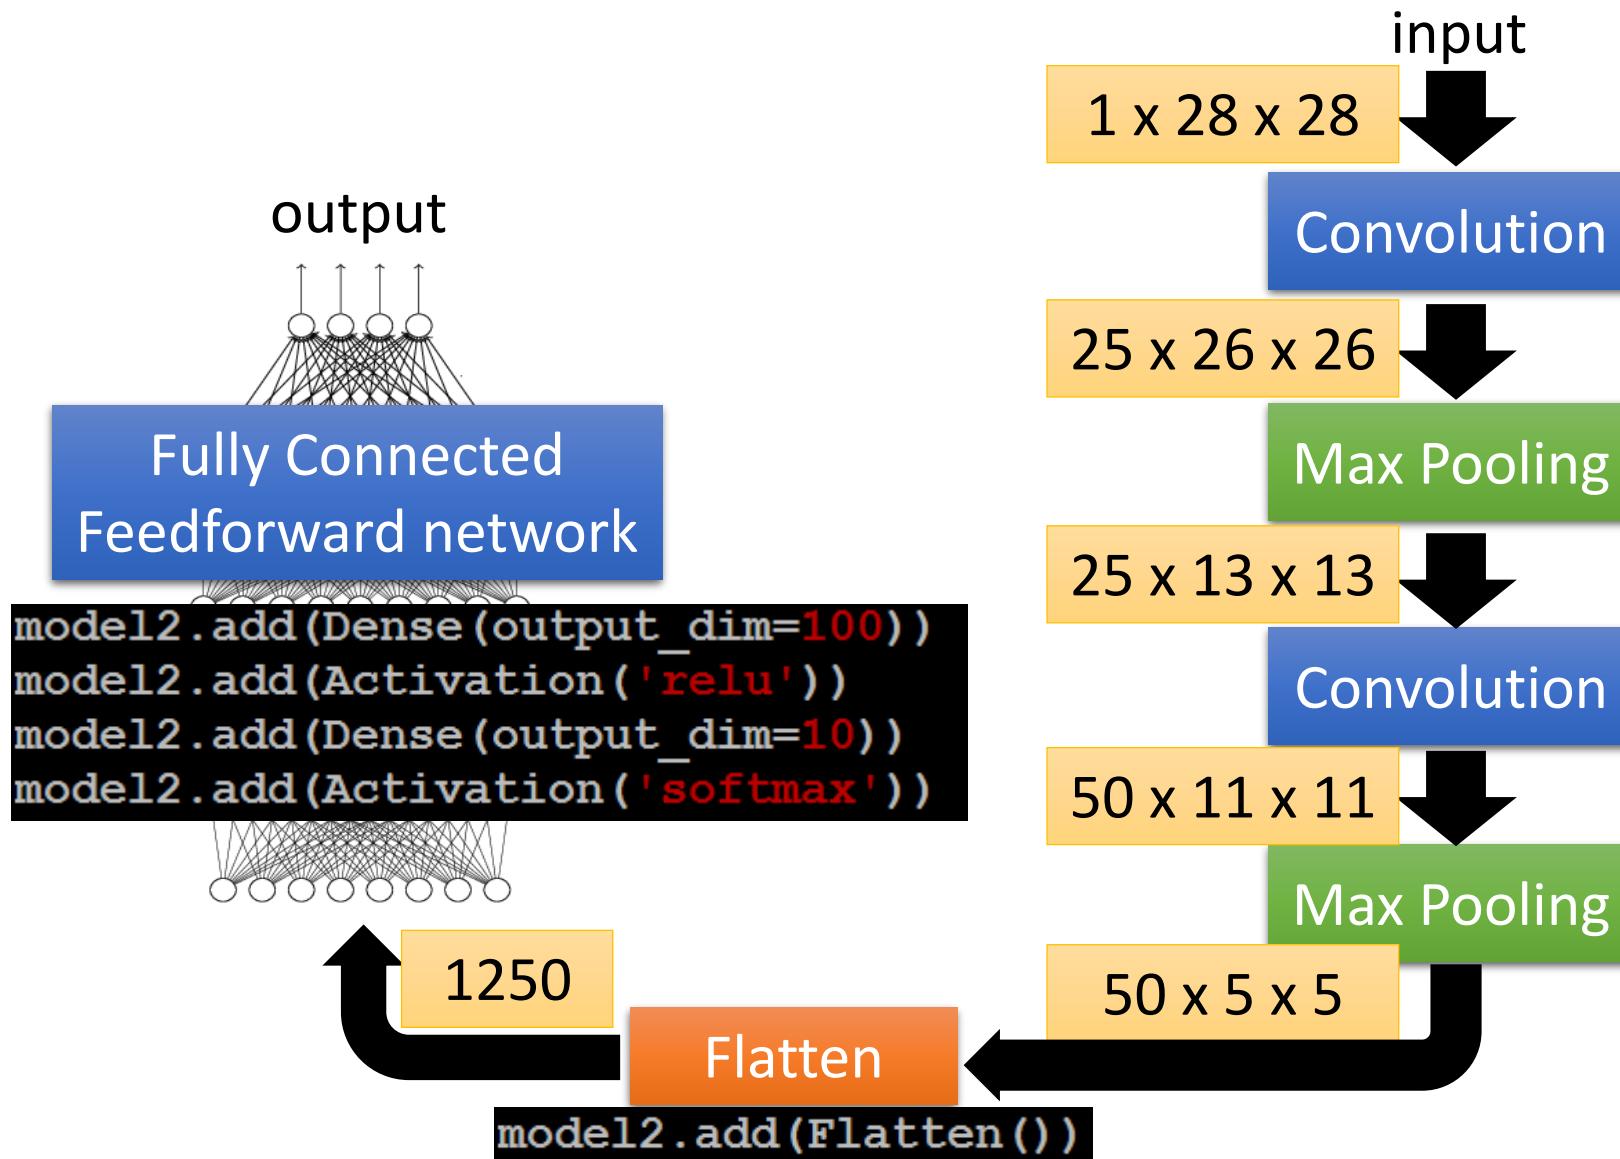
CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D tensor)*



CNN in Keras

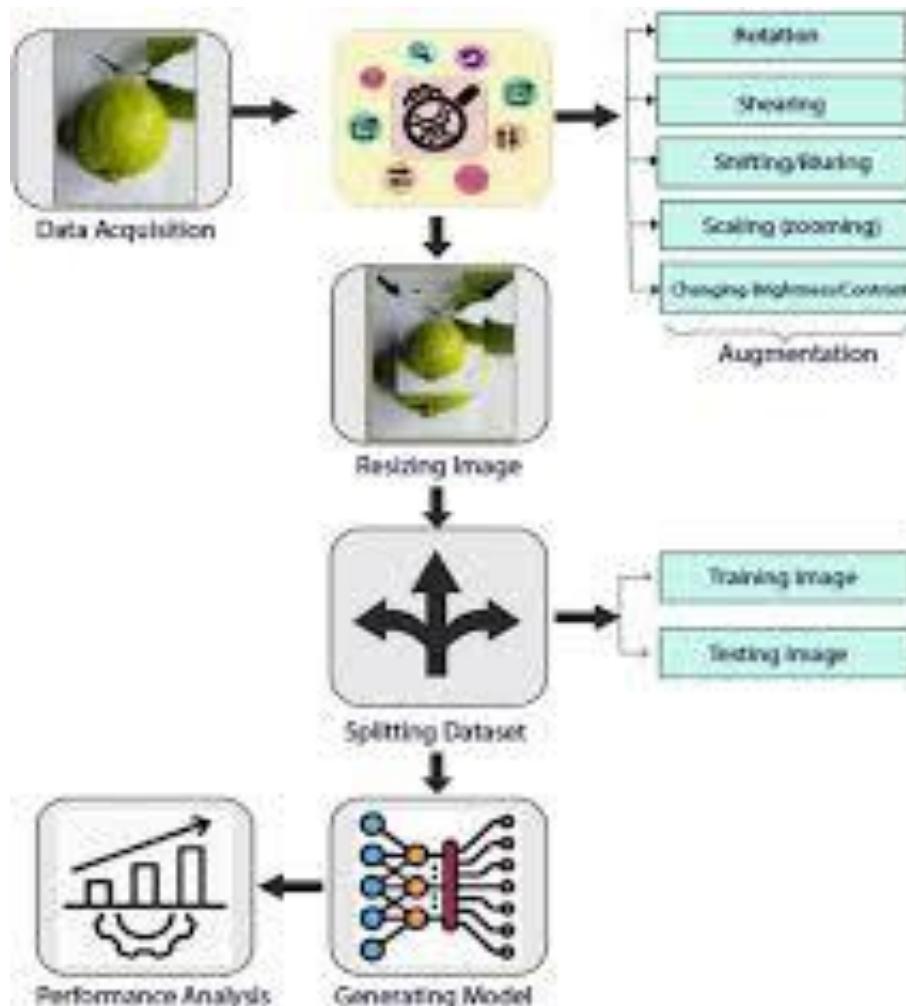
Only modified the ***network structure*** and
input format (vector -> 3-D tensor)



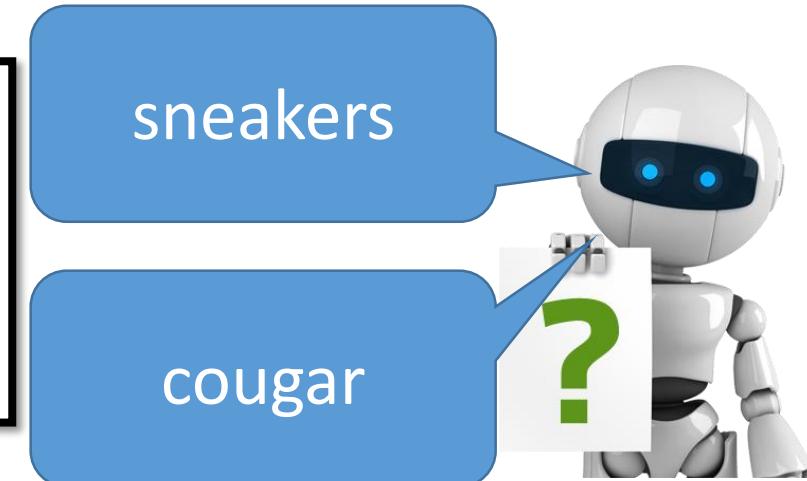
Next..



What does machine learn?



sneakers



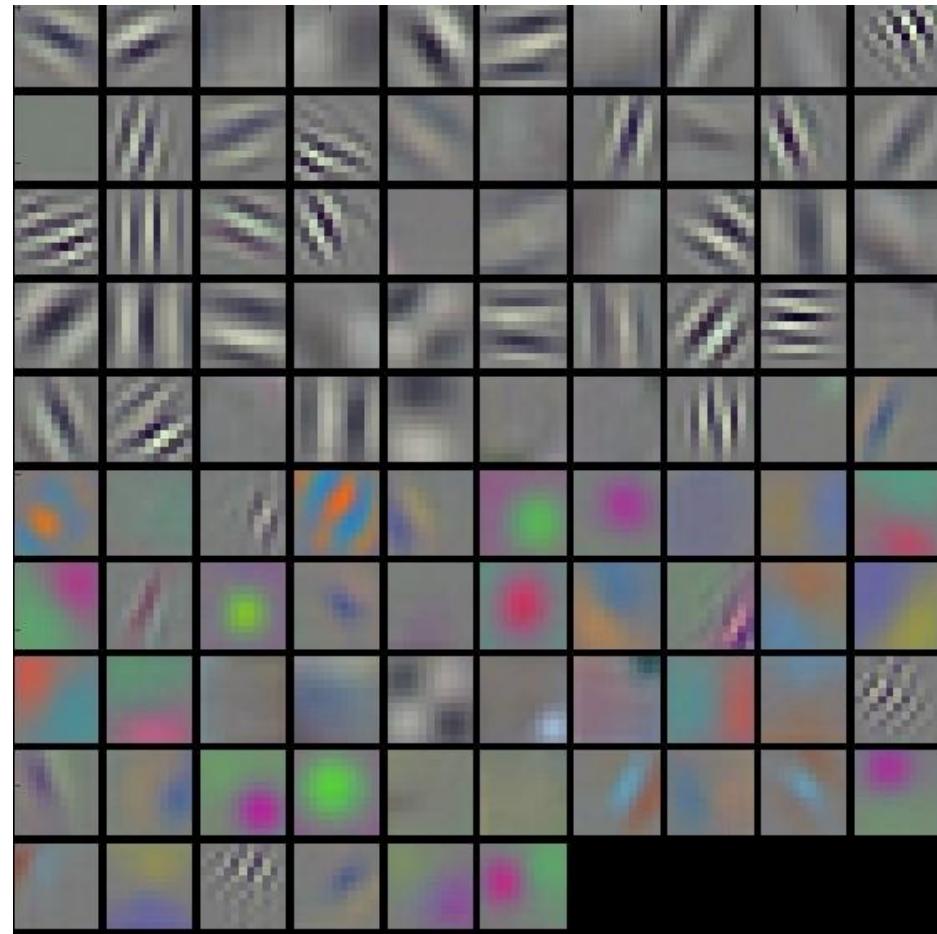
cougar

<http://newsneakernews.wpengine.netdna-cdn.com/wp-content/uploads/2016/11/rihanna-puma-creeper-velvet-release-date-02.jpg>

First Convolution Layer

- Typical-looking filters on the trained first layer

11 x 11
(AlexNet)

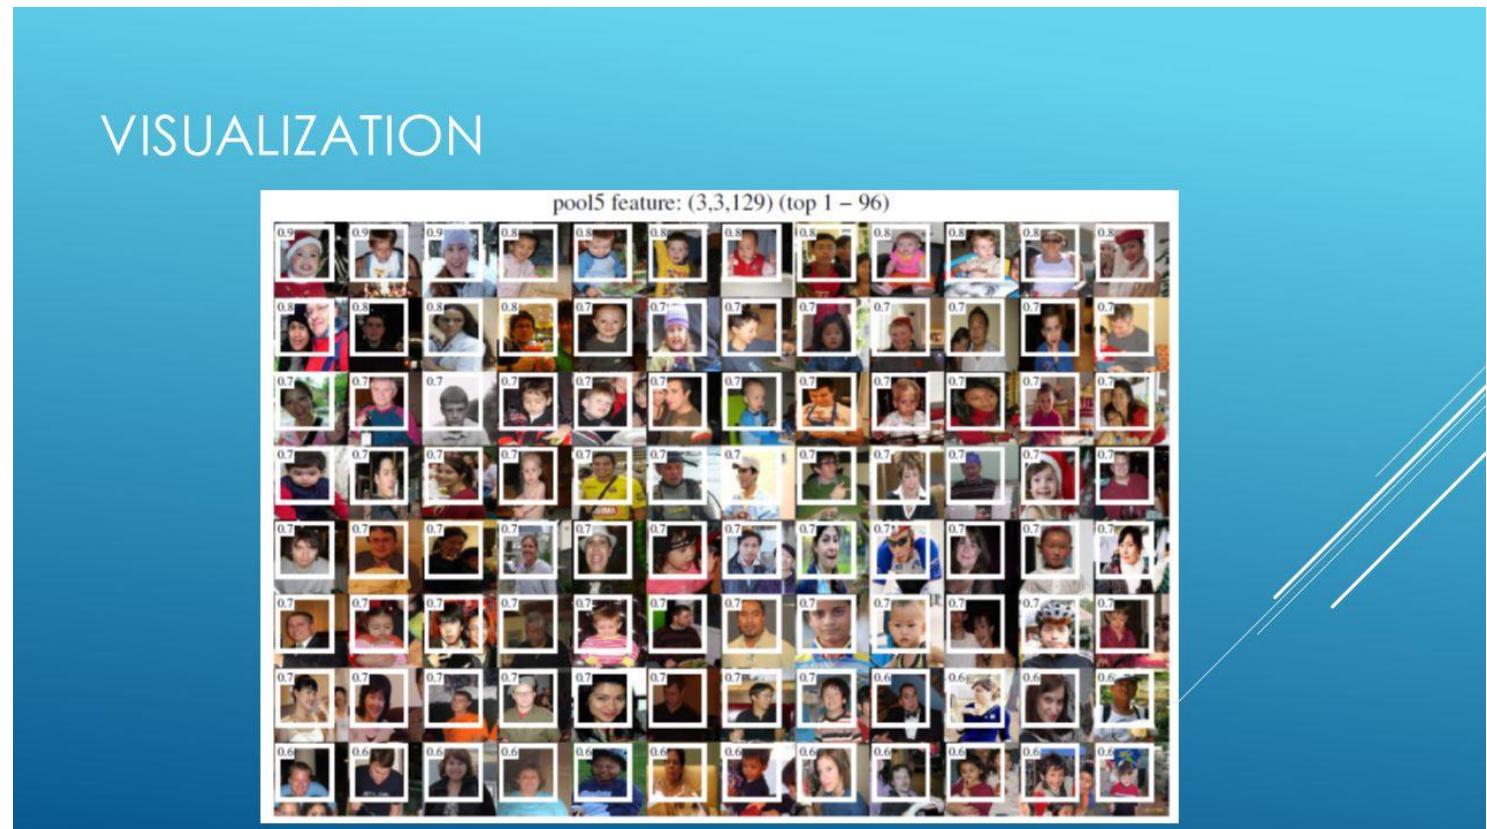


<http://cs231n.github.io/understanding-cnn/>

How about higher layers?

- Which images make a specific neuron activate

Rich feature Hierarchies for
Accurate object detection and
semantic segmentation Ross
Girshick, Jeff Donahue, Trevor
Darrell, Jitendra Malik (UC
Berkeley)

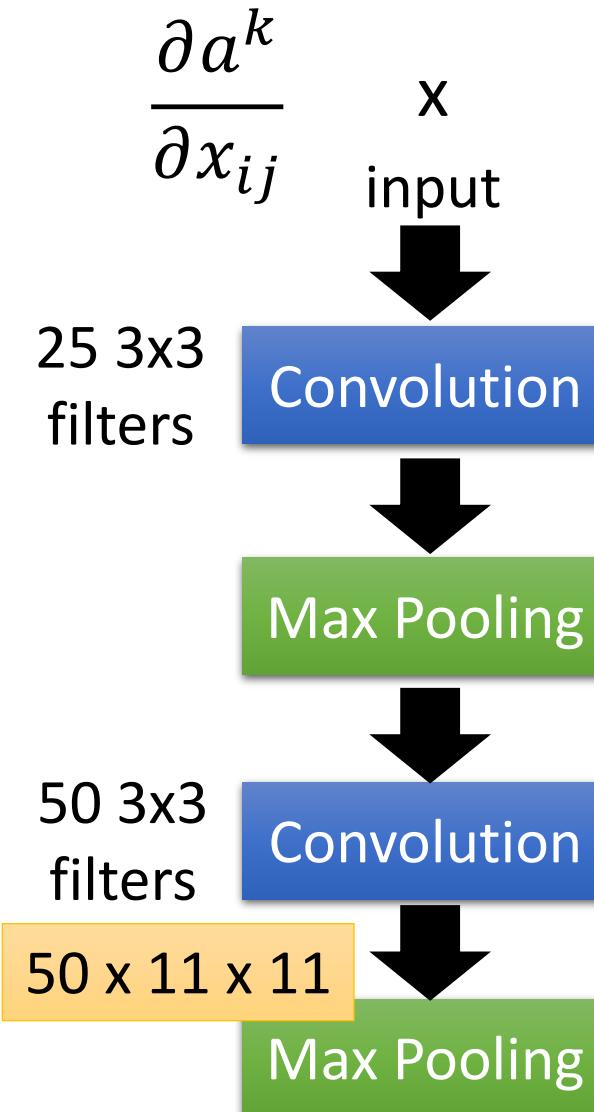
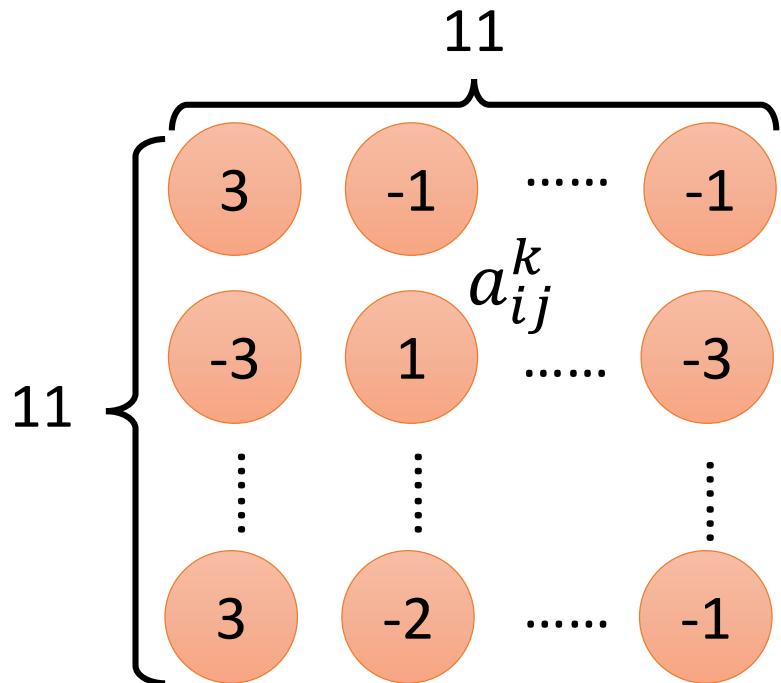


What does CNN learn?

The output of the k-th filter is a 11×11 matrix.

Degree of the activation of the k-th filter: $a^k = \sum_{i=1}^{11} \sum_{j=1}^{11} a_{ij}^k$

$$x^* = \arg \max_x a^k \text{ (gradient ascent)}$$

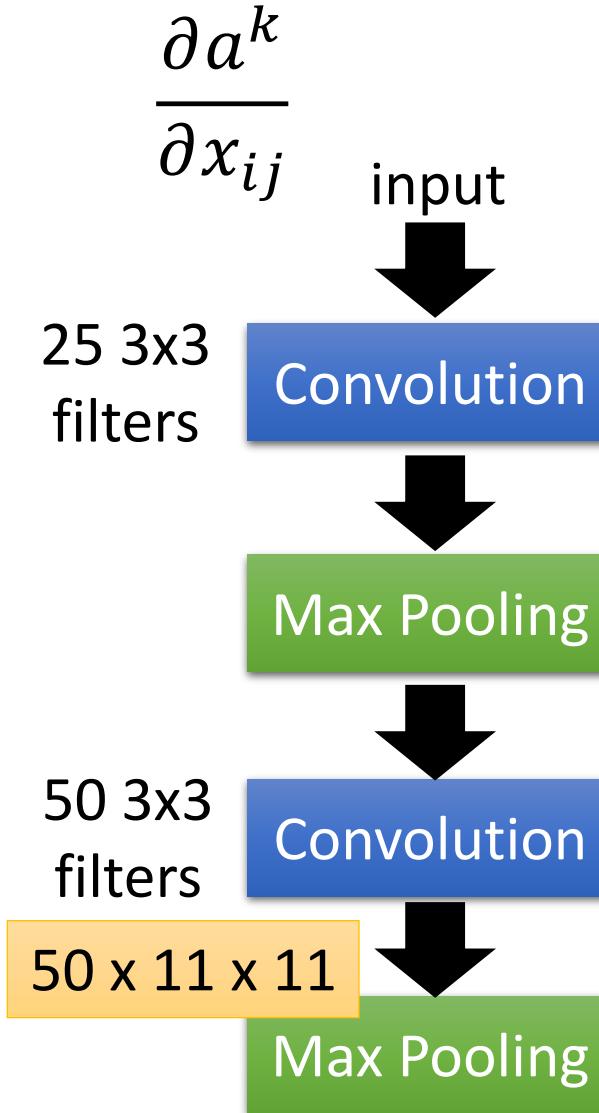
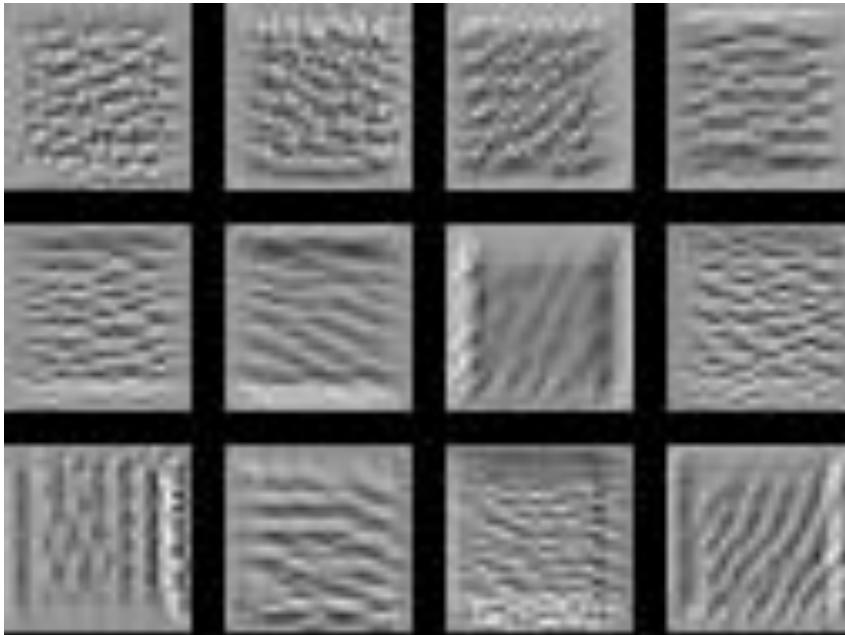


What does CNN learn?

The output of the k-th filter is a 11×11 matrix.

Degree of the activation of the k-th filter: $a^k = \sum_{i=1}^{11} \sum_{j=1}^{11} a_{ij}^k$

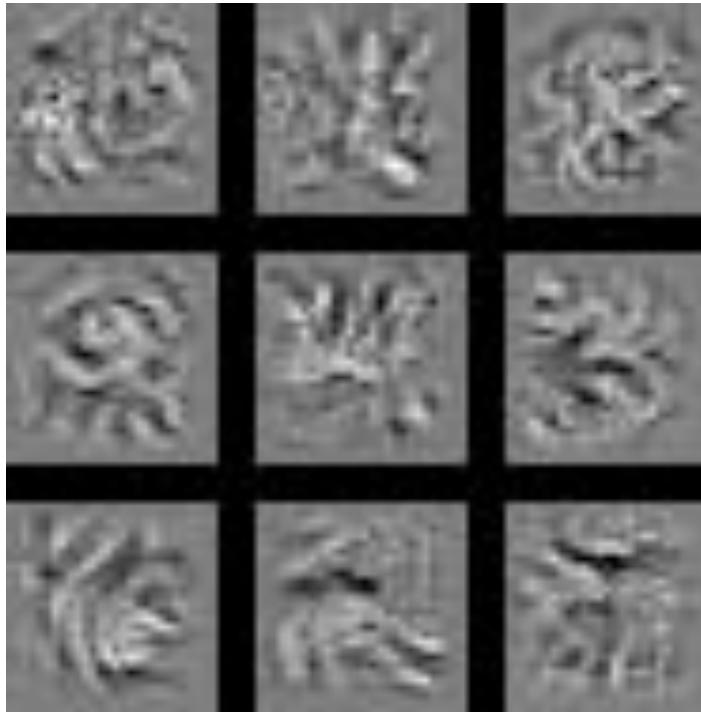
$$x^* = \arg \max_x a^k \text{ (gradient ascent)}$$



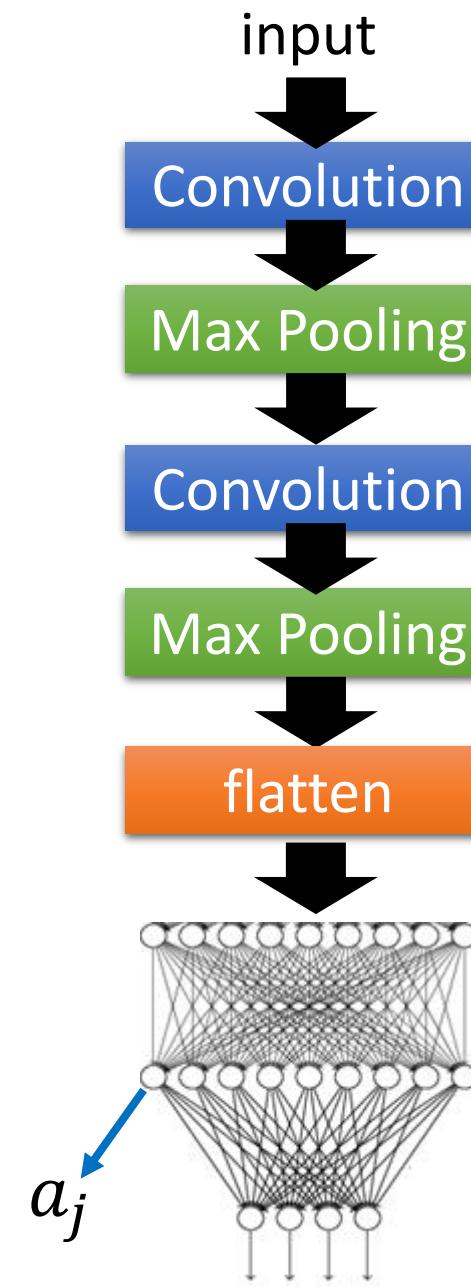
What does CNN learn?

Find an image maximizing the output of neuron:

$$x^* = \arg \max_x a_j$$



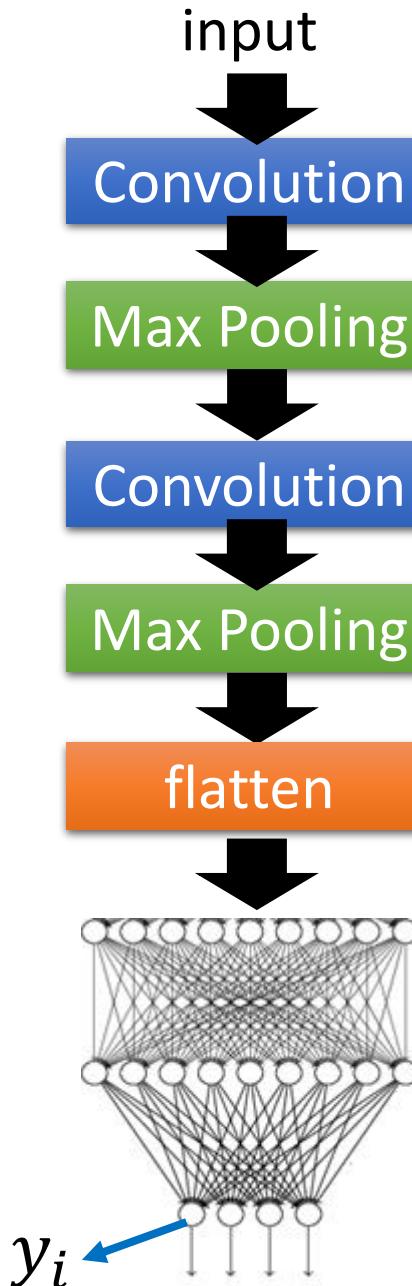
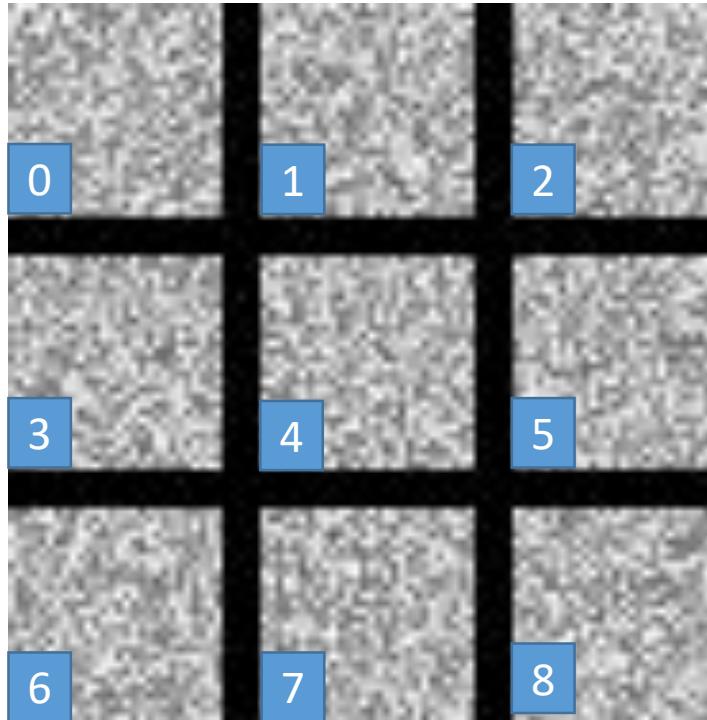
Each figure corresponds to a neuron



What does CNN learn?

$$x^* = \arg \max_x y^i$$

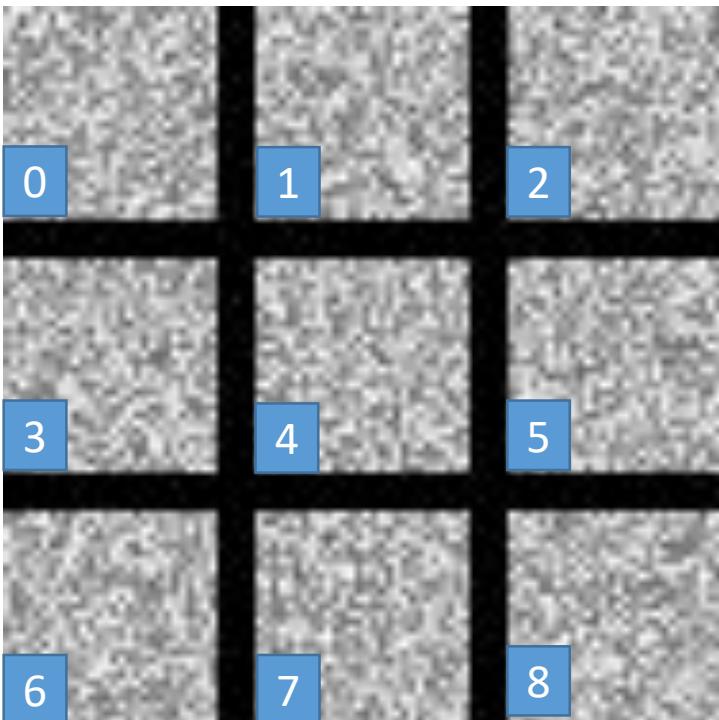
Can we see digits?



Deep Neural Networks are Easily Fooled
<https://www.youtube.com/watch?v=M2IebCN9Ht4>

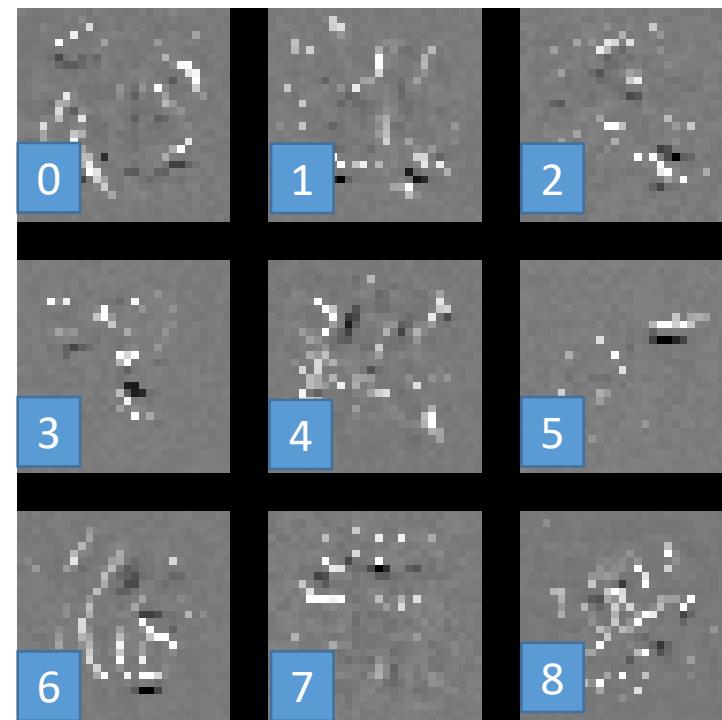
What does CNN learn?

$$x^* = \arg \max_x y^i$$



Over all
pixel values

$$x^* = \arg \max_x \left(y^i - \sum_{i,j} |x_{ij}| \right)$$





Karen Simonyan, Andrea Vedaldi, Andrew Zisserman, “Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps”, ICLR, 2014

$$\left| \frac{\partial y_k}{\partial x_{ij}} \right|$$

y_k : the predicted class of the model



Karen Simonyan, Andrea Vedaldi, Andrew Zisserman, "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps", ICLR, 2014

Review.....

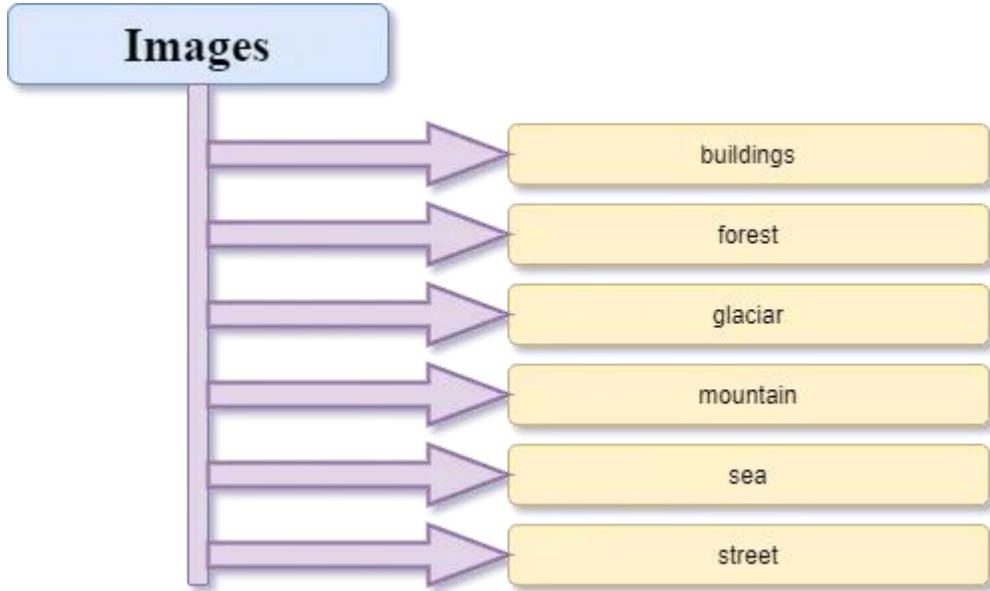
- Examples

Example 1:

CNN Model With PyTorch For Image Classification

- Preparing The Dataset
- Splitting Data and Prepare Batches
- Base Model For Image Classification
- Convolution, Padding, Stride, Pooling
- CNN Model For Classification
- Hyperparameters, Model Training, And Evaluation

Preparing the Dataset :



```
import torch
import torchvision
from torchvision import transforms
from torchvision.datasets import ImageFolder
```

```
#train and test data directory
data_dir = "../input/intel-image-
classification/seg_train/seg_train/"
test_data_dir = "../input/intel-image-
classification/seg_test/seg_test"
```

```
#load the train and test data
dataset = ImageFolder(data_dir,transform =
transforms.Compose([
    transforms.Resize((150,150)),transforms.ToTensor()
]))
test_dataset =
ImageFolder(test_data_dir,transforms.Compose([
    transforms.Resize((150,150)),transforms.ToTensor()
]))
```


CNN Model With PyTorch

The `torchvision.transforms` module provides various functionality to preprocess the images, here first we resize the image for $(150*150)$ shape and then transforms them into tensors.

```
img, label = dataset[0]
print(img.shape,label)
```

#output :

```
#torch.Size([3, 150, 150]) 0
```

first image in the dataset has a shape $(3,150,150)$ which means the image has 3 channels (RGB), height 150, and width 150. The image has a label 0, which represents the “buildings” class.

The image label set according to the class index in
`print("Follwing classes are there : \n",dataset.classes)`

#output:

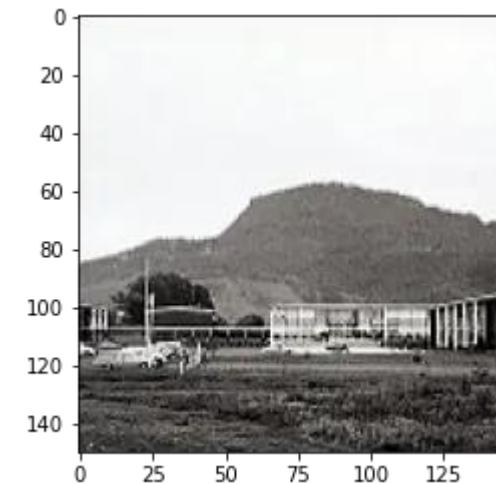
```
#Follwing classes are there :
# ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

dataset has 6 types of images in the dataset. dataset consists of images in form of *Tensors*, `imshow()` method of `matplotlib` python library can be used to visualize images.

```
def display_img(img,label):
    print(f"Label : {dataset.classes[label]}")
    plt.imshow(img.permute(1,2,0))
```

```
#display the first image in the dataset
display_img(*dataset[0])
```

permute method reshapes the image from $(3,150,150)$ to $(150,150,3)$



Splitting Data and Prepare Batches

```
from torch.utils.data.dataloader import DataLoader
from torch.utils.data import random_split

batch_size = 128
val_size = 2000
train_size = len(dataset) - val_size

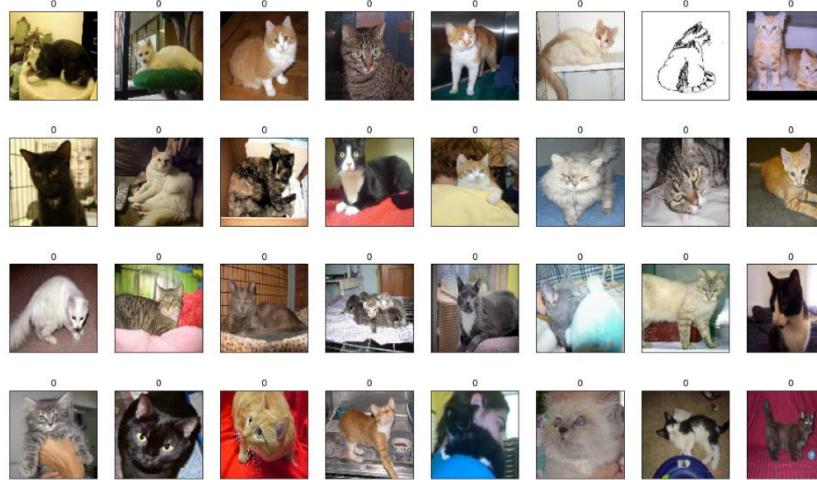
train_data, val_data =
random_split(dataset,[train_size, val_size])
print(f"Length of Train Data : {len(train_data)}")
print(f"Length of Validation Data : {len(val_data)}")

#output
#Length of Train Data : 12034
#Length of Validation Data : 2000

#load the train and validation into batches.
train_dl = DataLoader(train_data, batch_size, shuffle = True,
num_workers = 4, pin_memory = True)
val_dl = DataLoader(val_data, batch_size*2, num_workers = 4,
pin_memory = True)
```

CNN Model With PyTorch

Visualizing the images



```
from torchvision.utils import make_grid  
import matplotlib.pyplot as plt
```

```
def show_batch(dl):  
    """Plot images grid of single batch"""  
    for images, labels in dl:  
        fig,ax = plt.subplots(figsize = (16,12))  
        ax.set_xticks([])  
        ax.set_yticks([])  
        ax.imshow(make_grid(images,nrow=16).permute(1,2,0))  
        break  
  
show_batch(train_dl)
```

CNN Model With PyTorch

Base Model For Image Classification

```
import torch.nn as nn  
import torch.nn.functional as F
```

```
class ImageClassificationBase(nn.Module):  
  
    def training_step(self, batch):  
        images, labels = batch  
        out = self(images)          # Generate predictions  
        loss = F.cross_entropy(out, labels) # Calculate loss  
        return loss  
  
    def validation_step(self, batch):  
        images, labels = batch  
        out = self(images)          # Generate predictions  
        loss = F.cross_entropy(out, labels) # Calculate loss  
        acc = accuracy(out, labels)      # Calculate accuracy  
        return {'val_loss': loss.detach(), 'val_acc': acc}
```

```
def validation_epoch_end(self, outputs):  
    batch_losses = [x['val_loss'] for x in outputs]  
    epoch_loss = torch.stack(batch_losses).mean() # Combine losses  
    batch_accs = [x['val_acc'] for x in outputs]  
    epoch_acc = torch.stack(batch_accs).mean() # Combine accuracies  
    return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}
```

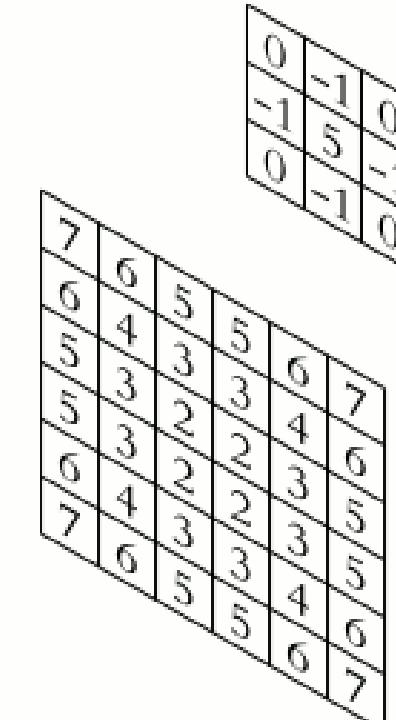
```
def epoch_end(self, epoch, result):  
    print("Epoch [{}], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(  
        epoch, result['train_loss'], result['val_loss'],  
        result['val_acc']))
```

Convolution, Padding, Stride, Maxpooling

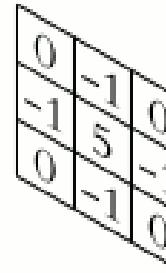
3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

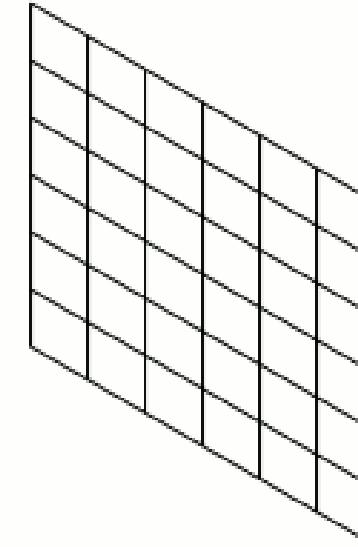
Convolution



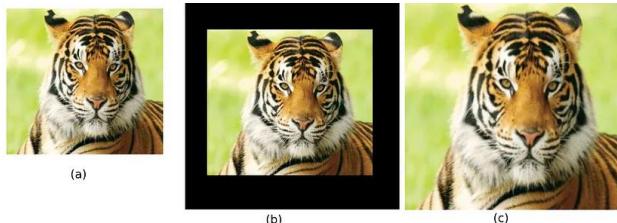
input



output



Convolution, Padding, Stride, Maxpooling

A small image (a) resized to 180×180 pixels using zero-padding (b) and interpolation (c)

$$((W - F * 2P)/S) + 1 \text{ and } ((H - F * 2P)/S) + 1$$

n = number of pixels of an input image

p = padding number

f = filter size

from these $(n*n) + (p*p)$ convolution $(f*f)$ gives $\rightarrow (n+2p-f+1) * (n+2p-f+1)$

Zerro Padding

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

0	-1	0
-1	5	-1
0	-1	0

114			

Convolution, Padding, Stride, Maxpooling

Stride

Strided convolution

2	3	3	4	7	4	4	6	2	9
6	1	6	0	9	2	8	7	4	3
3	-1	4	0	8	3	3	8	9	7
7	8	3	6	6	6	3	4		
4	2	1	8	3	4	6			
3	2	4	1	9	8	3			
0	1	3	9	2	1	4			

7x7

*

3	4	4
1	0	2
-1	0	3

3x3

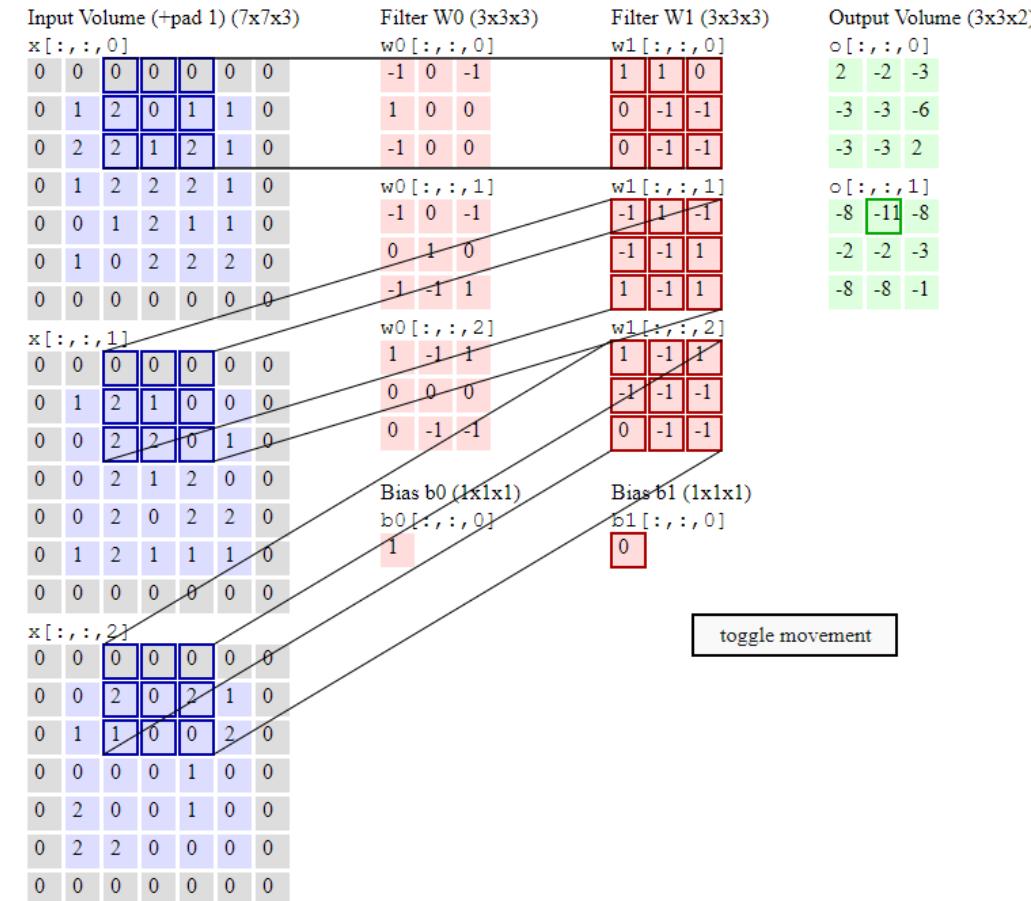
stride = 2

=

q1		

Convolution, Padding, Stride, Maxpooling

MaxPooling



CNN Model With PyTorch

```
Model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), activation='relu',  
input_shape=(32, 32, 3)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model.summary()
```

model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
<hr/>		
<hr/>		
Total params: 56,320		
Trainable params: 56,320		
Non-trainable params: 0		

```
model.compile(optimizer='adam',  
loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits  
=True),  
metrics=['accuracy'])  
  
history = model.fit(train_images, train_labels, epochs=10,  
validation_data=(test_images, test_labels))
```

Epoch 1/10

1563/1563 [=====] - 8s 4ms/step - loss: 1.4971 - accuracy: 0.4553 - val_loss: 1.2659 -
val_accuracy: 0.5492

Epoch 2/10

1563/1563 [=====] - 6s 4ms/step - loss: 1.1424 - accuracy: 0.5966 - val_loss: 1.1025 -
val_accuracy: 0.6098

Epoch 3/10

1563/1563 [=====] - 6s 4ms/step - loss: 0.9885 - accuracy: 0.6539 - val_loss: 0.9557 -
val_accuracy: 0.6629

Epoch 4/10

1563/1563 [=====] - 6s 4ms/step - loss: 0.8932 - accuracy: 0.6878 - val_loss: 0.8924 -
val_accuracy: 0.6935

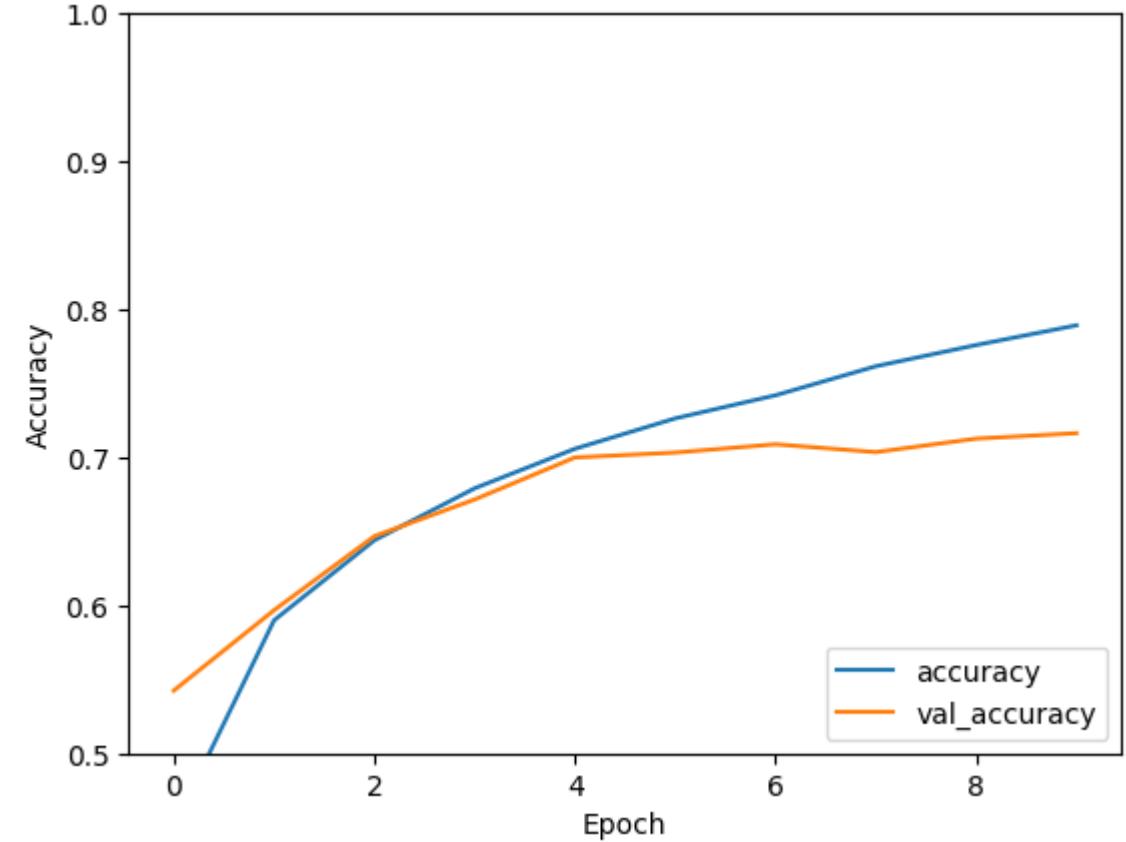
Epoch 5/10

1563/1563 [=====] - 6s 4ms/step - loss: 0.8222 - accuracy: 0.7130 - val_loss: 0.8679

Evaluasi model

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels,
verbose=2)
```



Example 2:

Flower Classification with CNN

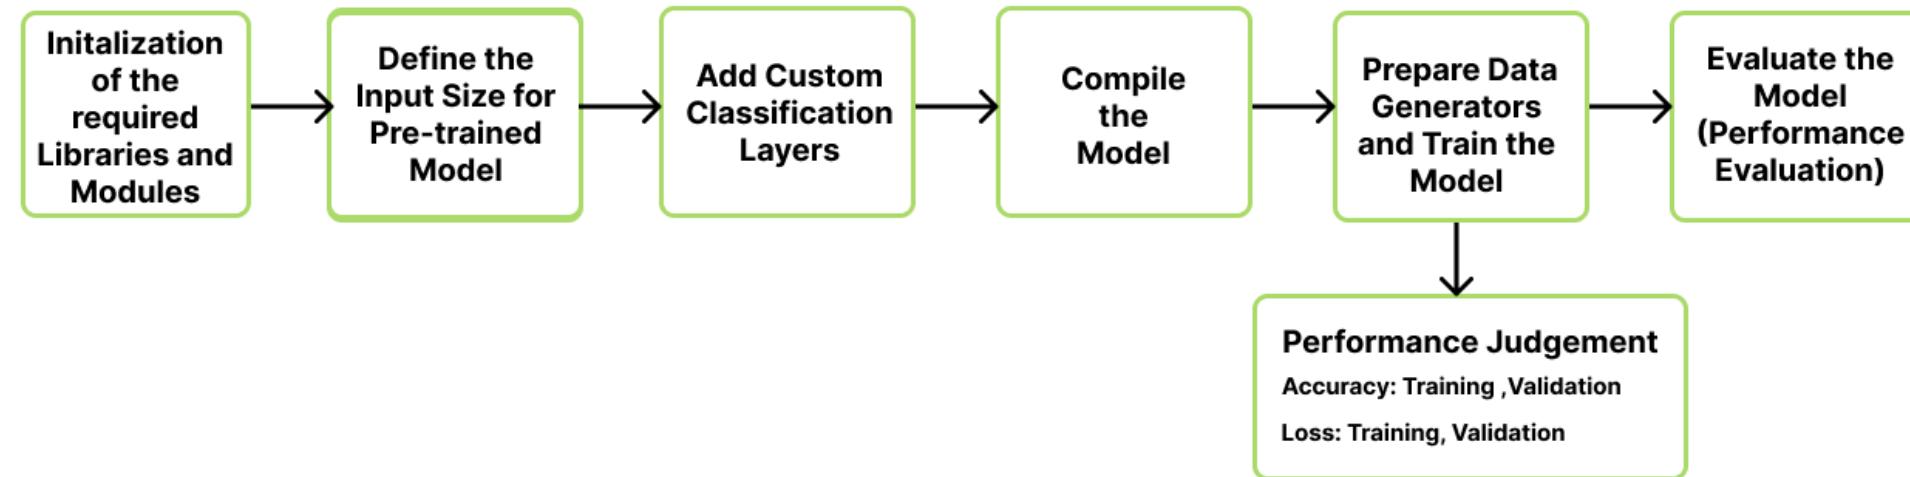


<https://github.com/siagianp/KerasTF-Flower-Classification-Convolutional-Neural-Network>

CNN Model With KTF

1. Importing Libraries
2. Data Collection and Preparation
3. Utilize Pre-trained Architectures (Xception)
4. Feature Extraction and Model Architecture
5. Data Augmentation and Preprocessing
6. Model Training
7. Model Evaluation
8. User Interaction and Engagement
9. Deployment and Integration

CNN Model With KTF



CNN Model With KTF

flower_images (5 directories)

... >

About this directory

The directory consists of five flower sub-directories; **Lilly, Lotus, Orchid, Sunflower and Tulip**. Each sub-directory has **1000 images** of the respective flowers.

-  Lilly
1000 files
-  Lotus
1000 files
-  Orchid
1000 files
-  Sunflower
1000 files
-  Tulip
1000 files

Data Explorer

261.51 MB

- flower_images
 - Lilly
 - Lotus
 - Orchid
 - Sunflower
 - Tulip

Summary

- 5000 files

CNN Model With KTF

```
# Setting the data
```

```
"""
```

```
import os  
import numpy as np  
import matplotlib.pyplot as plt  
import cv2  
import pickle
```

```
data_dir = "../data/train/"
```

```
categories = ["daisy", "dandelion", "rose", "sunflower", "tulip"]
```

```
data = []
```

```
def make_data():
```

```
    for category in categories:
```

```
        path = os.path.join(data_dir, category)
```

```
        label = categories.index(category)
```

```
        for img_name in os.listdir(path):
```

Example 3:

Pytorch : Cat Dog Classification with CNN

 checkpoint	commit2
 sample/sample	Add files via upload
 .gitattributes	Initial commit
 README.md	Update README.md
 preprocessing.py	commit2
 testSample.py	Add files via upload
 train.py	commit2

Checkpoint: epoch60_93.pt
sample/sample : flower1.jpg

CNN Model With PyTorch

Preprocessing.py

```
import torchvision.datasets as dset  
import torchvision.transforms as transforms  
import torch
```

batch_size = 64

```
def trainDataset():
```

```
dataset = dset.ImageFolder(root="dataset/train",
                           transform=transforms.Compose([
                               transforms.RandomRotation(30),
                               transforms.RandomHorizontalFlip(),
                               transforms.RandomResizedCrop(224),
                               transforms.ToTensor(),
                               transforms.Normalize([0.485, 0.45, 0.406], [0.229, 0.224, 0.225])]))
```

scale=(0.96, 1.0), ratio=(0.95, 1.05)),

`transforms.ToTensor(),`

`transforms.Normalize([0.485, 0.456, 0.406],`

[0.229, 0.224, 0.225])

1))

##Test Sample.py

```
import torchvision.datasets as dset
import torchvision.transforms as transforms
import torch
import numpy
from torchvision import models
```

```
dataset = dset.ImageFolder(root="sample",
                           transform=transforms.Compose([
                               transforms.Resize([224, 224]),
                               transforms.ToTensor(),
                               transforms.Normalize([0.485, 0.456, 0.406],
[0.229, 0.224, 0.225]),
                           ]))
```

```
dataloader = torch.utils.data.DataLoader(dataset,  
6, 0.406], batch_size=1,  
shuffle=False,  
num_workers=0)
```

##train dir

```
from prep  
import tim
```

```
import torch  
import torch  
import torch  
from torch import  
import torch  
from torch
```

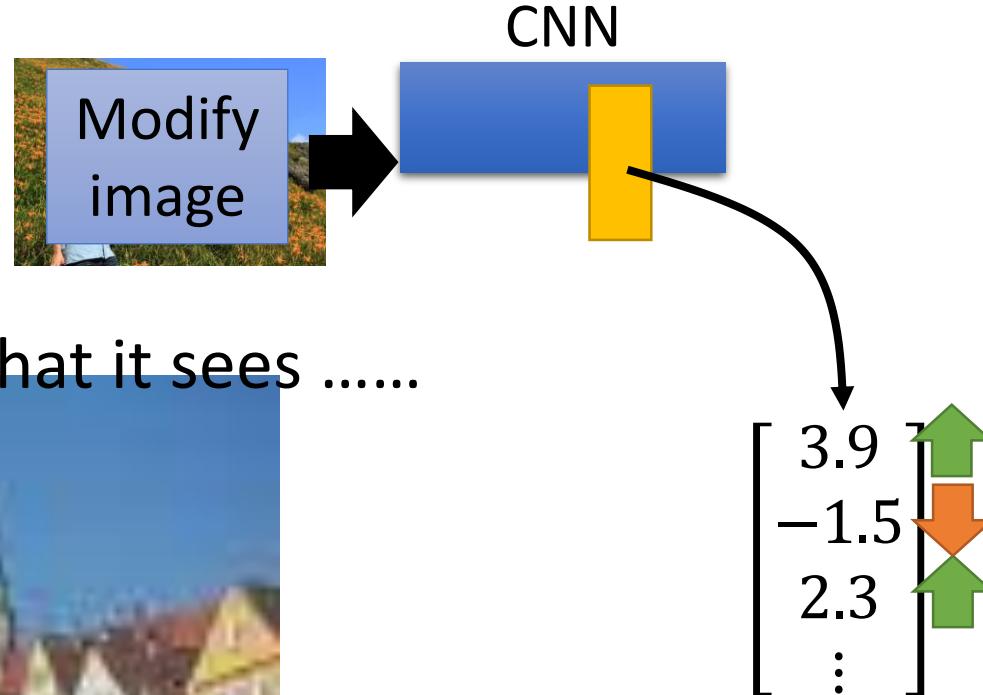
$$\text{lr} = 0.003$$

device = tc

train_load
test_load

Deep Dream

- Given a photo, machine adds what it sees



Deep Dream

- Given a photo, machine adds what it sees



Deep Style

- Given a photo, make its style like famous paintings



<https://dreamscopeapp.com/>

Deep Style

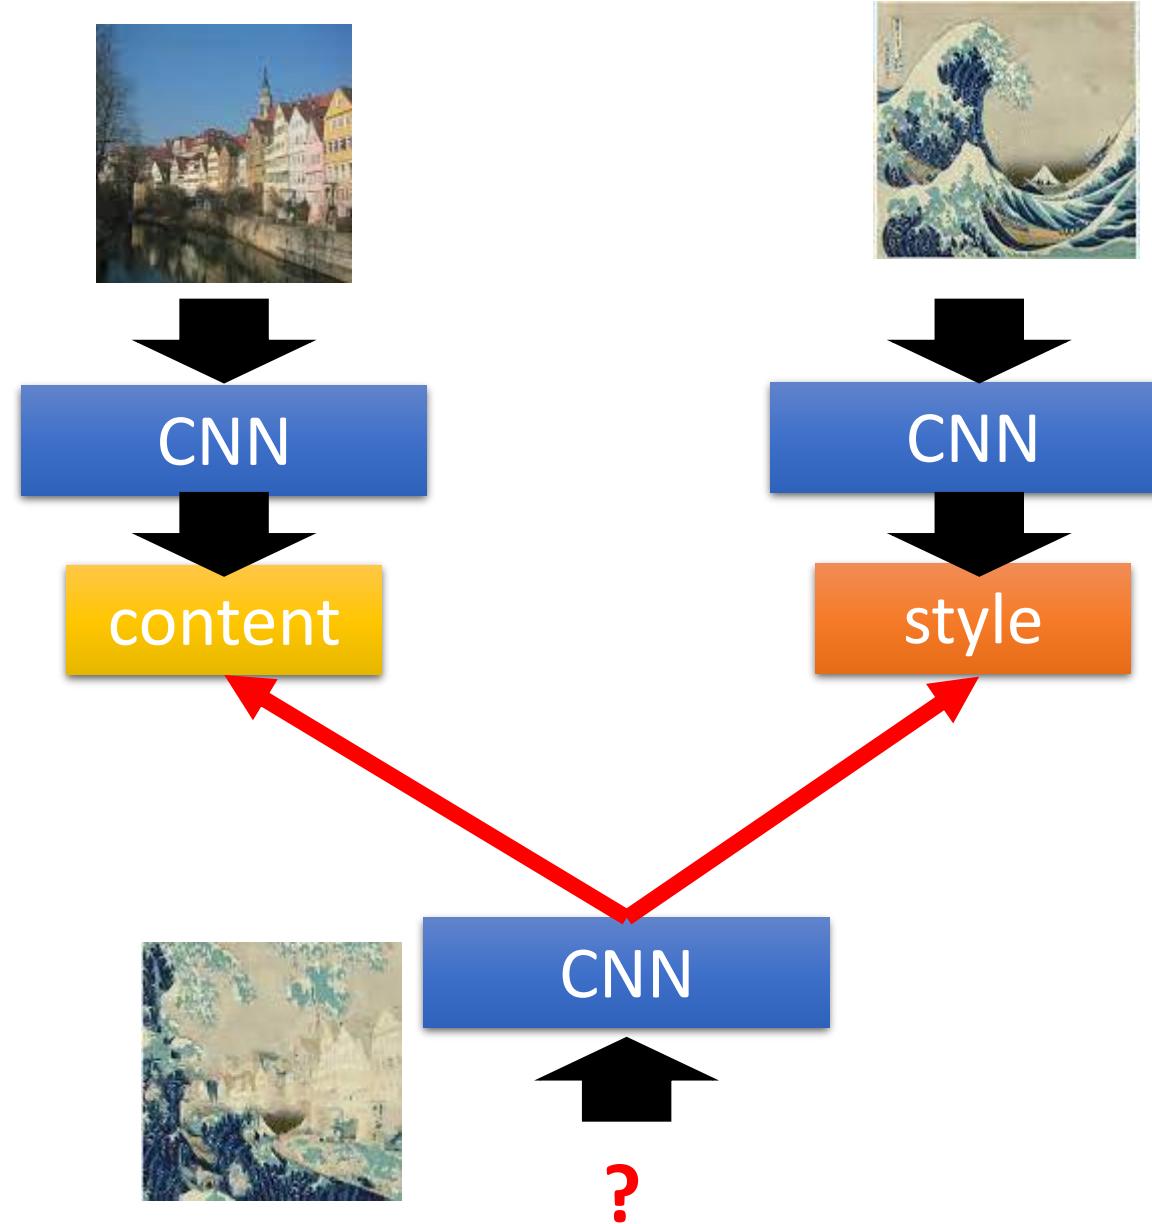
- Given a photo, make its style like famous paintings



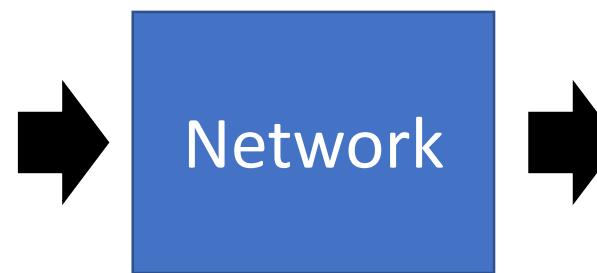
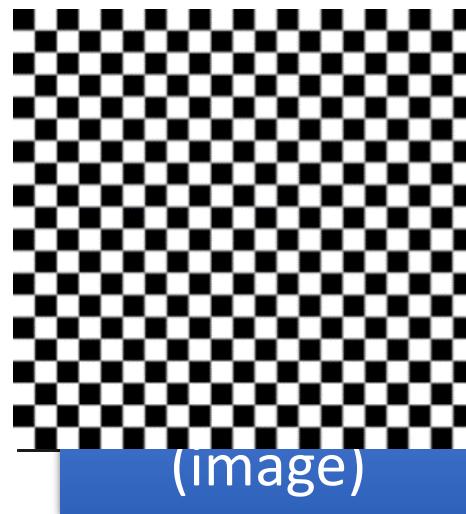
<https://dreamscopeapp.com/>

Deep Style

A Neural Algorithm
of Artistic Style



More Application: Chess



Next move
(19×19
positions)

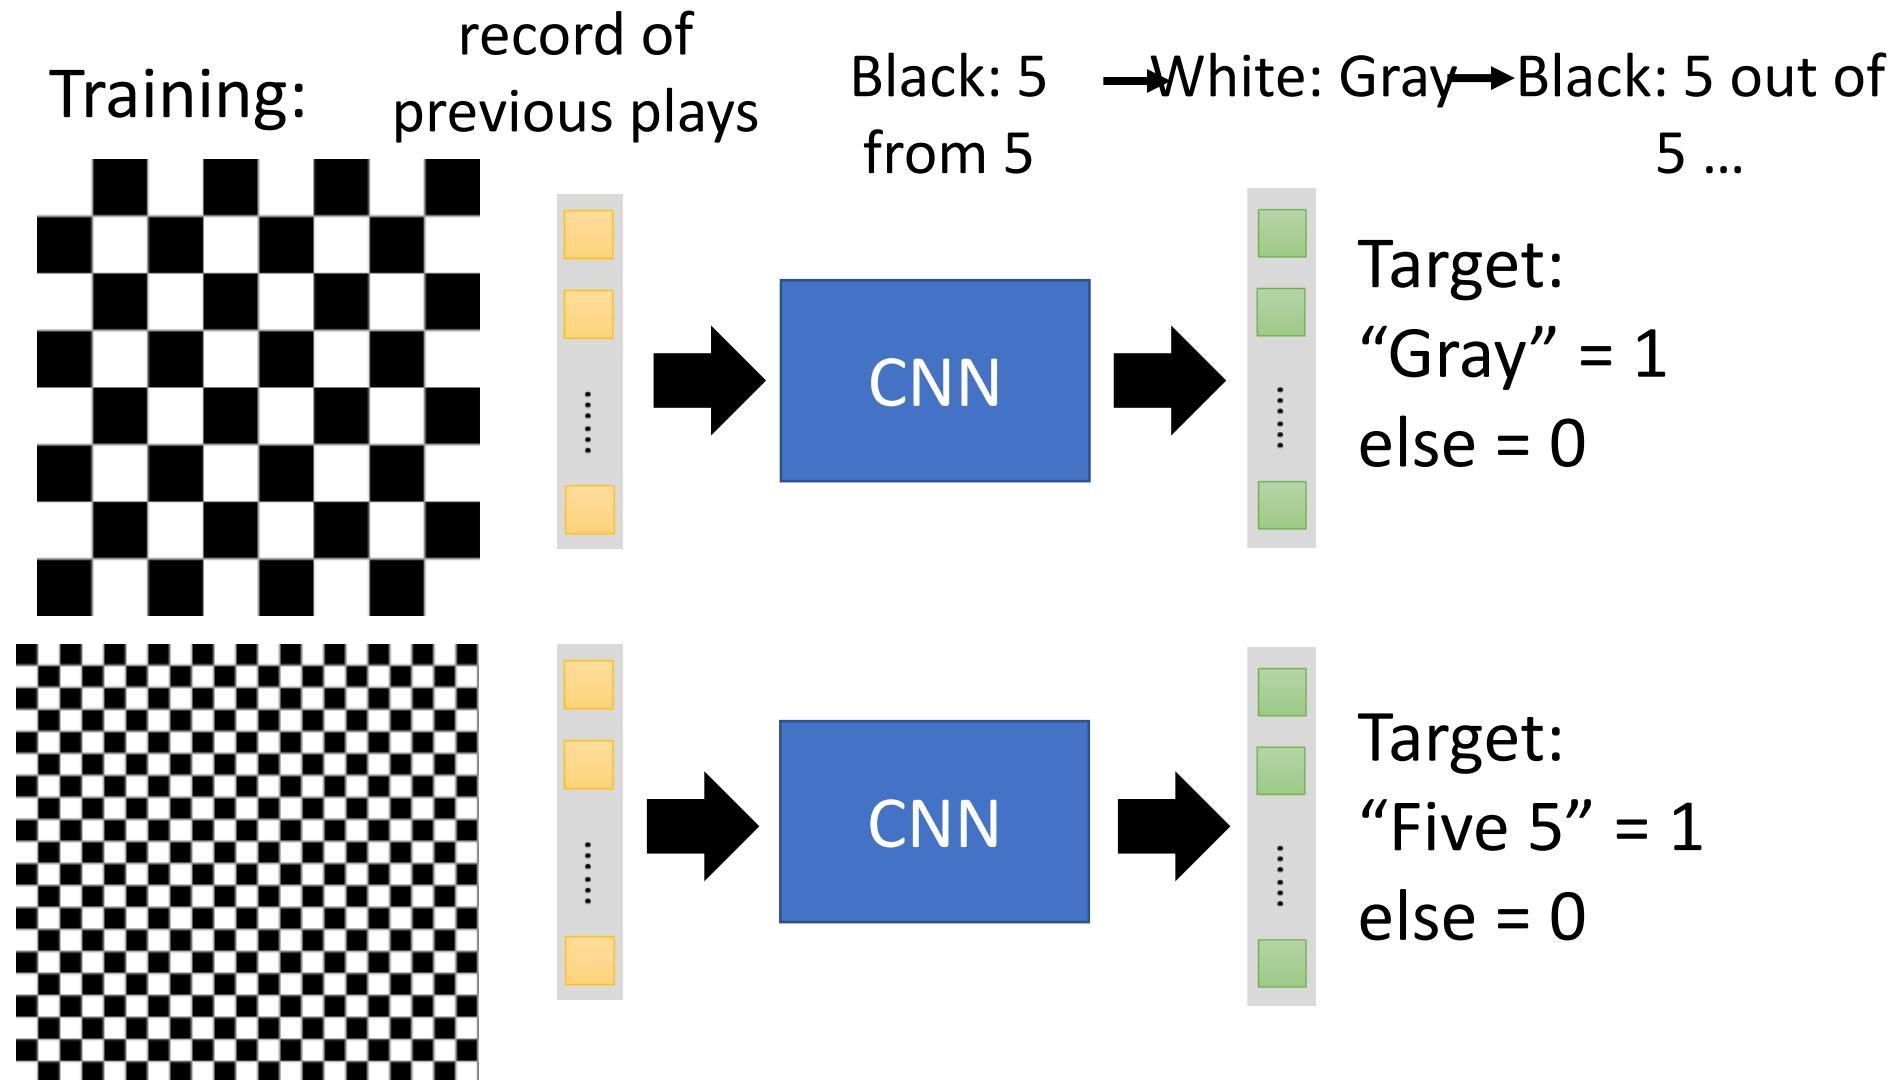
19×19 vector

Black: 1
white: -1
none: 0

Fully-connected feedforward
network can be used

But CNN performs much better.

More Application: Catur



Why CNN for playing Chess Go?

- Some patterns are much smaller than the whole image

Alpha Go uses 5×5 for first layer



- The same patterns appear in different regions.



Why CNN for playing Chess Go?

- Subsampling the pixels will not change the object

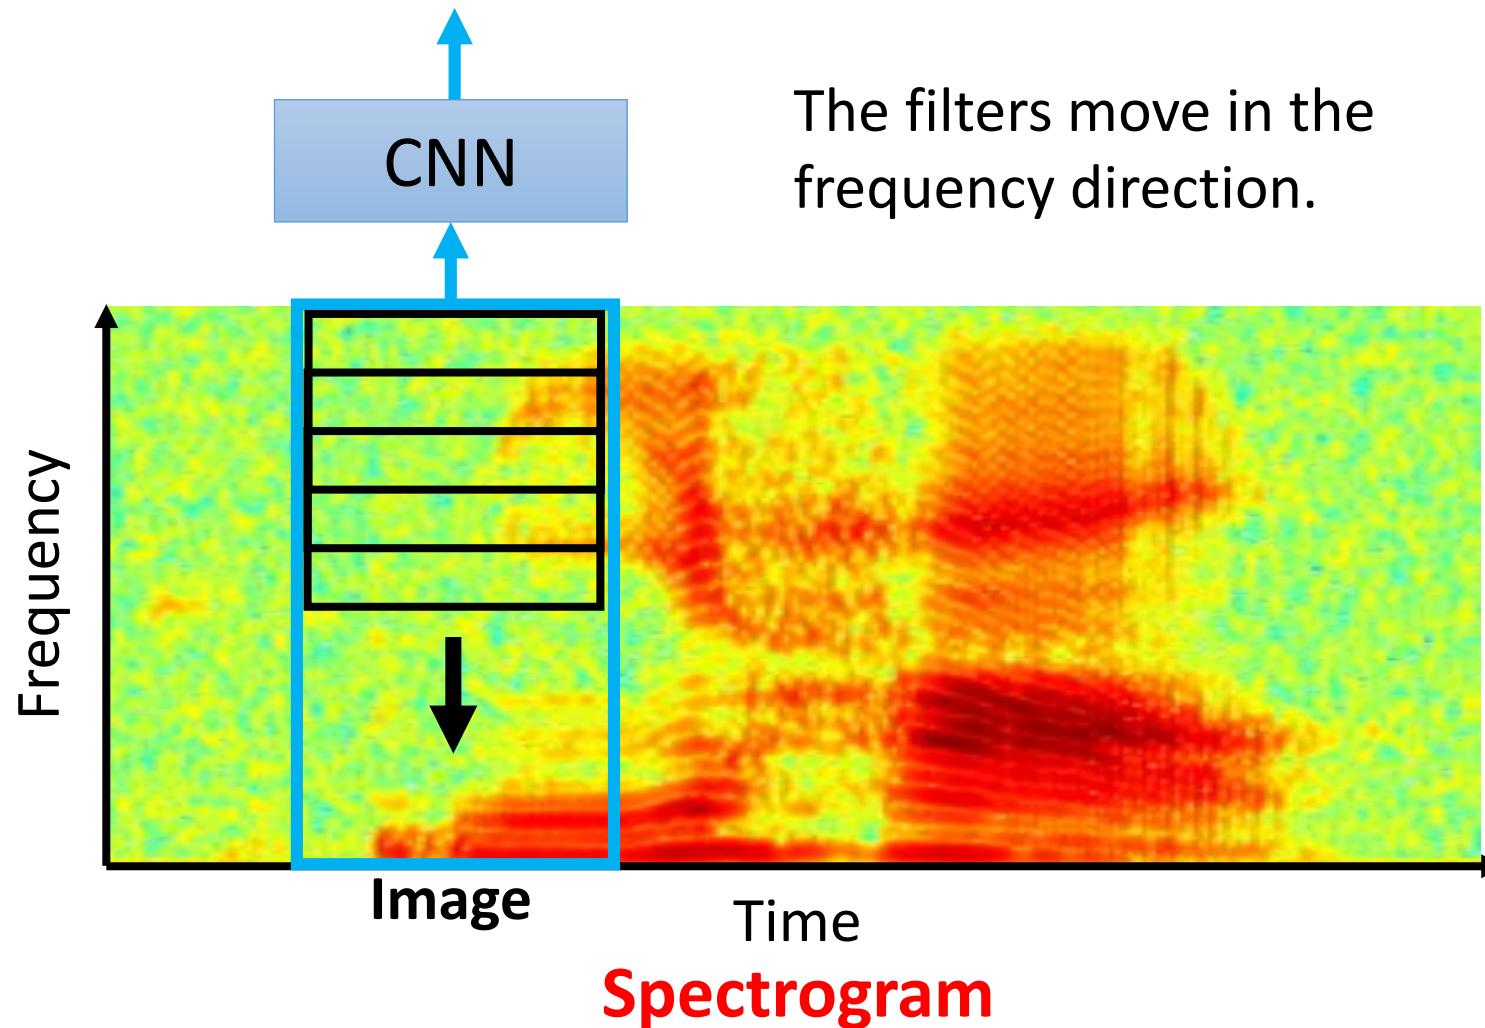


Max Pooling

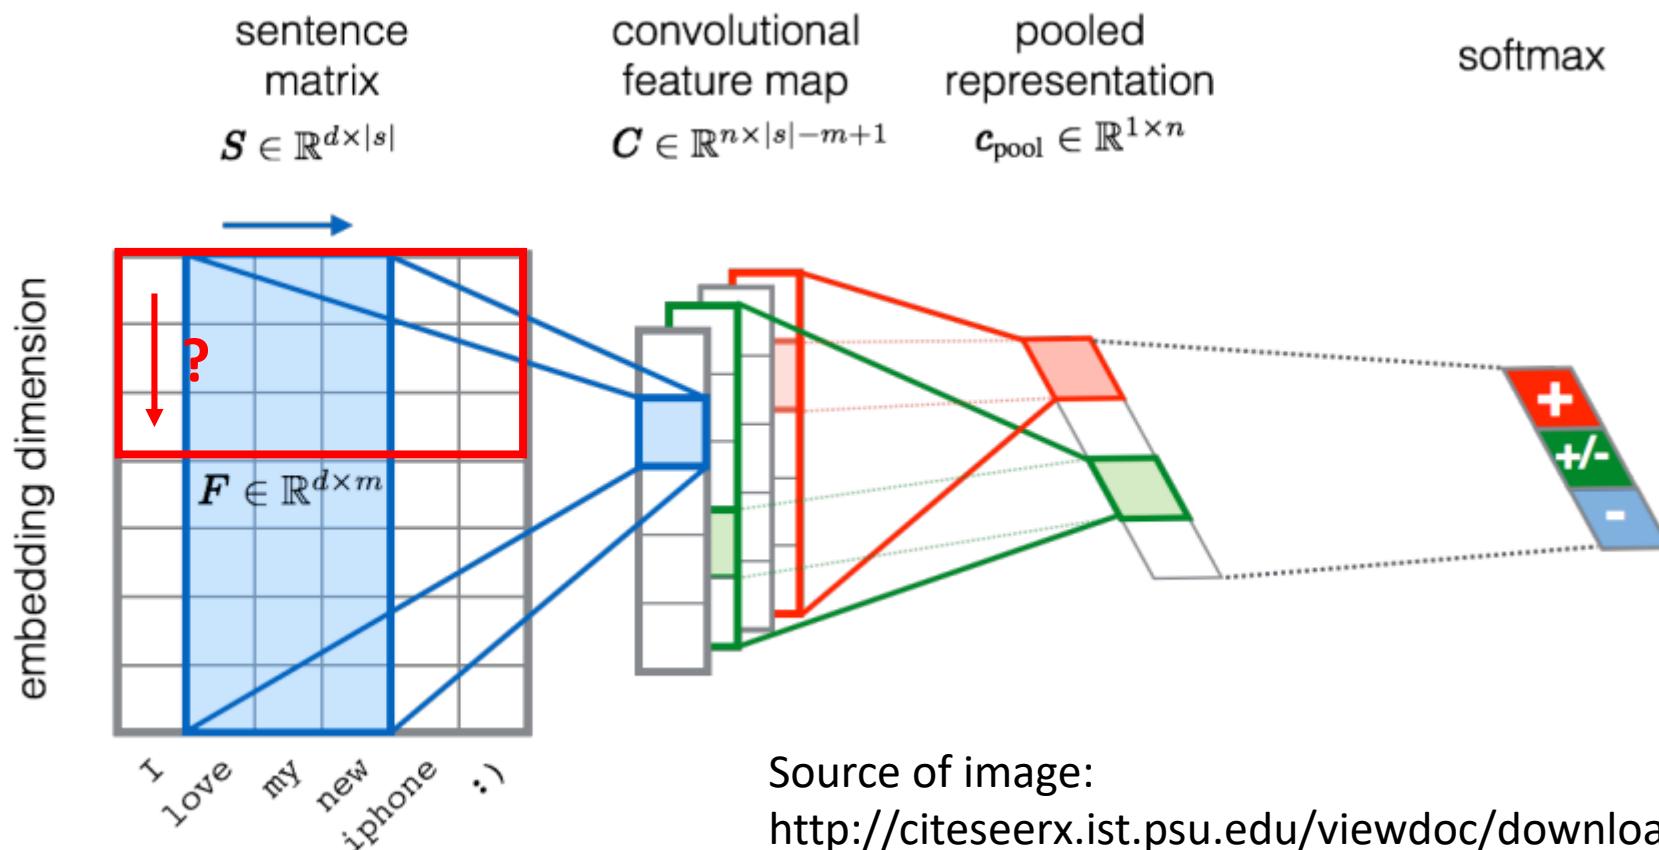
How to explain this???

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1 with a different bias for each position, and applies a softmax function. The Alpha Go does not use Max Pooling Extended Data Table 3 additionally show the results of training with $k = 128, 256$ and 384 filters.

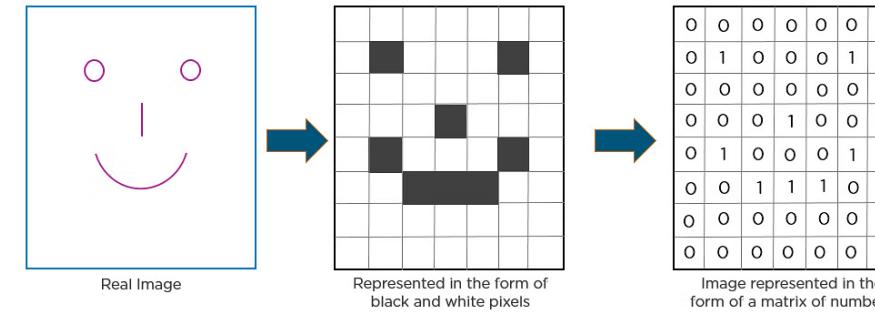
More Application: Speech



More Application: Text



CNN



- Understanding Convolutions

<http://colah.github.io/posts/2014-07-Understanding-Convolutions/>

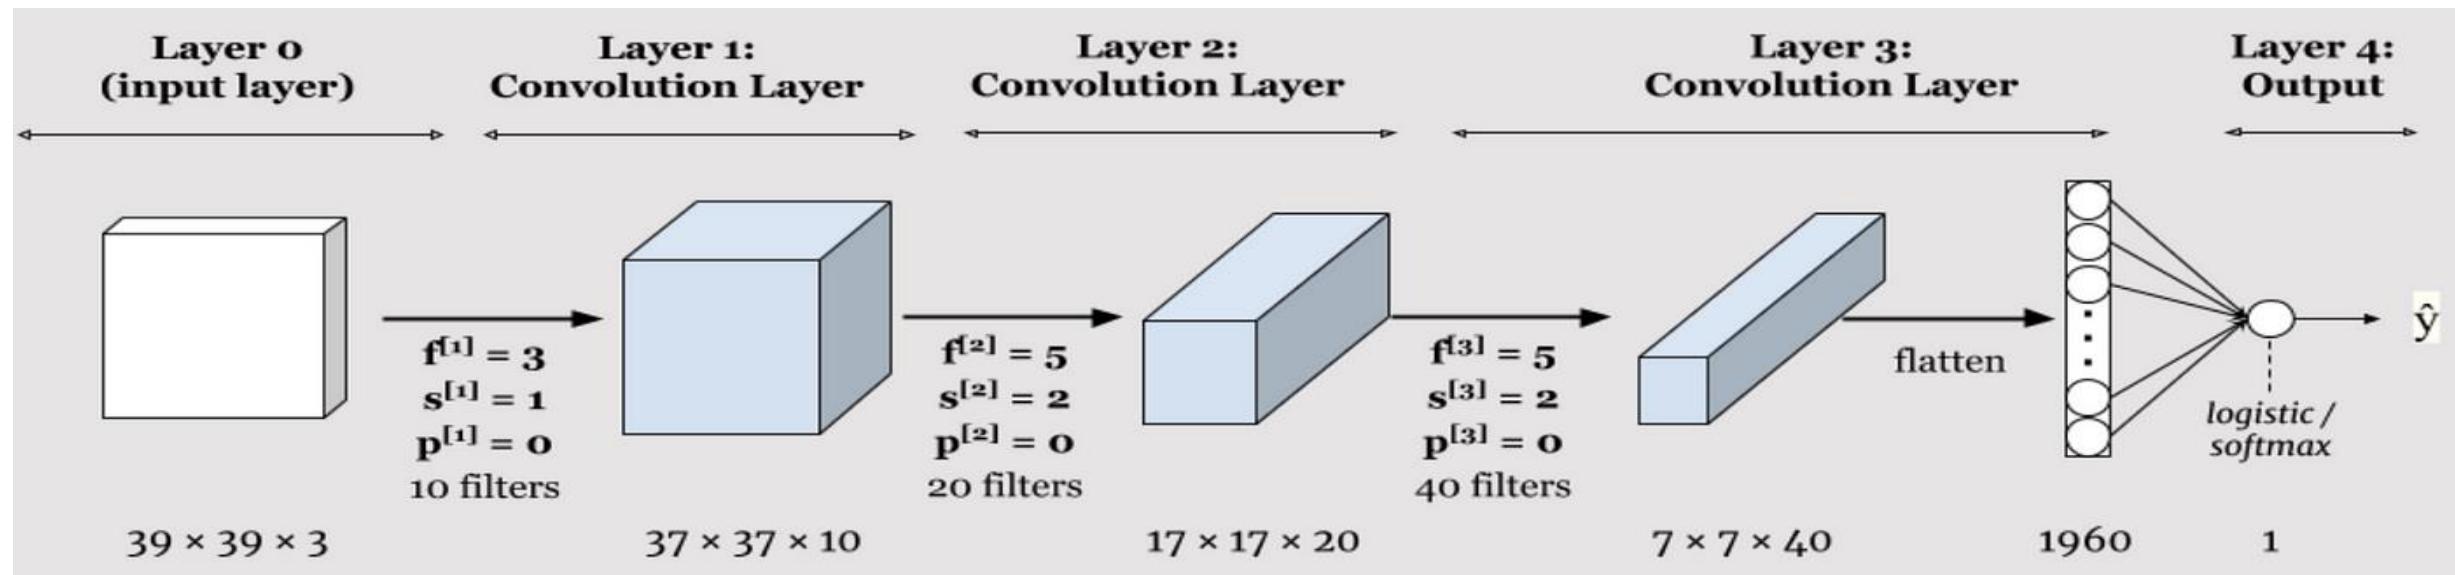
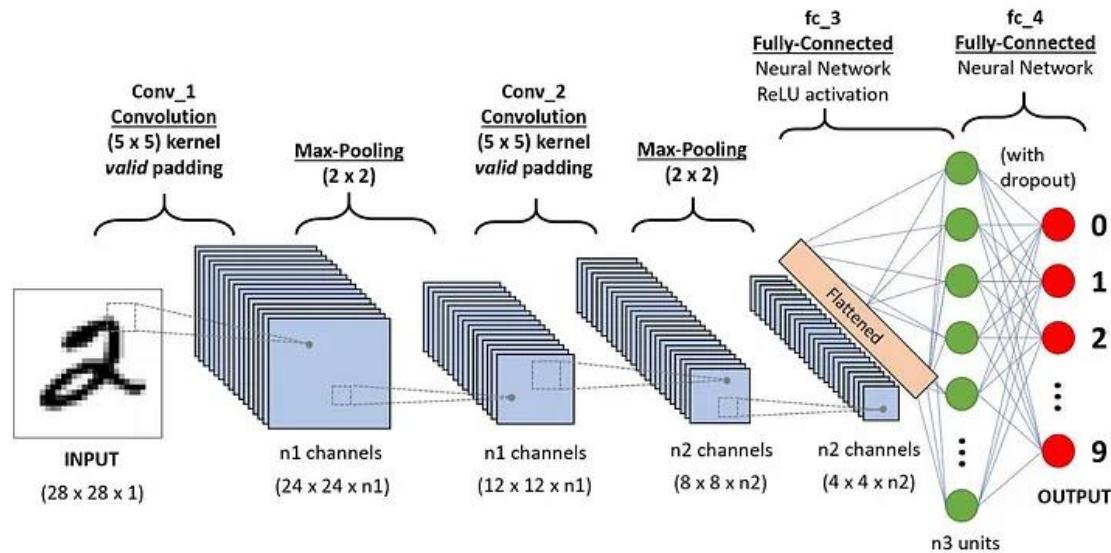
<https://python.plainenglish.io/cat-dog-classification-with-cnn-84af3ae98c44>

https://betterexplained.com/articles/intuitive-convolution/

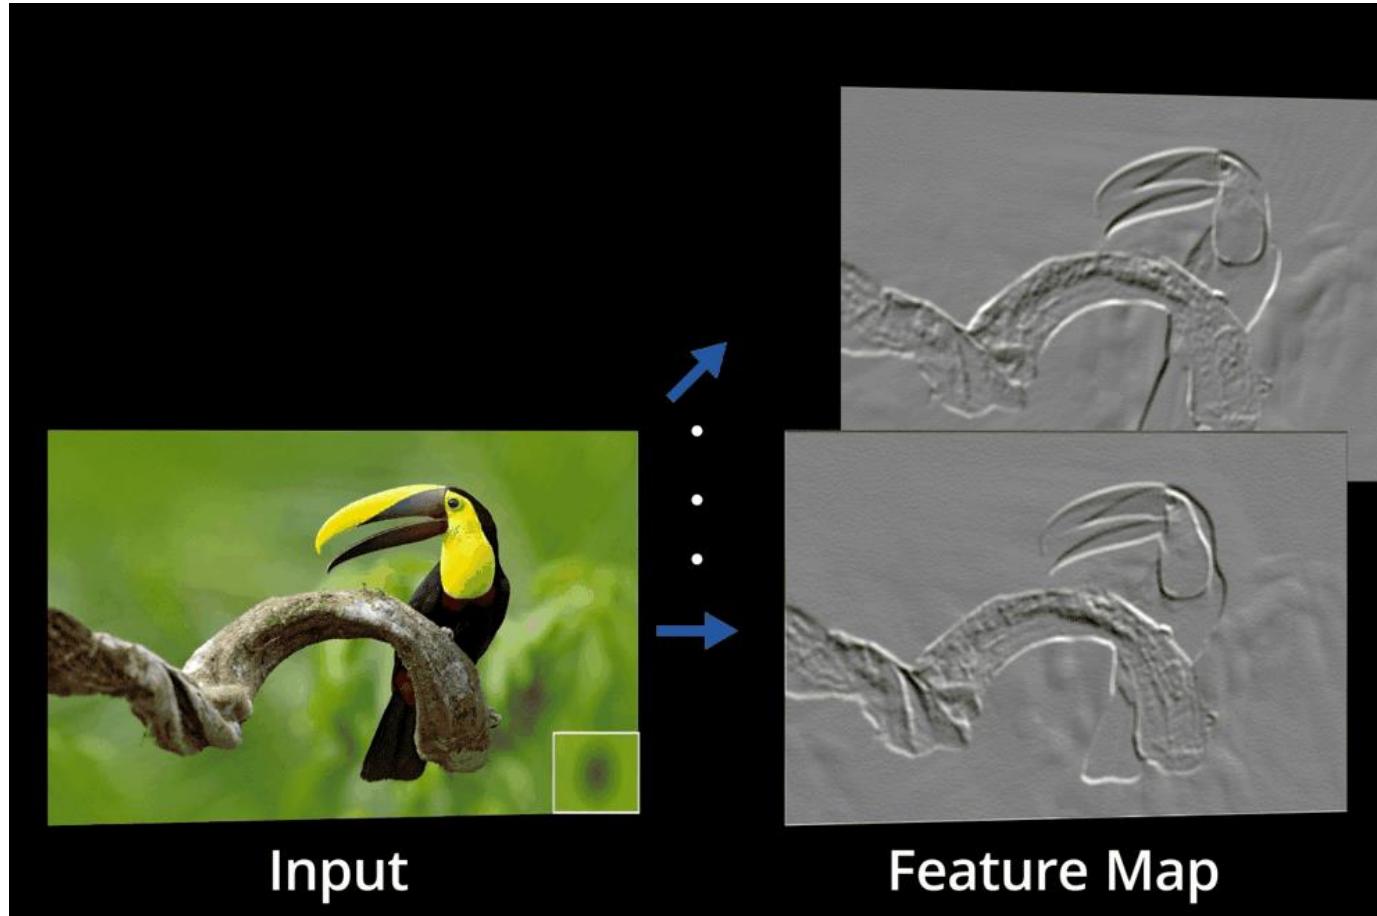
https://www.tomasbeuzen.com/deep-learning-with-pytorch/chapters/chapter5_cnns-pt1.html

<https://poloclub.github.io/cnn-explainer/>

<https://medium.com/thedeepphub/convolutional-neural-networks-a-comprehensive-guide-5cc0b5eae175>



Feature ?



Next.....

The First Step:

- Understanding “convolution” operations in CNN
- Understanding “Image RGB
- Understanding “Image Greyscale
- Mathematical understanding

I.....

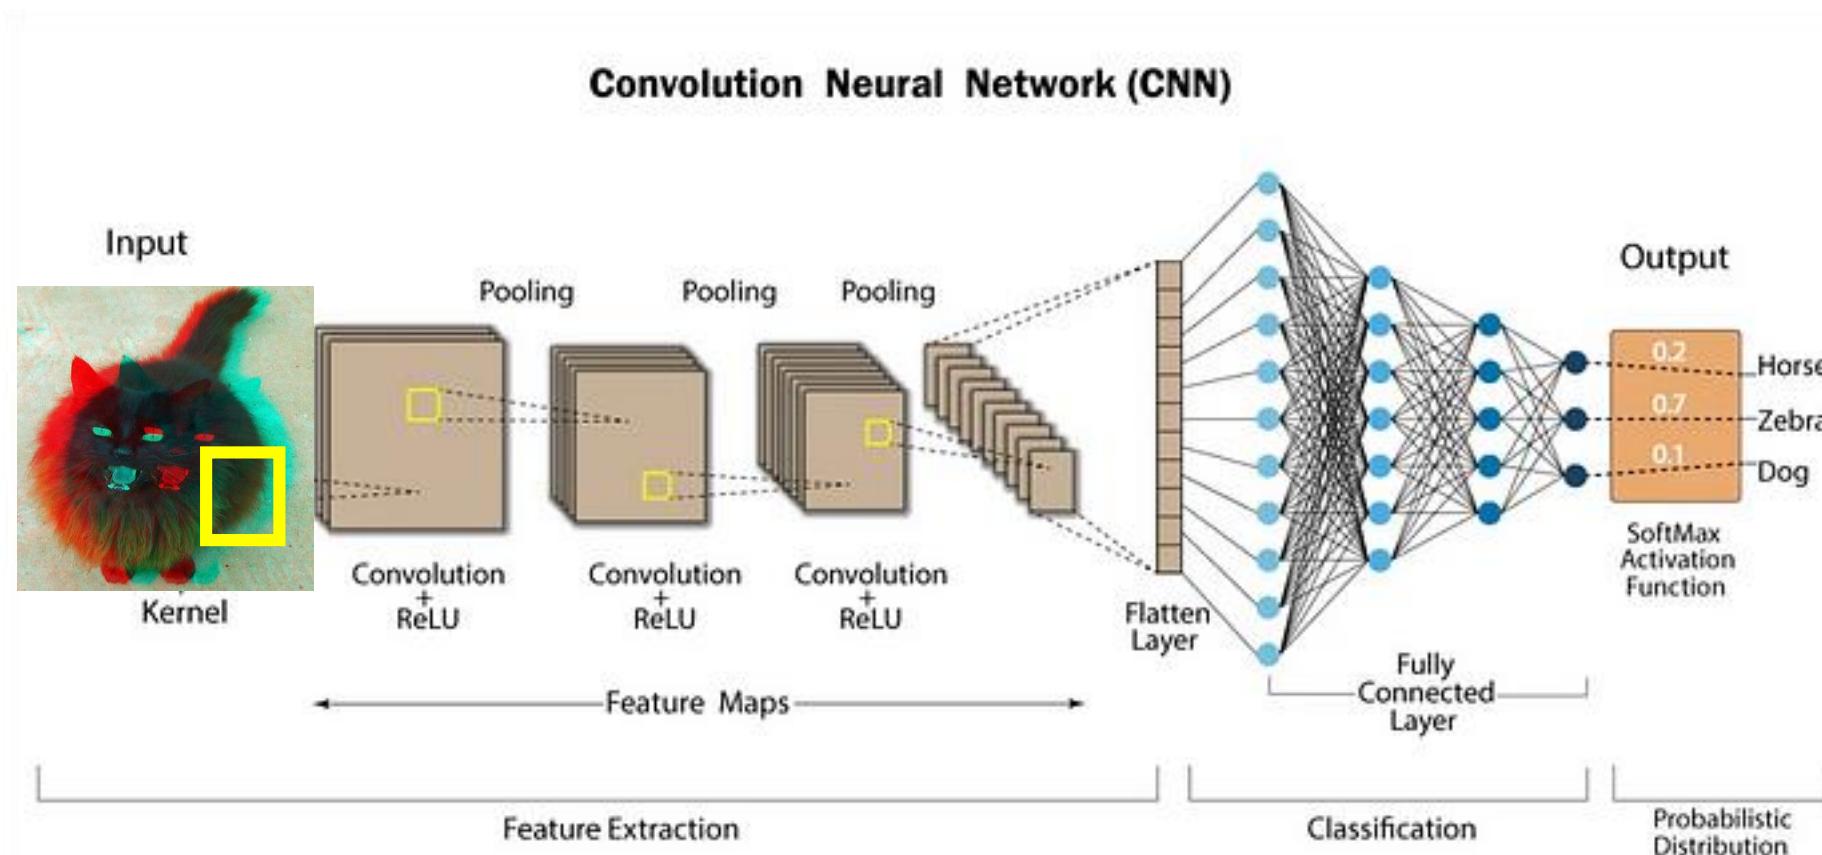
II.....

The Second Step:

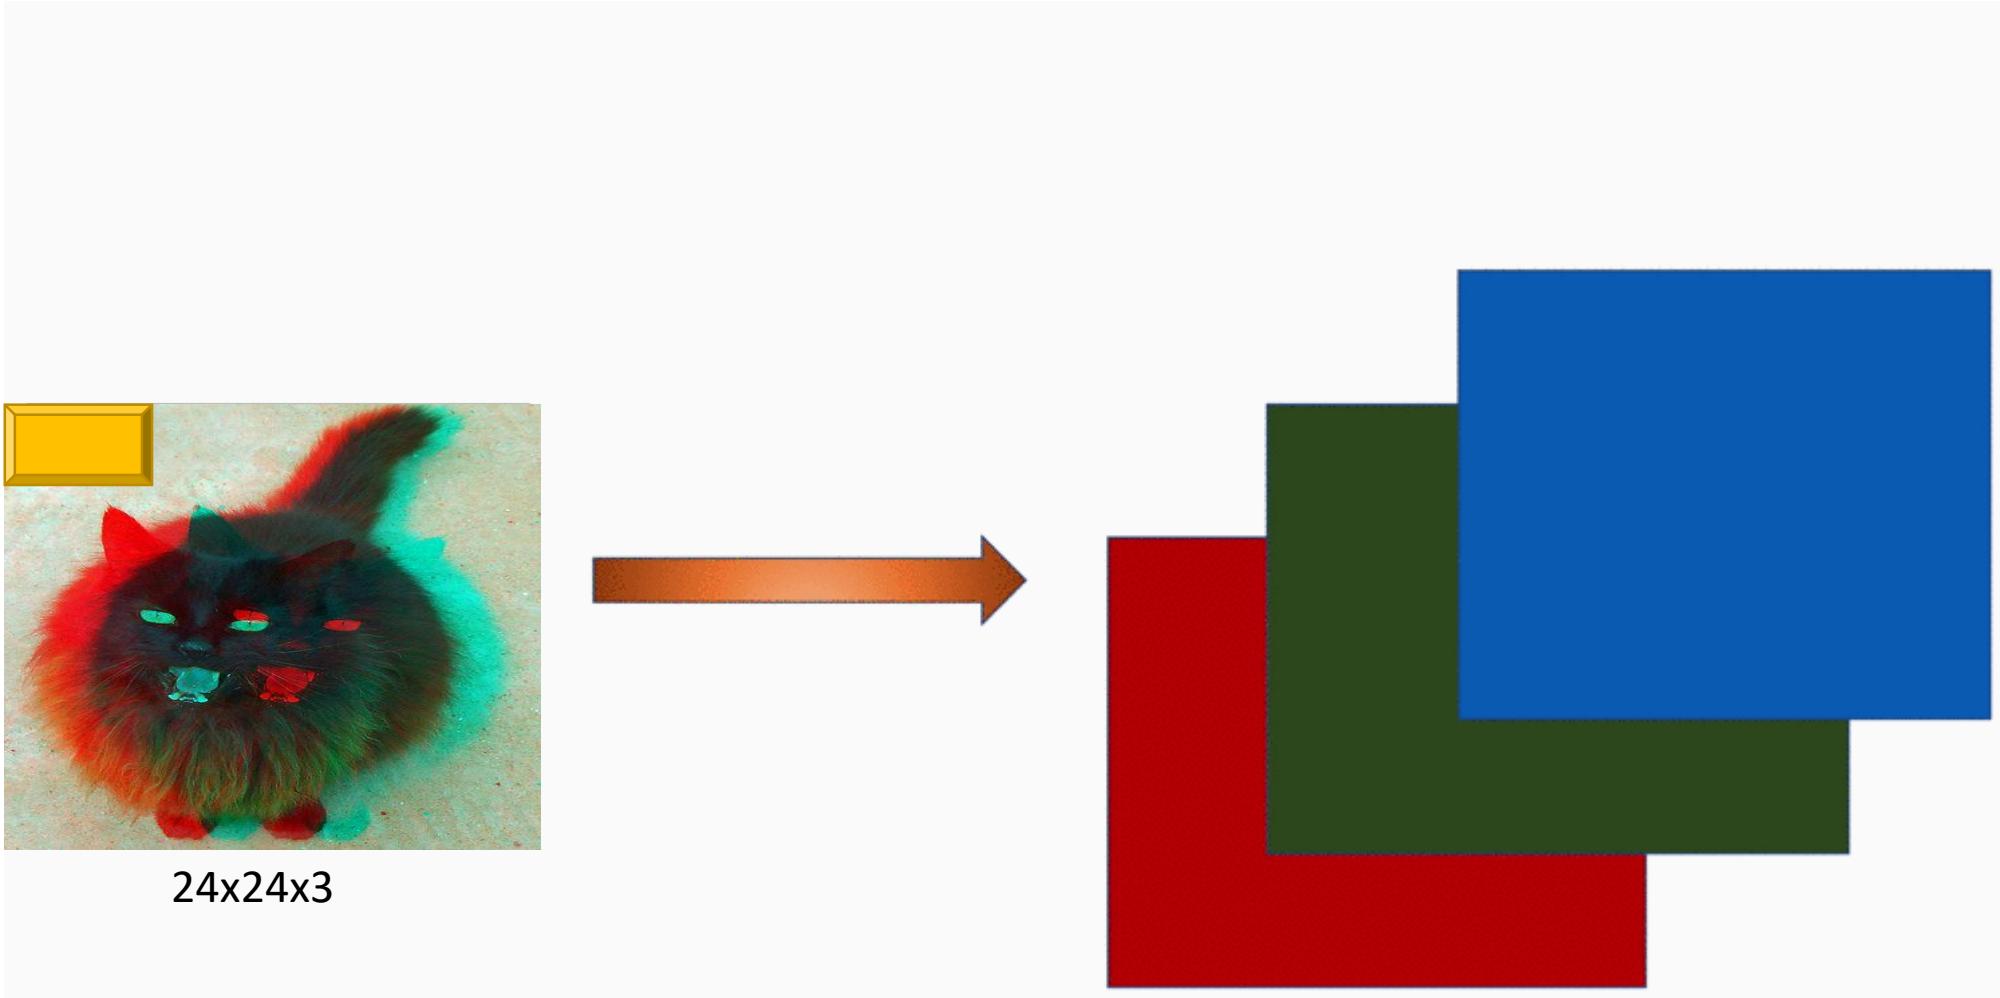
- What is Convolutional Neural Network?
- Convolutional Layer
- Padding and Stride
- Pooling
- ReLU
- Basic Python Implementation

The First Step:

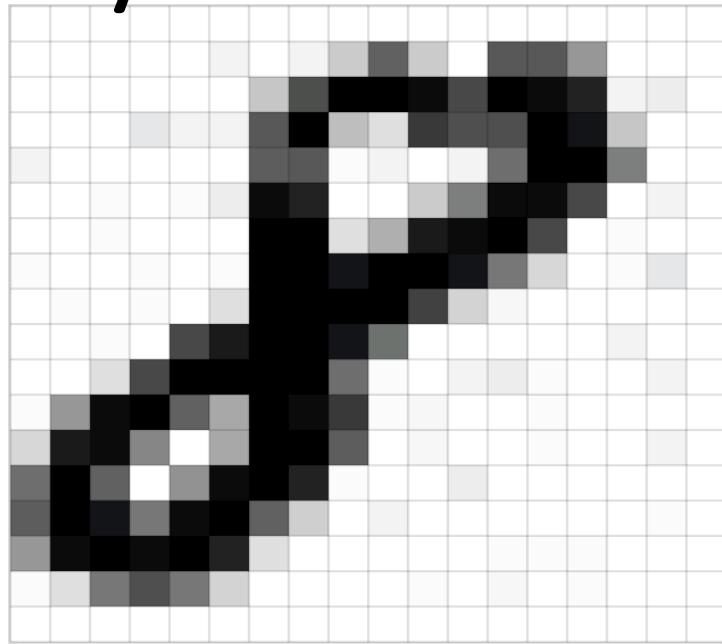
Understanding “convolution” operations in CNN



Understanding “Image RGB”



Understanding “Image Greyscale



24x24x1

$I_x \rightarrow 0 - 255$

Understanding “Image Binary

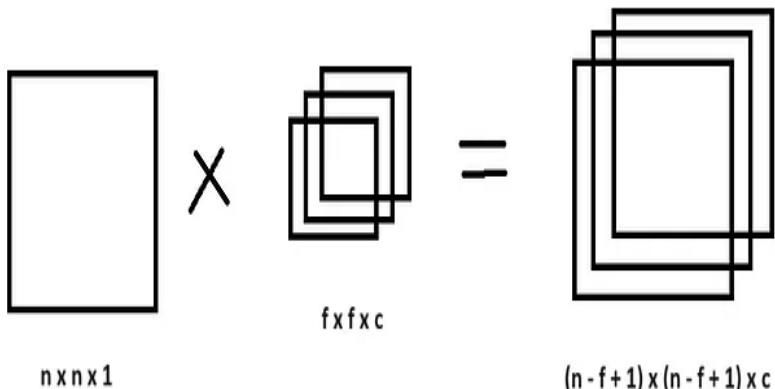


24x24x1

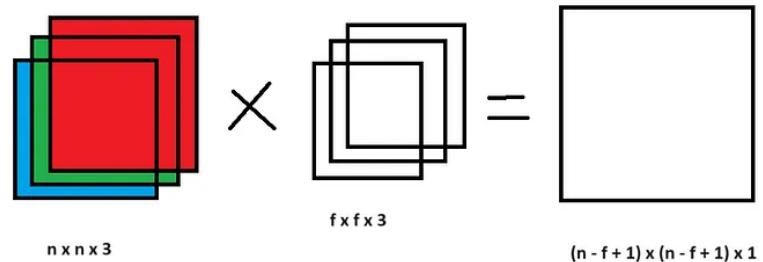
$I_x \rightarrow 0 - 1$

Mathematical understanding

Grayscale image, lets say the image size is $24 \times 24 \times 1$ (24 is the width, 24 is the height, 1 is the number of color channels, since its a grayscale image, it will have 1 channel; if it would have been a coloured image, number of channels would have been 3) and the kernel size is $3 \times 3 \times 64$ (3 is the width, 3 is the height, 64 is the number of such kernels); then after the convolution operation the size of the output image will be $22 \times 22 \times 64$ ($(24 - 3 + 1) \times (24 - 3 + 1) \times 64$)



The image is a coloured one, lets consider the size of the image as $24 \times 24 \times 3$, (for performing convolution operation on a coloured image, the kernel should also have 3 channels), kernel size as $4 \times 4 \times 3$ (where 3 is the number of channels). The number of channels of the image and the kernel should match incase of coloured image for the convolution operation to happen. The convolution operation will happen by super imposing the kernel on the image. The size of the image after the convolution operation will be $(21 \times 21 \times 1)$. Here 1 denotes the number of colour channels, which means after the convolution operation, grayscale image is obtained. This is a common scenario in convolutional neural networks, where intermediate layers may reduce the number of channels while preserving spatial information. **If there would have been multiple kernels lets say 64, then the size of the image would have been $21 \times 21 \times 1 \times 64$.**



Feature Map = Input Image x Feature Detector

- Size of Image: $N \times N$
- Size of Filter : $F \times F$
- $(N \times N)(F \times F) = (N-F+1) \times (N-F+1)$ = Convolution Operation
- Let us consider dimensions of feature map be $H \times W \times C$
- H — Height of feature map
- W — Width of feature map
- C — Number of channels in feature map
- F — Filter / Kernel
- S — Stride
- The output image will be $\{(H-F)/S+1\} \times \{(W-F)/S+1\} \times C$

Convolution process

$$I_{y,x} \star h = \sum_{i=-n_2}^{n_2} \sum_{j=-m_2}^{m_2} I(y-i, x-j) \cdot h(i, j)$$

n_2 is half the height of the filter

m_2 is half the length of the filter

x is the column position of a certain pixel in the image

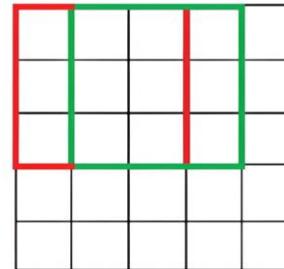
y is the row position of a certain pixel in the image.

$n_2 = [n/2]$, $m_2 = [m/2]$ above, n is the height of the filter, and m is the length of the filter h .

$$I = \begin{bmatrix} I_{00} & I_{01} & I_{02} & I_{03} \\ I_{10} & I_{11} & I_{12} & I_{13} \\ I_{20} & I_{21} & I_{22} & I_{23} \\ I_{30} & I_{31} & I_{32} & I_{33} \end{bmatrix} = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \\ 3 & 6 & 9 & 12 \end{bmatrix} h = \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

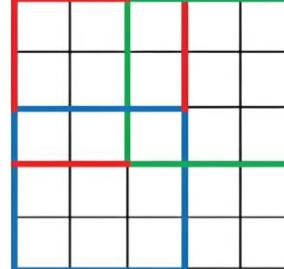
Stride

Convolution with Stride=1



Output

Convolution with Stride=2



Output

Padding

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

0	-1	0	114			
-1	5	-1				
0	-1	0				

i.e. **padding=1**, and **stride=1**,

$$n_{out} = \left\lfloor \frac{n_{in} + 2padding - n}{stride} \right\rfloor + 1$$

$$m_{out} = \left\lfloor \frac{m_{in} + 2padding - m}{stride} \right\rfloor + 1$$

Where n_{in} , m_{in} – height and length of the input image,
 n_{out} , m_{out} – are the height and length of the output image.

$$I = \begin{bmatrix} I_{00} & I_{01} & I_{02} & I_{03} \\ I_{10} & I_{11} & I_{12} & I_{13} \\ I_{20} & I_{21} & I_{22} & I_{23} \\ I_{30} & I_{31} & I_{32} & I_{33} \end{bmatrix} \quad I = \begin{bmatrix} I_{00} & I_{01} & I_{02} & I_{03} \\ I_{10} & I_{11} & I_{12} & I_{13} \\ I_{20} & I_{21} & I_{22} & I_{23} \\ I_{30} & I_{31} & I_{32} & I_{33} \end{bmatrix}$$

$$I = \begin{bmatrix} I_{00} & I_{01} & I_{02} & I_{03} \\ I_{10} & I_{11} & I_{12} & I_{13} \\ I_{20} & I_{21} & I_{22} & I_{23} \\ I_{30} & I_{31} & I_{32} & I_{33} \end{bmatrix} \quad I = \begin{bmatrix} I_{00} & I_{01} & I_{02} & I_{03} \\ I_{10} & I_{11} & I_{12} & I_{13} \\ I_{20} & I_{21} & I_{22} & I_{23} \\ I_{30} & I_{31} & I_{32} & I_{33} \end{bmatrix}$$

$$I = \begin{bmatrix} I_{00} & I_{01} & I_{02} & I_{03} \\ I_{10} & I_{11} & I_{12} & I_{13} \\ I_{20} & I_{21} & I_{22} & I_{23} \\ I_{30} & I_{31} & I_{32} & I_{33} \end{bmatrix} = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \\ 3 & 6 & 9 & 12 \end{bmatrix}$$

$$\begin{aligned} I_{y=1,x=1} \star h &= \sum_{i=-n_2}^{n_2} \sum_{j=-m_2}^{m_2} I(x-i, y-j) \cdot h(i, j) \\ &= \sum_{i=-1}^1 \sum_{j=-1}^1 I(1-i, 1-j) \cdot h(i, j) \\ &= I_{00} \cdot h_{00} + I_{01} \cdot h_{01} + I_{02} \cdot h_{02} + \\ &\quad I_{10} \cdot h_{10} + I_{11} \cdot h_{11} + I_{12} \cdot h_{12} + \\ &\quad I_{20} \cdot h_{20} + I_{21} \cdot h_{21} + I_{22} \cdot h_{22} \\ &= 0 \cdot 0 + 3 \cdot 1 + 6 \cdot 0 + \\ &\quad 1 \cdot 0 + 4 \cdot 1 + 7 \cdot 0 + \\ &\quad 2 \cdot 0 + 5 \cdot 1 + 8 \cdot 0 \\ &= 3 + 4 + 5 = 12 \end{aligned}$$

$$\begin{aligned} I_{y=1,x=2} \star h &= \sum_{i=-1}^1 \sum_{j=-1}^1 I(2-i, 1-j) \cdot h(i, j) \\ &= 3 \cdot 0 + 6 \cdot 1 + 9 \cdot 0 + \\ &\quad 4 \cdot 0 + 7 \cdot 1 + 10 \cdot 0 + \\ &\quad 5 \cdot 0 + 8 \cdot 1 + 11 \cdot 0 \\ &= 6 + 7 + 8 = 21 \end{aligned}$$

$$I \star h = \begin{bmatrix} I_{y=1,x=1} & I_{y=1,x=2} \\ I_{y=2,x=1} & I_{y=2,x=2} \end{bmatrix} = \begin{bmatrix} 12 & 21 \\ 15 & 24 \end{bmatrix}$$

```

sum them

for i in range(k):
    for j in range(k):
        a = img[y+i+k_start][x+j+k_start] * filter_[i][j]
        print(f"i[{y+i+k_start},{x+j+k_start}]*f[{i},{j}]={a}")
    end='\\t')
    sum += a
    print("")
# Return the result of convolution for a given point
return sum

# Kernel size
k = 3

# Filter for convolution
filter_ =
[ [0, 1, 0],
  [0, 1, 0],
  [0, 1, 0],
]
# Input image
img =
[ [0, 3, 6, 9 ],
  [1, 4, 7, 10],
  [2, 5, 8, 11],
  [3, 6, 9, 12]
]

i[0,0]*f[0,0]=0          i[0,1]*f[0,1]=3          i[0,2]*f[0,2]=0
i[1,0]*f[1,0]=0          i[1,1]*f[1,1]=4          i[1,2]*f[1,2]=0
i[2,0]*f[2,0]=0          i[2,1]*f[2,1]=5          i[2,2]*f[2,2]=0
process (x=1, y=1) image point

i[0,1]*f[0,0]=0          i[0,2]*f[0,1]=6          i[0,3]*f[0,2]=0
i[1,1]*f[1,0]=0          i[1,2]*f[1,1]=7          i[1,3]*f[1,2]=0
i[2,1]*f[2,0]=0          i[2,2]*f[2,1]=8          i[2,3]*f[2,2]=0
process (x=2, y=1) image point

i[1,0]*f[0,0]=0          i[1,1]*f[0,1]=4          i[1,2]*f[0,2]=0
i[2,0]*f[1,0]=0          i[2,1]*f[1,1]=5          i[2,2]*f[1,2]=0
i[3,0]*f[2,0]=0          i[3,1]*f[2,1]=6          i[3,2]*f[2,2]=0
process (x=1, y=2) image point

i[1,1]*f[0,0]=0          i[1,2]*f[0,1]=7          i[1,3]*f[0,2]=0
i[2,1]*f[1,0]=0          i[2,2]*f[1,1]=8          i[2,3]*f[1,2]=0
i[3,1]*f[2,0]=0          i[3,2]*f[2,1]=9          i[3,3]*f[2,2]=0
process (x=2, y=2) image point

array([[12., 21.],
       [15., 24.]])

```

```

# Determine the size of the kernel
filter2 = np.zeros((k, k), dtype=int)
filter2[k//2, :] = 1
filter2[:, k//2] = 1
filter2[k//2, k//2] = 1

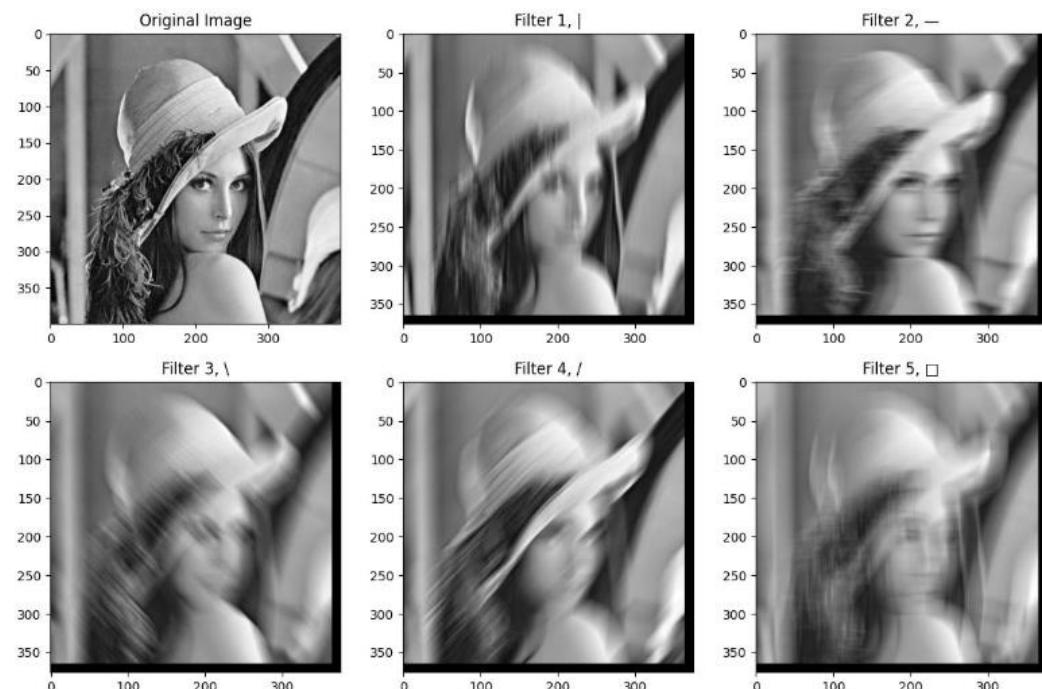
Convolution for Black and White Images
# Define the boundaries within which convolution will be
# Filter 3
performed
      """
row_start = int(np.floor(k/2))
[[1, 0, 0],
row_end = height-k+2
[0, 1, 0],
col_start = int(np.floor(k/2))
[0, 0, 1]],
col_end = width-k+2

# Create a new image filled with zeros, smaller than the
filters3 = np.zeros((k, k), dtype=int)
original by the size of the kernel
np.fill_diagonal(filter3, 1)
new_image = np.zeros((height-k+1, width-k+1))
for y in range(row_start, row_end):
    # for x in range(col_start, col_end):
        # Filter 4
        # for x in range(col_start, col_end):
            # Convolve for each point in the area defined by the
            [[0, 0, 1],
borders
            [0, 1, 0],
            new_image[y-row_start][x-col_start] = convolve(img,
filter_, k, x, y)
            """
return new_image
filter4 = np.zeros((k, k), dtype=int)
np.fill_diagonal(np.fliplr(filter4), 1)

# Load the image
img = cv2.imread('Lenna_(test_image).png', 0)
# Filter 5
height, width = img.shape

# Create a figure with 2 graphs
[[1, 1, 1],
[[1, 0, 1],
fig, axs = plt.subplots(ncols=2, nrows=1, figsize=(14, 6))

```



$$Y = \max_{i \in I, j \in J} X[i, j]$$

$$X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix}$$

$$X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix} X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix} X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix}$$

$$X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix} X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix} X = \begin{bmatrix} 0 & 3 & 6 & 9 \\ 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \end{bmatrix}$$

$$Y = \begin{bmatrix} 4 & 7 & 10 \\ 5 & 8 & 11 \end{bmatrix}$$

The Second Step:

The Second Step:

- What is Convolutional Neural Network?
- Convolutional Layer
- Padding and Stride
- Pooling
- ReLU
- Basic Python Implementation

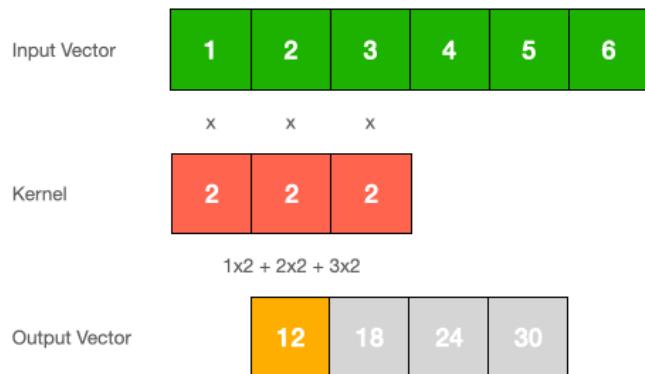
Kernel/filter

1	0	1
0	1	0
1	0	1

Filter 1
Matrix

0	1	0
0	1	0
0	1	0

Filter 2
Matrix



Kernel 3

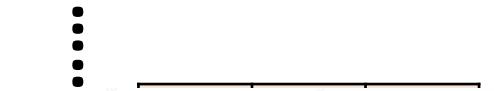
```
nn.Conv1d(in_channels=1, out_channels=1,
          kernel_size=3,
          bias=False)
```

Size

$$L_{\text{out}} = \left\lfloor \frac{L_{\text{in}} + 2 \times \text{padding} - \text{dilation} \times (\text{kernel_size} - 1) - 1}{\text{stride}} + 1 \right\rfloor$$

$$4 = \frac{6 + (2 * 0) - 1 * (3 - 1) - 1}{1} + 1$$

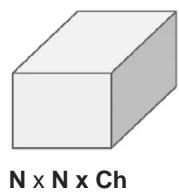
jinglescode.github.io



Original	Gaussian Blur	Sharpen	Edge Detection
$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$



Stride



Filter : $f \times f \times Ch$
Stride : S
Padding : P

$$= \left(\frac{n+2p-f}{S}+1\right) \times \left(\frac{n+2p-f}{S}+1\right) \times NCh$$

Input

a	b	c	d	e	f
g	h	i	j	k	l
m	n	o	p	q	r
s	t	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Kernel

U'	V'	W'
X'	Y'	Z'
L'	M'	N'

Sum

$$\begin{aligned}
 &= a.U' + b.V' + c.W' + \\
 &\quad g.X' + h.Y' + i.Z' + \\
 &\quad m.L' + n.M' + o.N' +
 \end{aligned}$$

⋮

⋮

$$h' = \left\lfloor \frac{h-f+p}{s} + 1 \right\rfloor$$

$$w' = \left\lfloor \frac{w-f+p}{s} + 1 \right\rfloor$$

$$(n \times n) * (f \times f) = \left(\frac{n+2p-f}{s} + 1 \right) \times \left(\frac{n+2p-f}{s} + 1 \right)$$

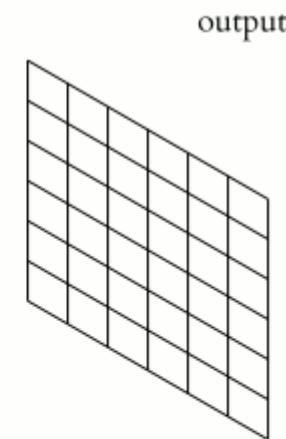
padding = p stride = s

resultant dimension

Stride

0	-1	0				
-1	5	-1				
0	-1	0				
7	6	5	5	6	7	
6	4	3	3	4	6	
5	3	2	2	3	5	
5	3	2	2	3	5	
6	4	3	2	3	5	
7	6	5	5	6	7	

input



output

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Those are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

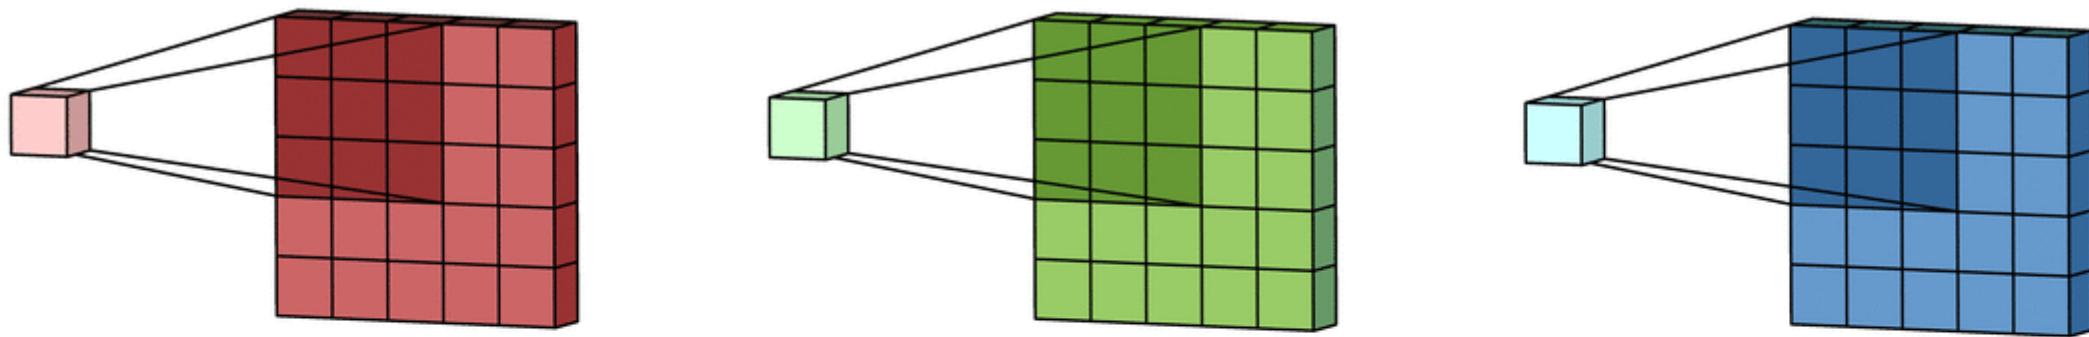
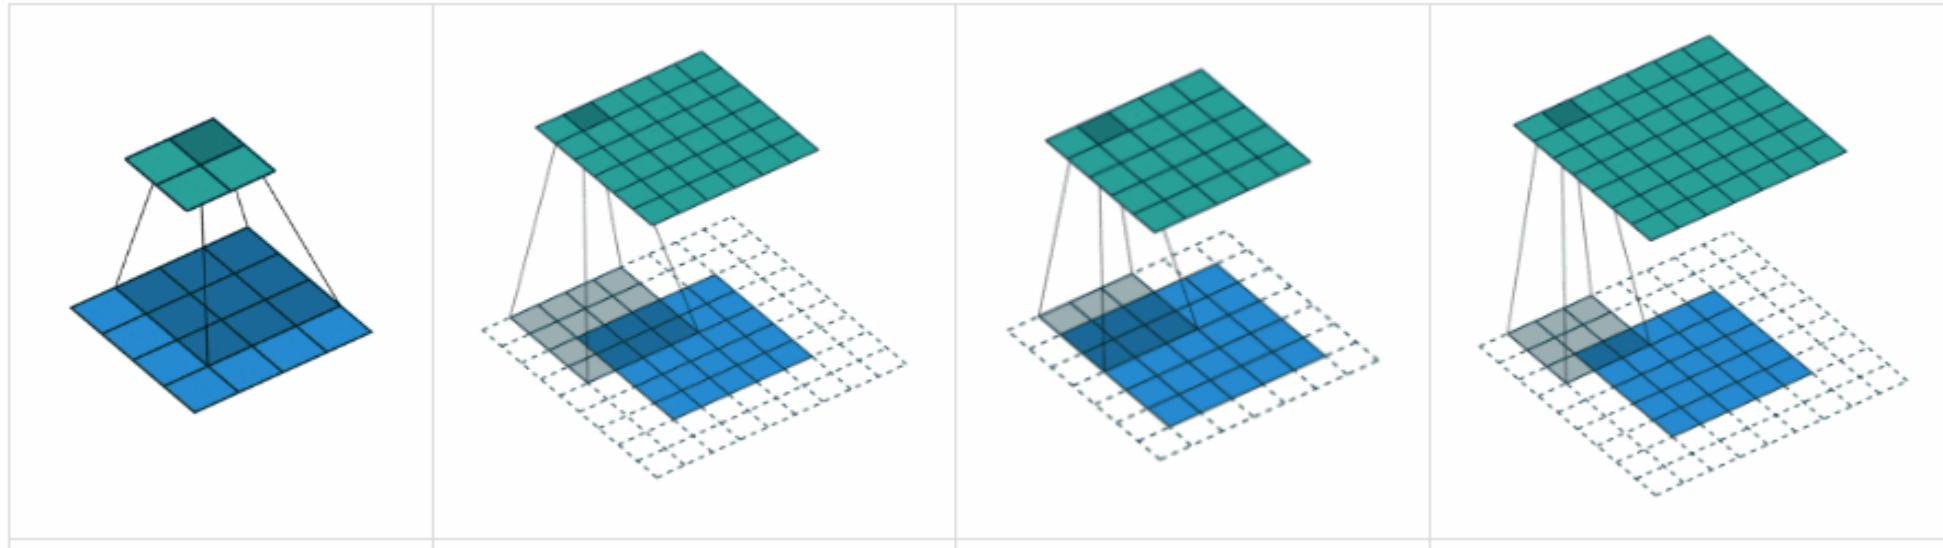
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

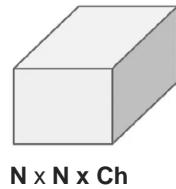


Padding , Stride

0	0	0	0	0	0	0	0
0	60	113	56	139	85	0	0
0	73	121	54	84	128	0	0
0	131	99	70	129	127	0	0
0	80	57	115	69	134	0	0
0	104	126	123	95	130	0	0
0	0	0	0	0	0	0	0

Kernel	0	-1	0
-1	5	-1	
0	-1	0	

114				



Filter : f x f Ch
Stride : S
Padding : P

$$= \left(\frac{n+2p-f}{S} + 1 \right) \times \left(\frac{n+2p-f}{S} + 1 \right) \times NCh$$

$$[(W-K+2P)/S]+1$$

$$output = \frac{W - N + 2P}{S} + 1$$

W = Panjang/Tinggi Input

N = Panjang/Tinggi Filter

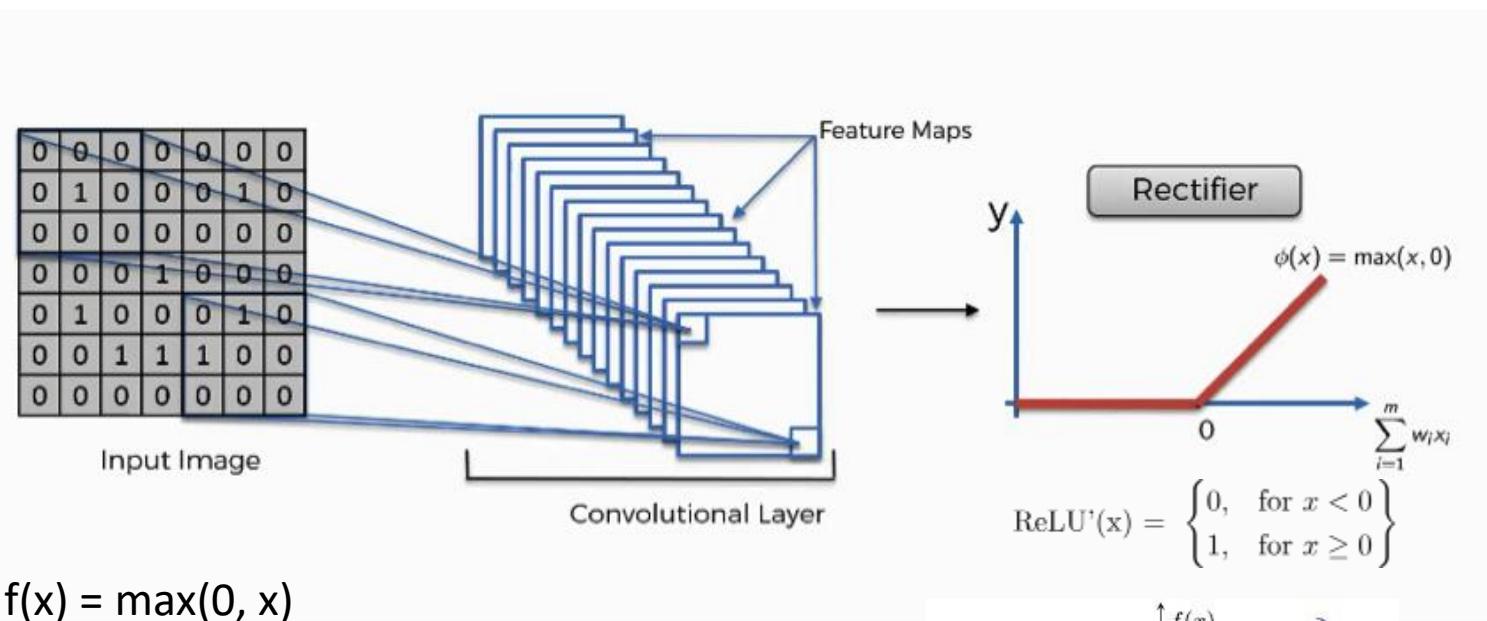
P = Zero Padding

S = Stride

ReLU

The rectified linear activation function

ReLU merupakan fungsi non-linear yang sederhana dan efektif dalam mempercepat proses pelatihan jaringan saraf tiruan, Goodfellow et al. (2016)

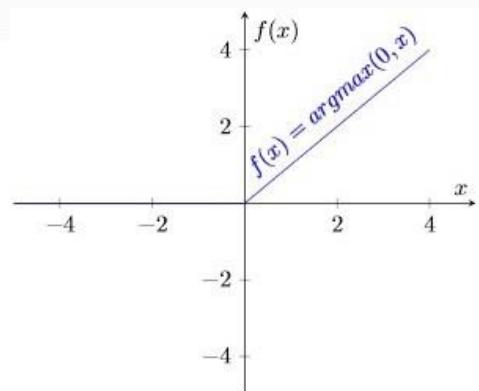


$$f(x) = \max(0, x)$$

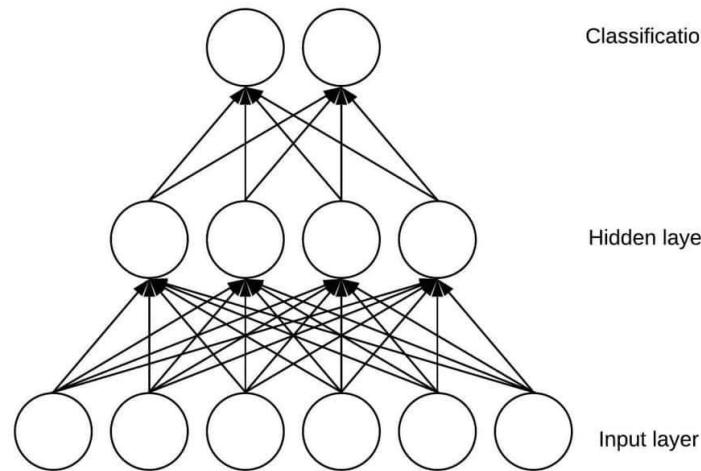
f(x): fungsi aktivasi ReLU

x: input ke neuron dalam jaringan saraf

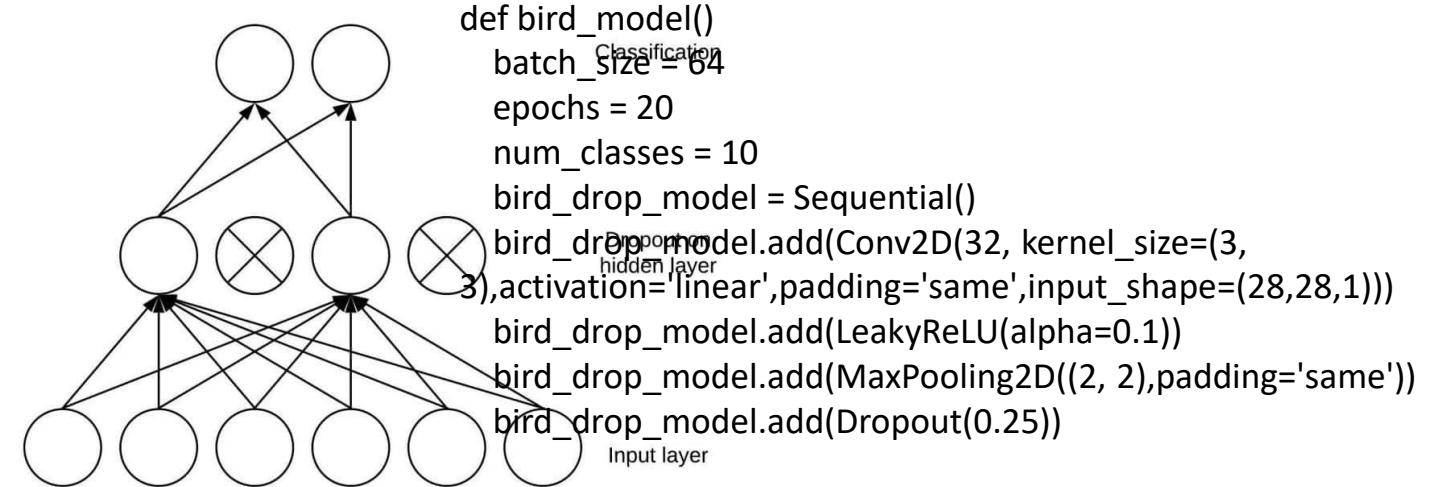
```
model = Sequential()  
model.add(Dense(64, input_dim=100))  
model.add(Activation('relu'))  
model.add(Dense(10))  
model.add(Activation('softmax'))
```



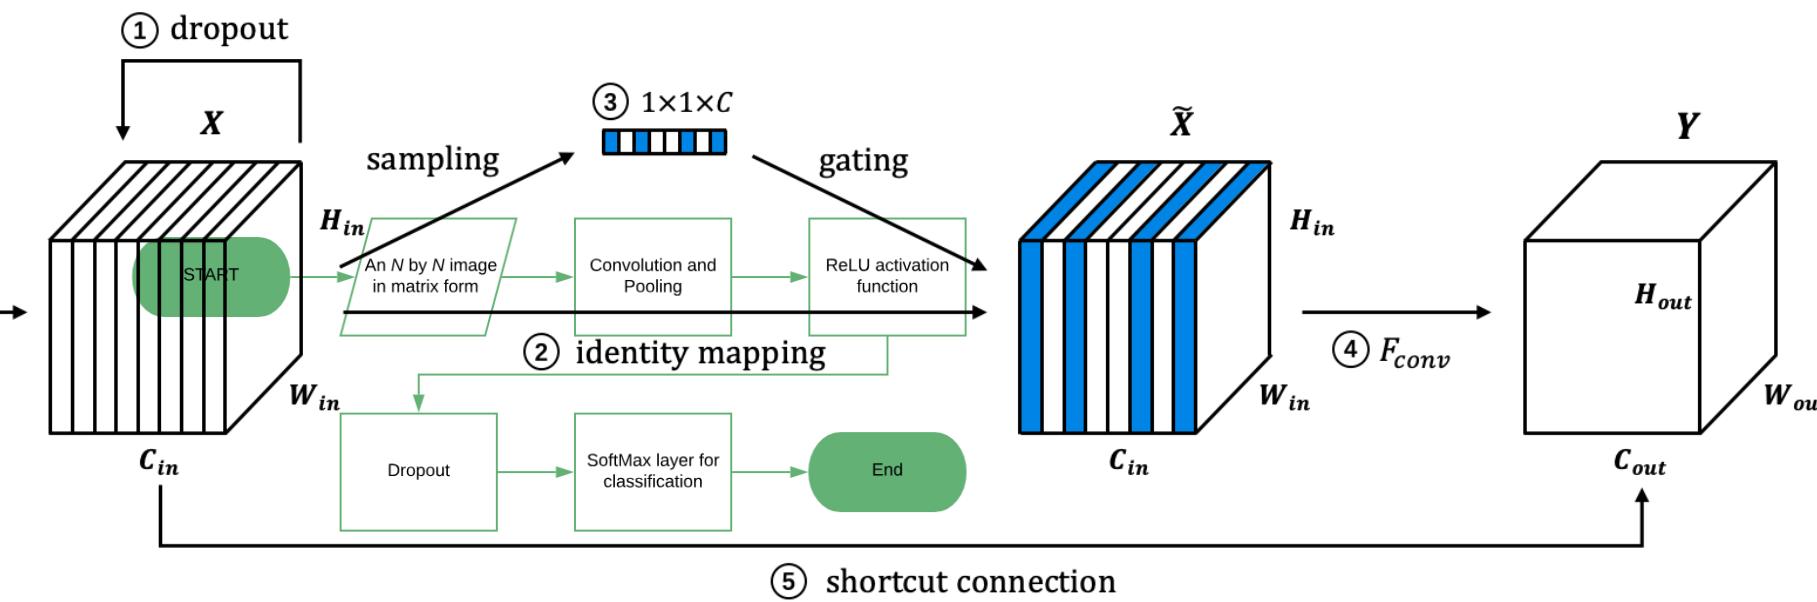
Dropout



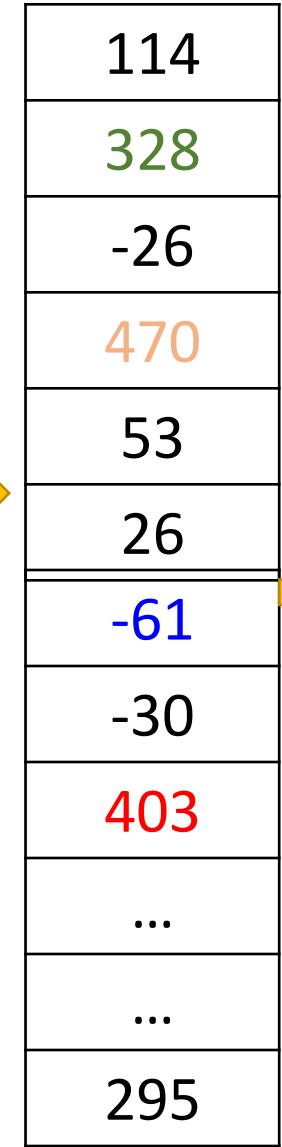
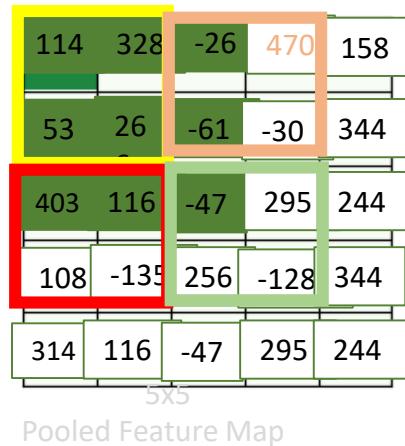
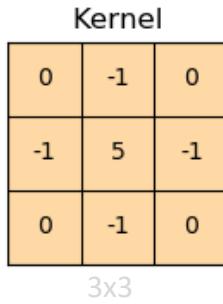
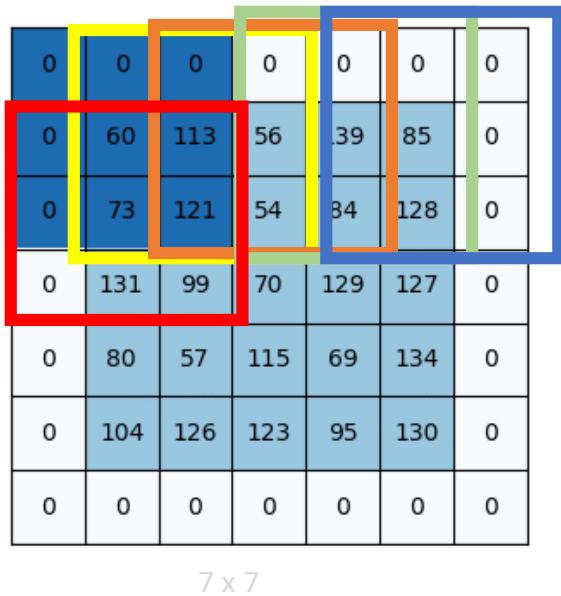
Without Dropout



With Dropout

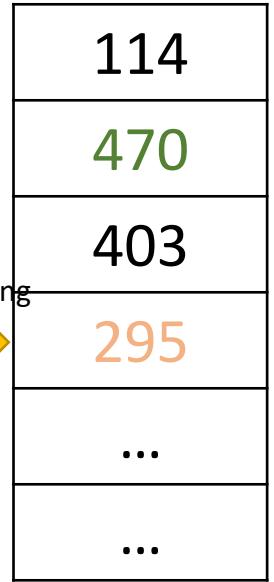


Pooling, Flatten

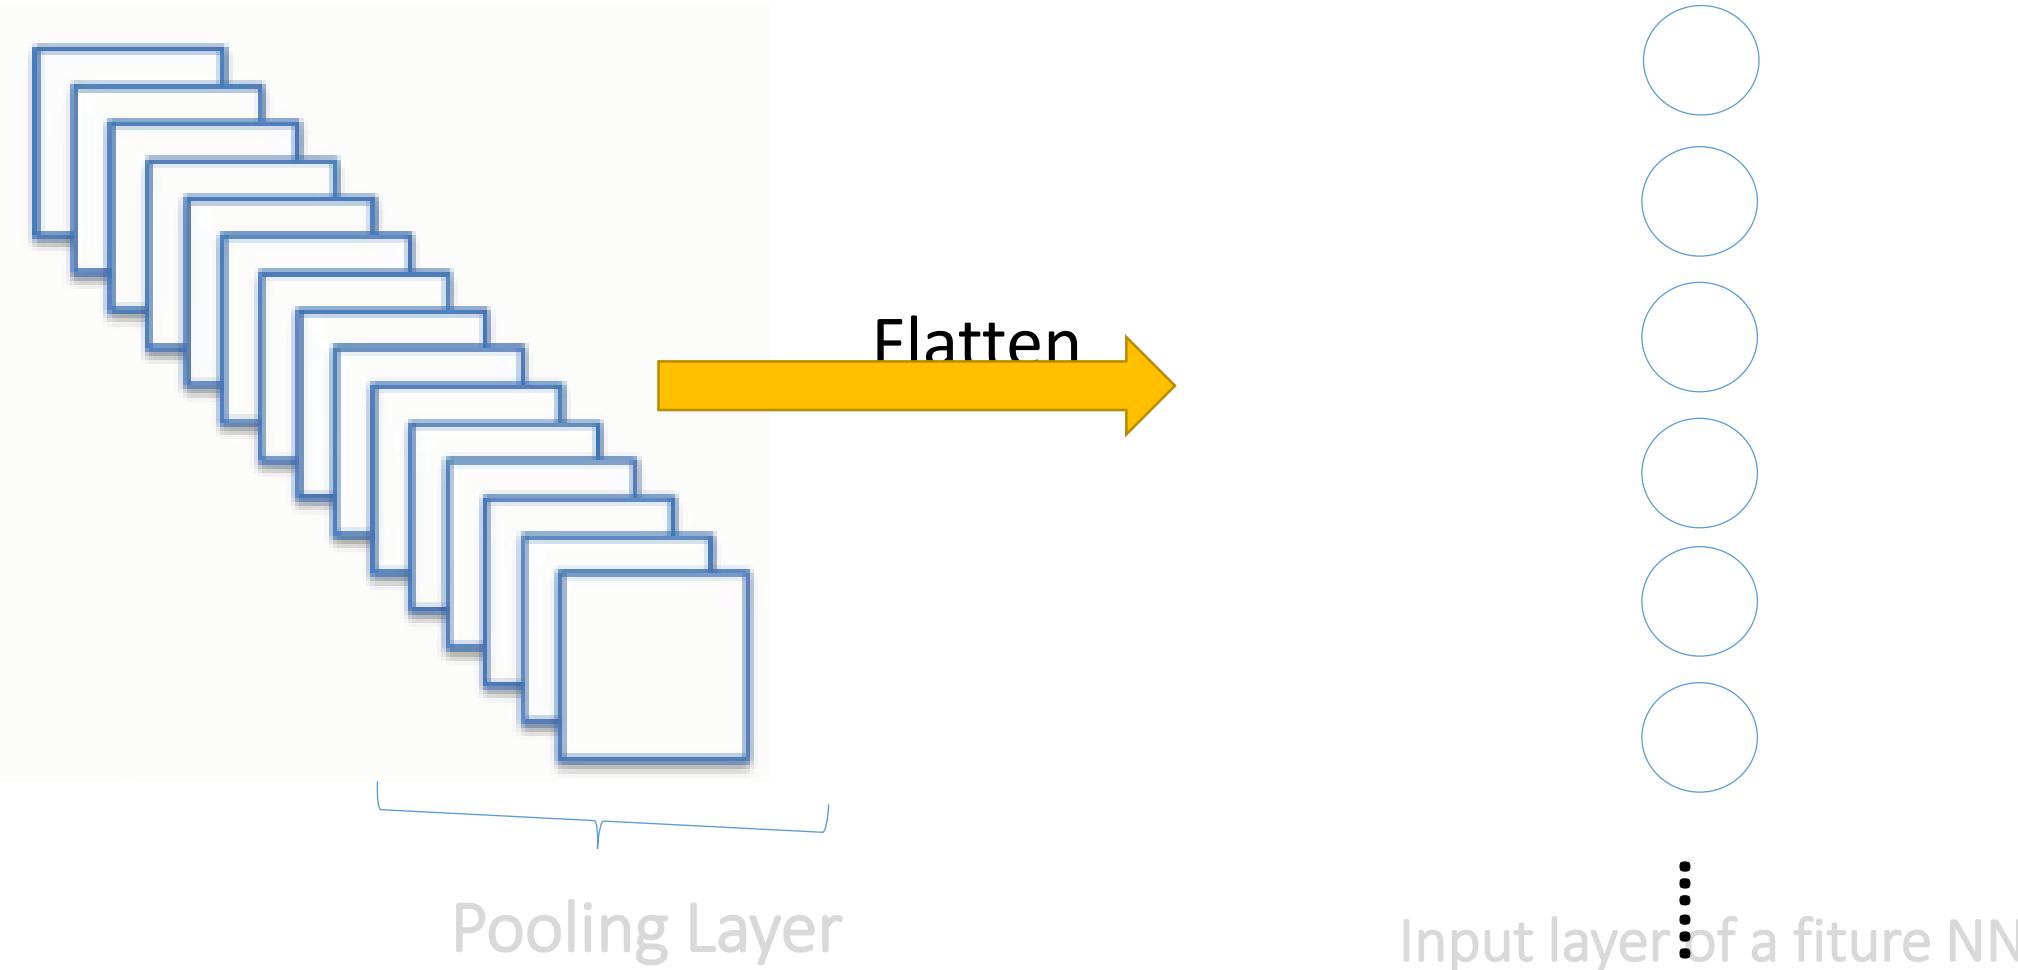


Pooled Feature Map

Flatten
Max-pooling

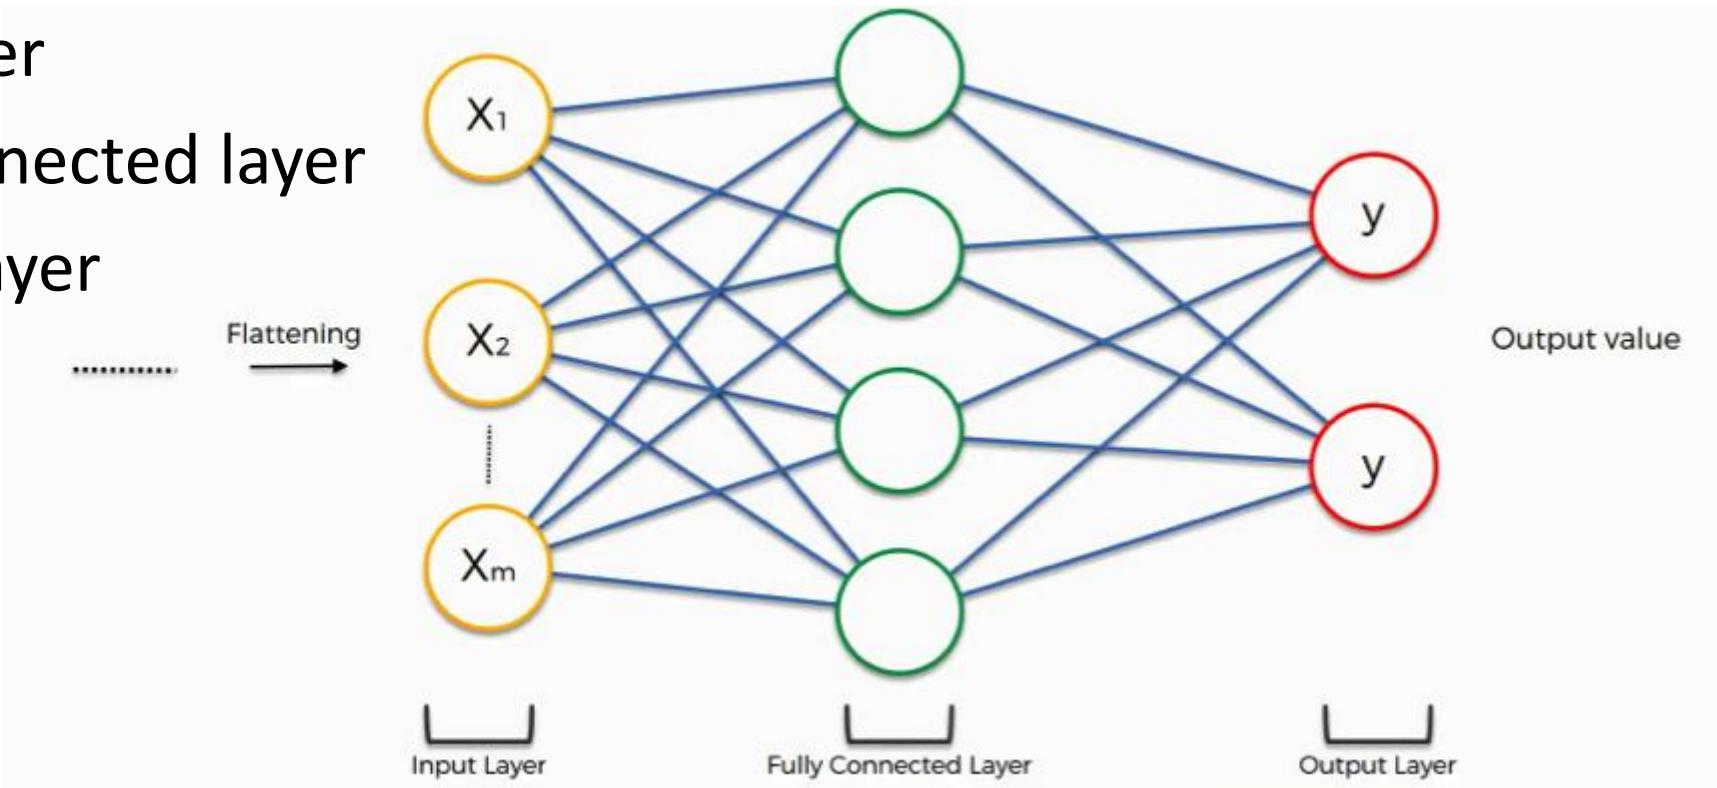


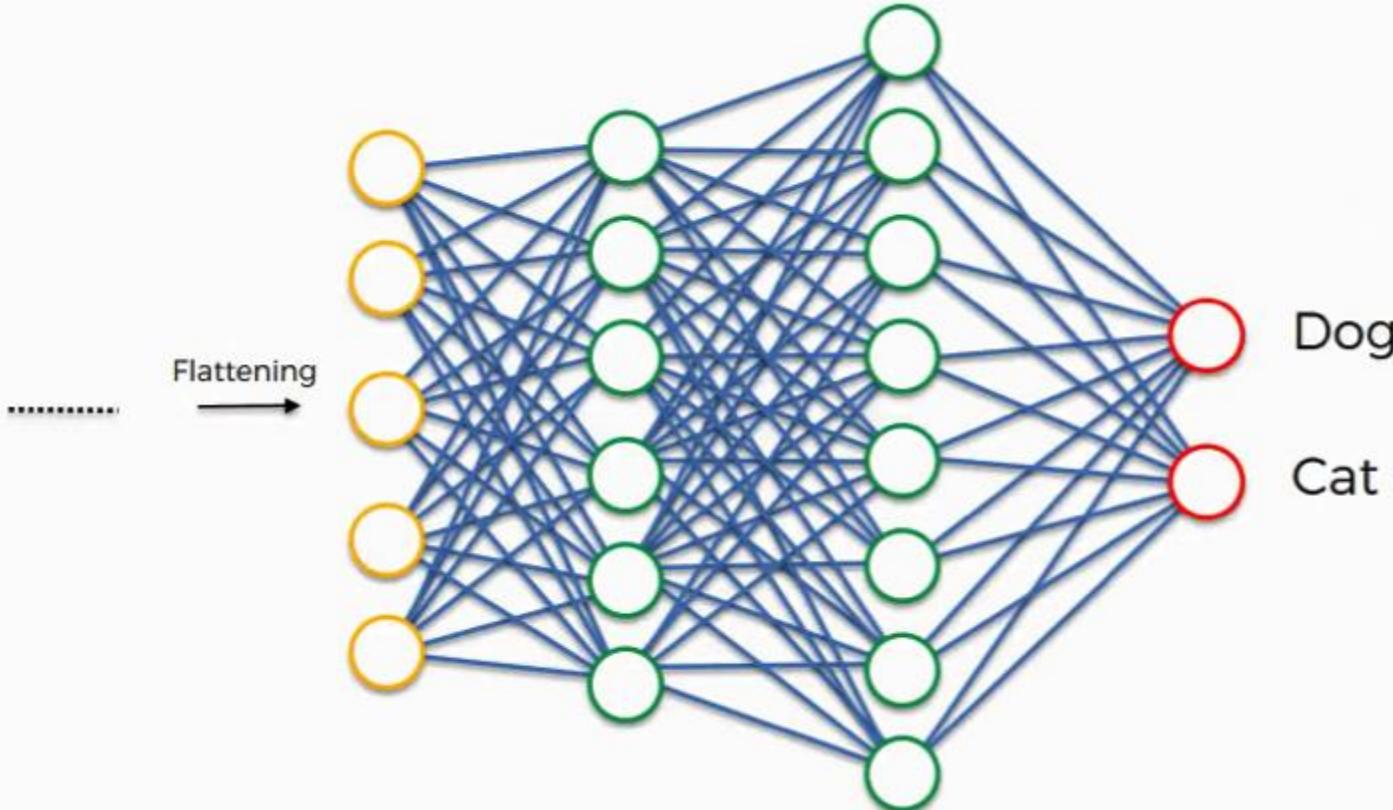
Pooling, Flatten



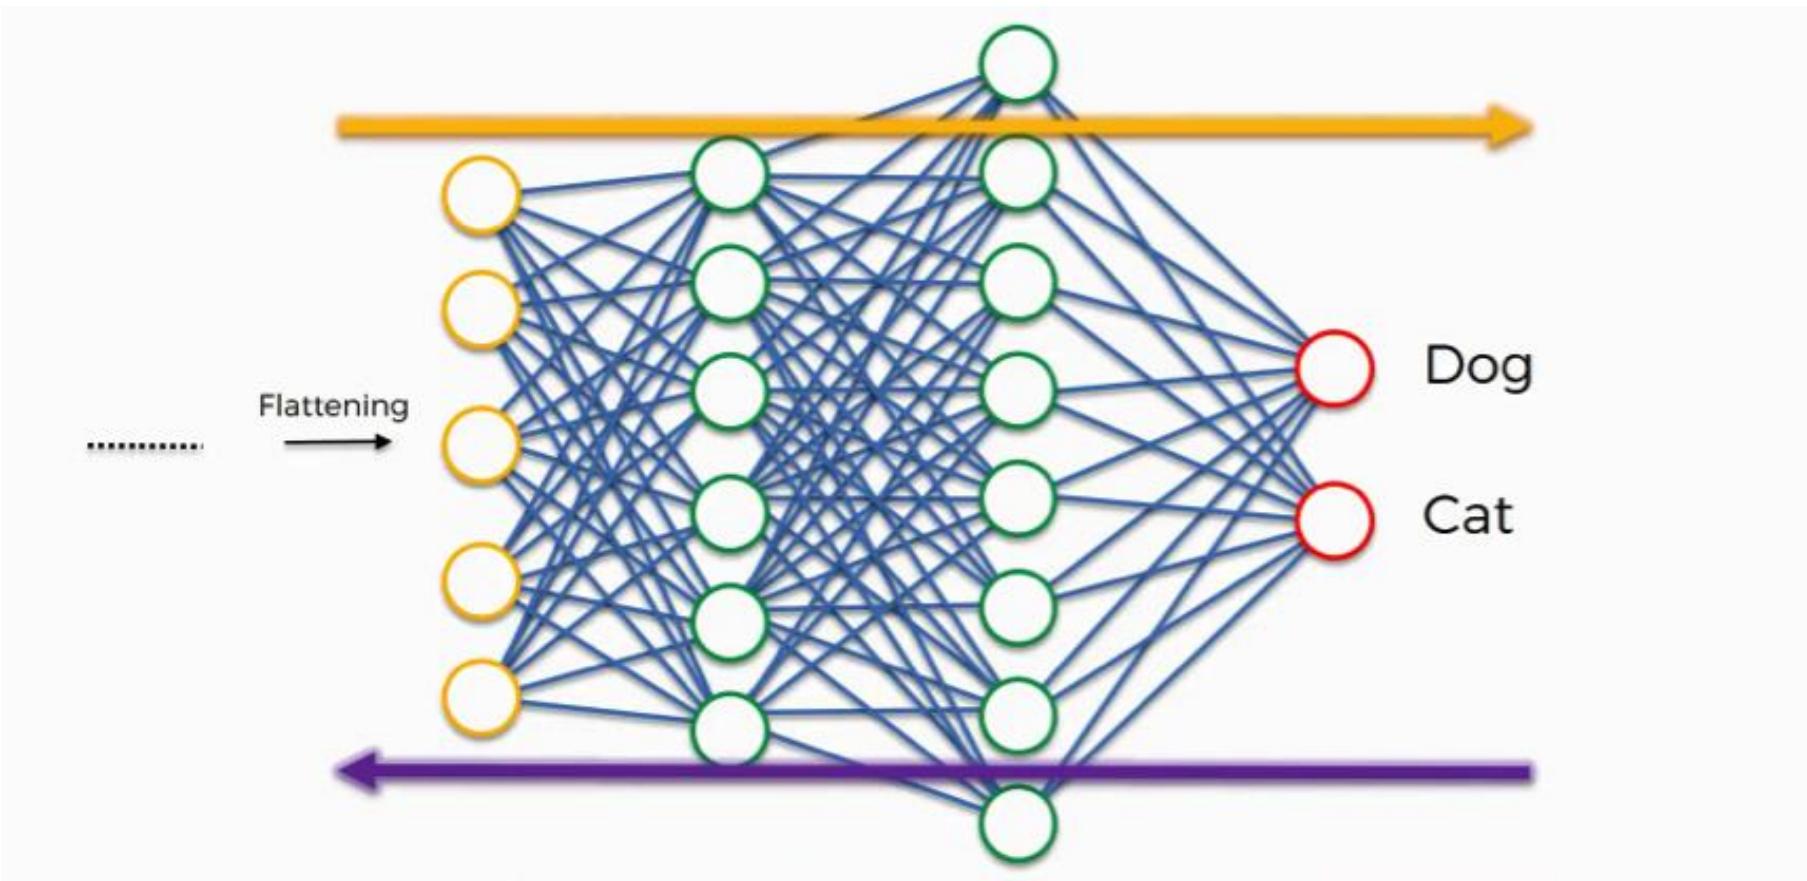
Full connection step:

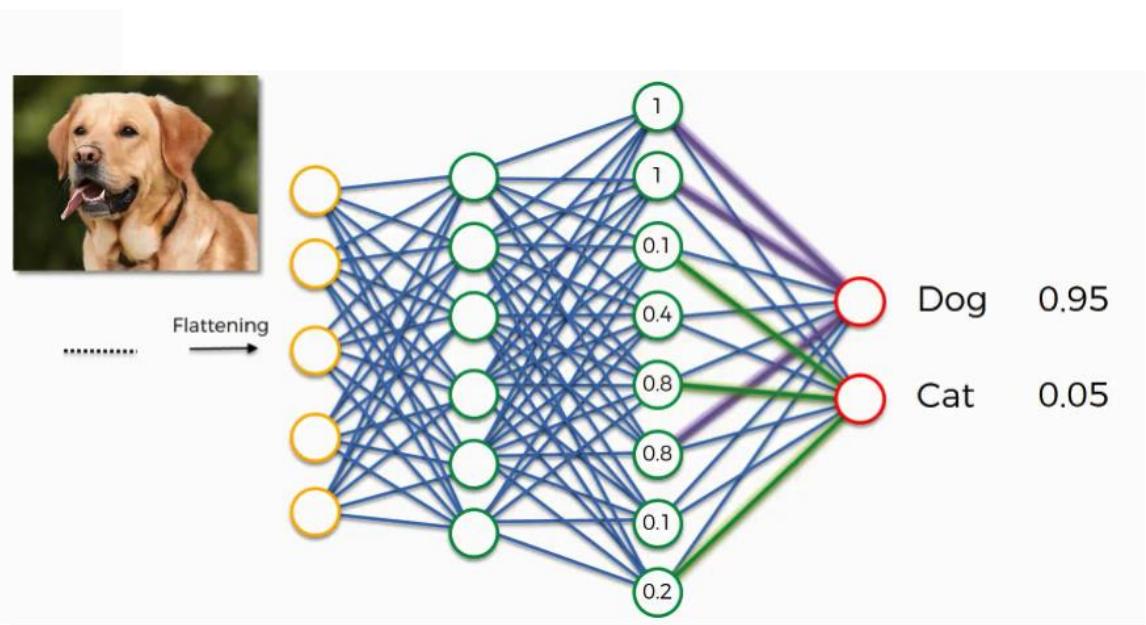
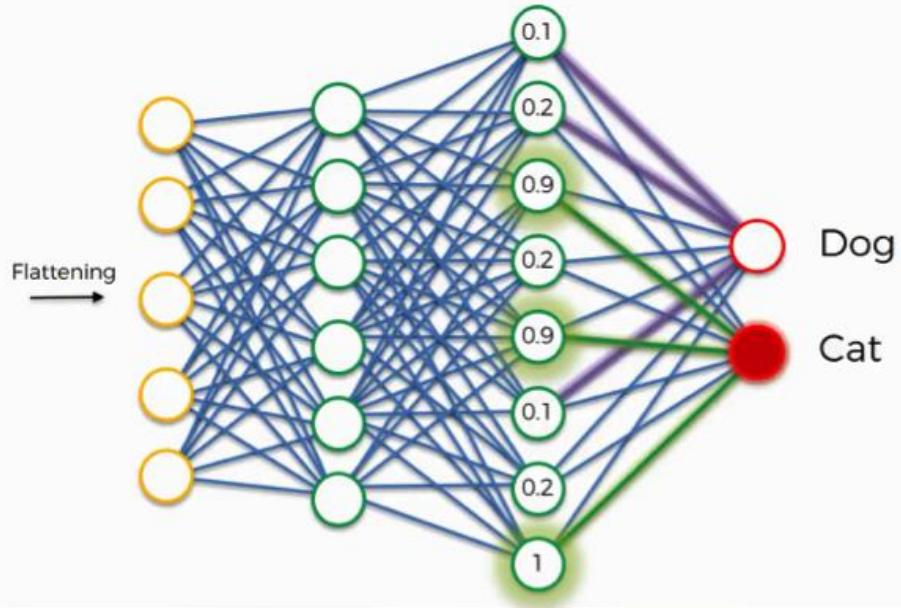
- Input layer
- Fully-connected layer
- Output layer

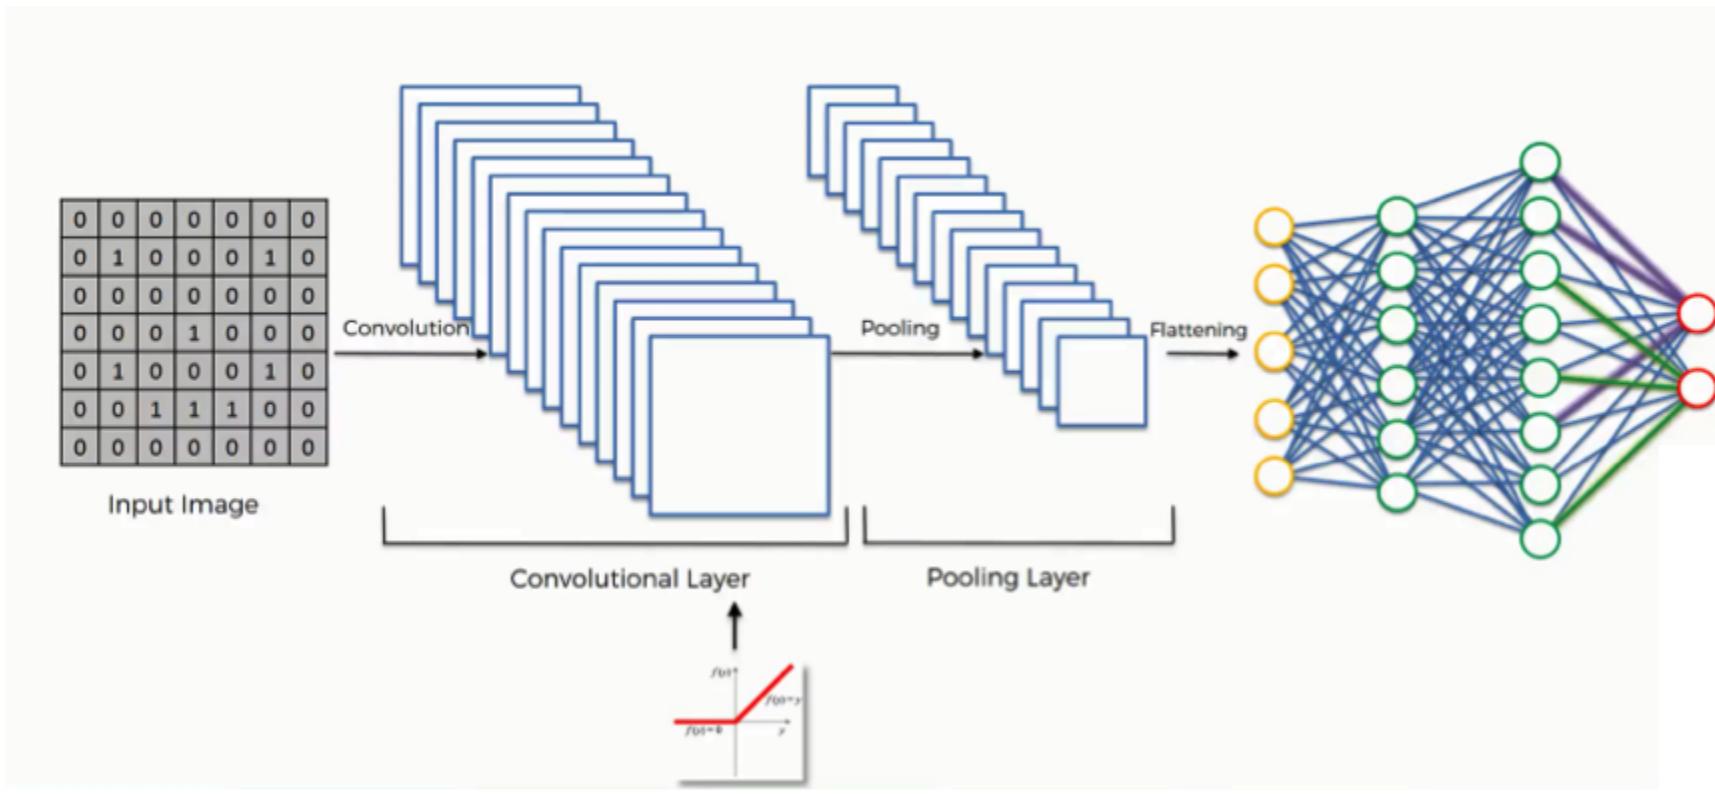




Class Recognition



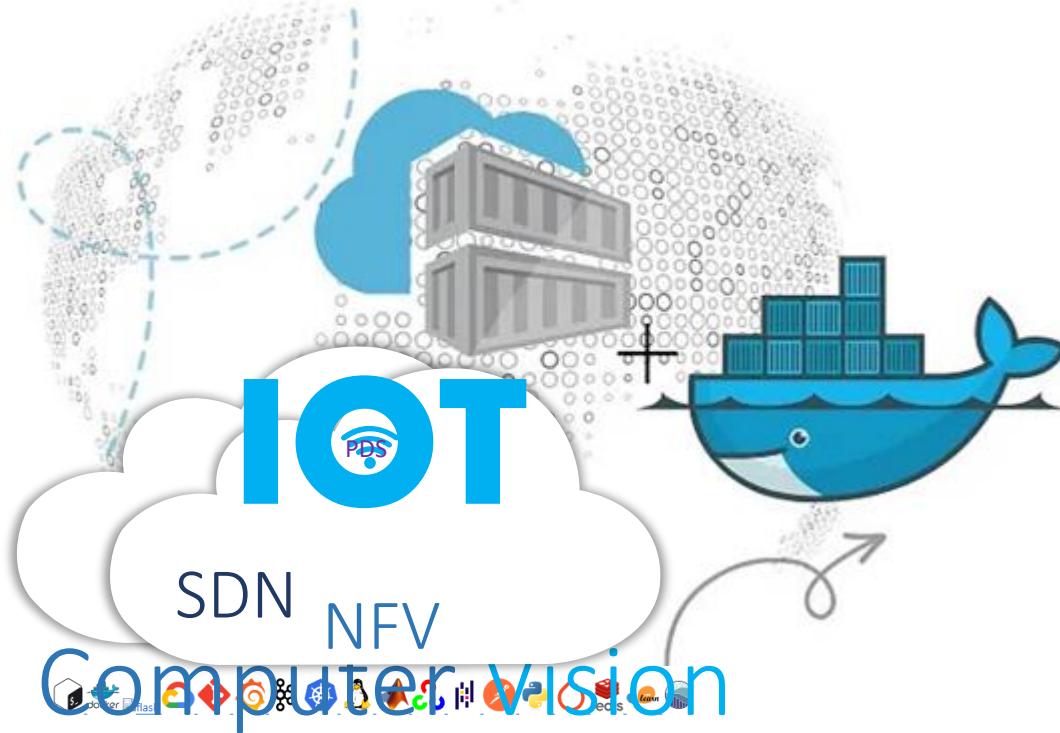




$$z =: \sum_{i=1}^n w_i x_i$$

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$$

$$p_i = \text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$



Thank you!

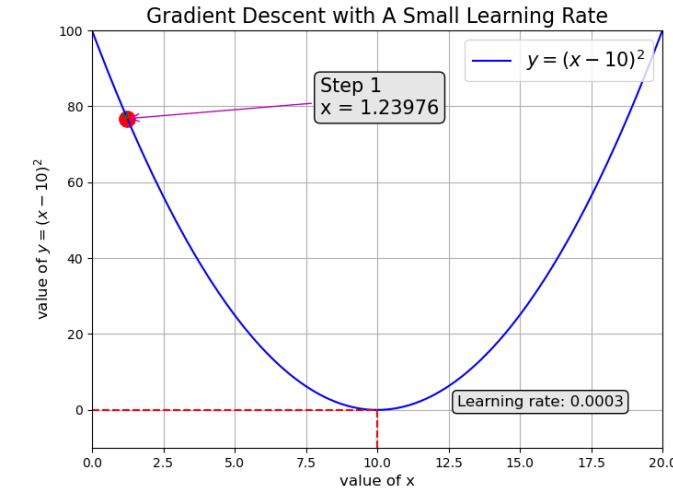
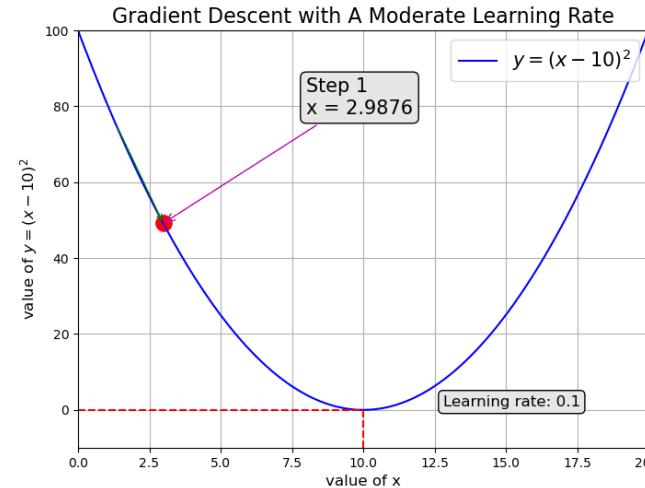
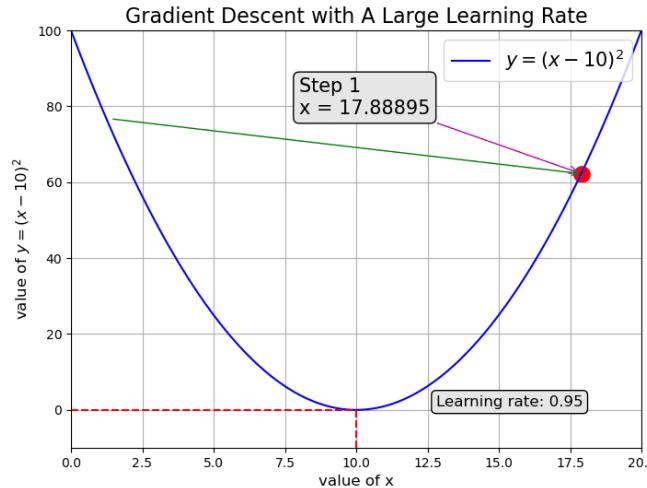
Thank you For being a great class!

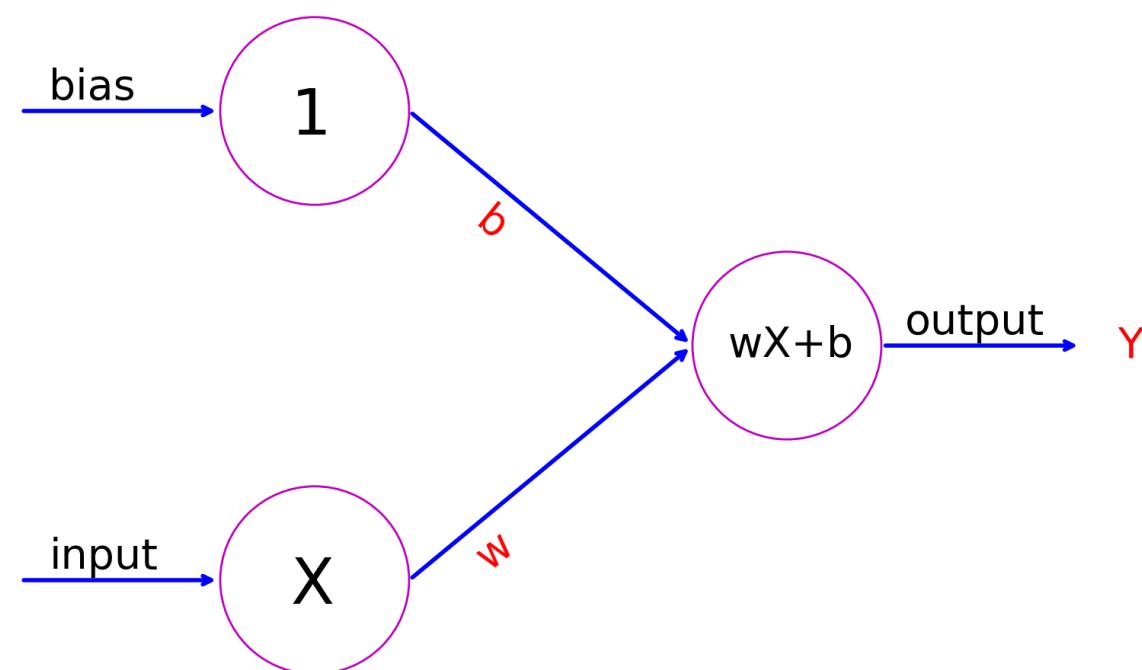
<https://github.com/siagianp>

https://github.com/amelcharolinesgn2/IoT_simulator-mqtt-NodeRed

<https://www.youtube.com/@AmelOline/videos>

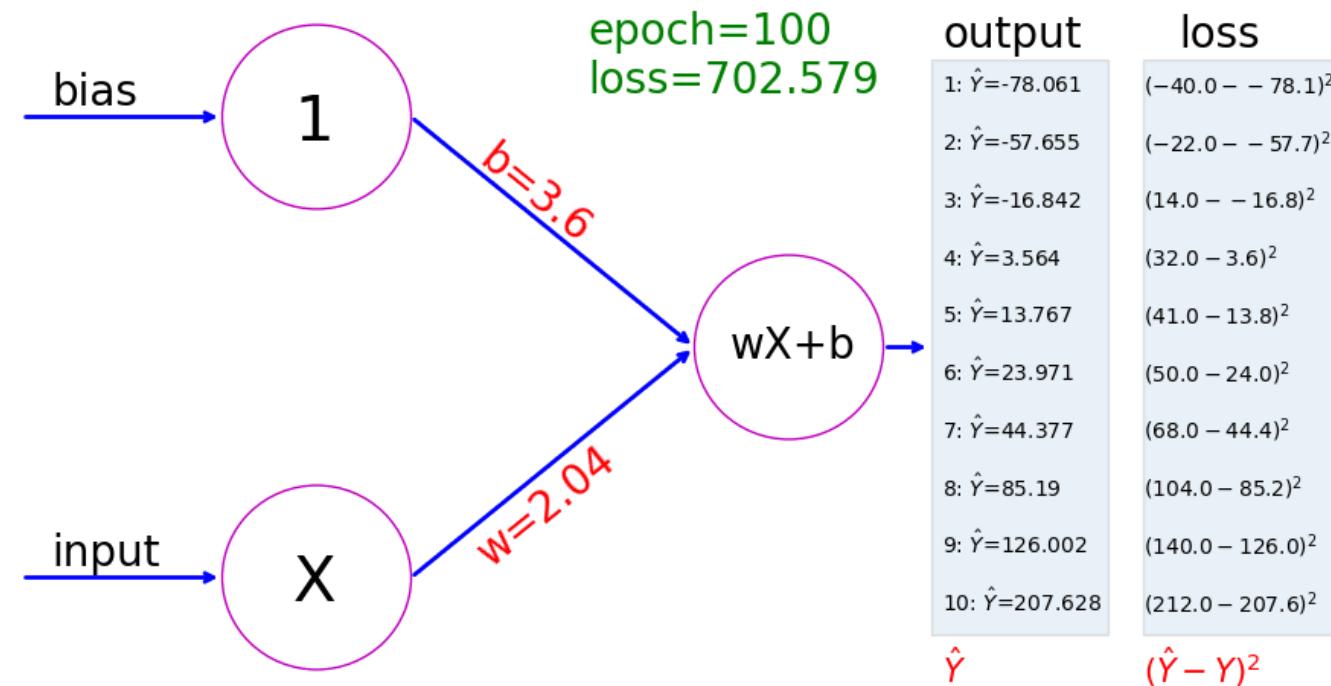
Gradient Descent -- Where Magic Happens



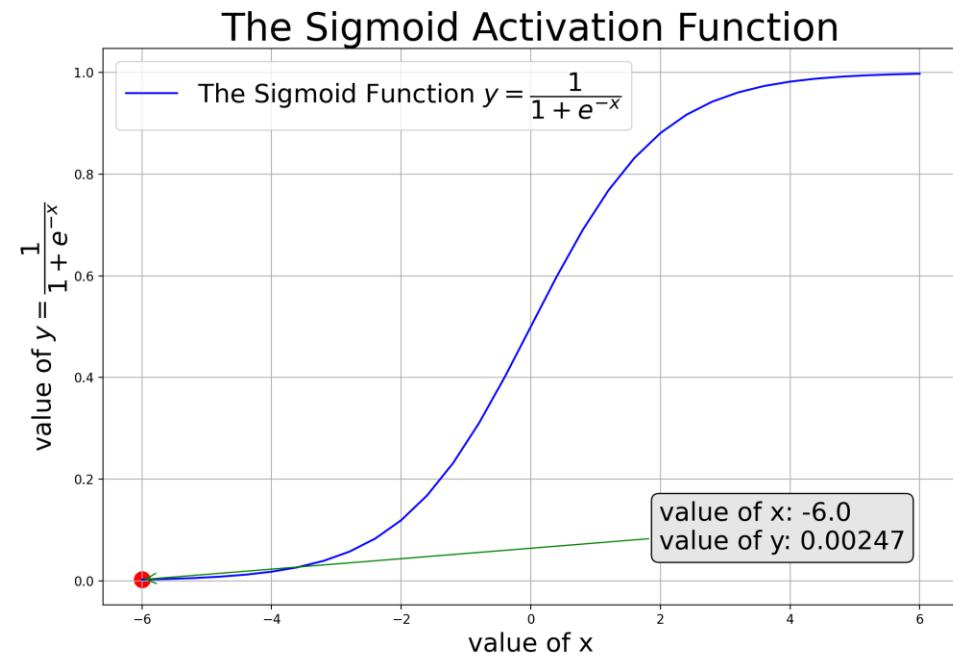
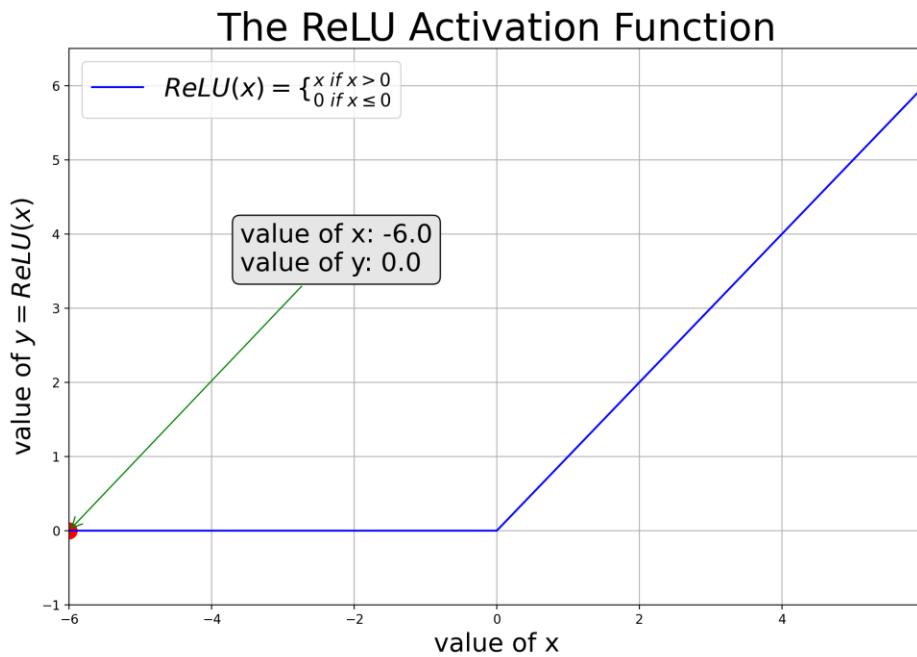


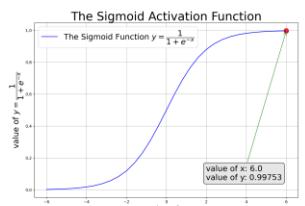
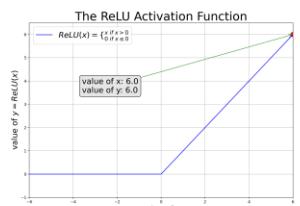
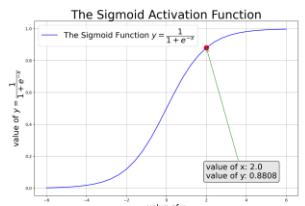
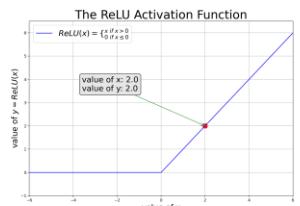
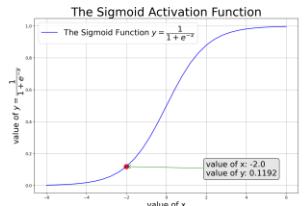
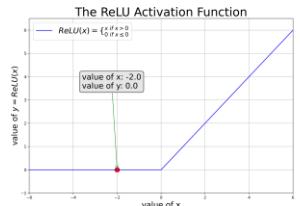
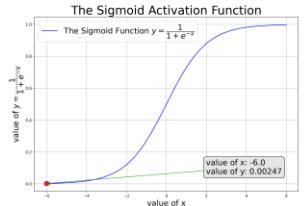
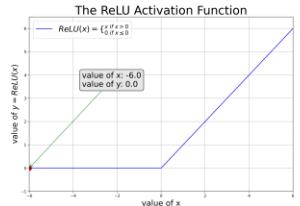
Introduction to Neural Networks

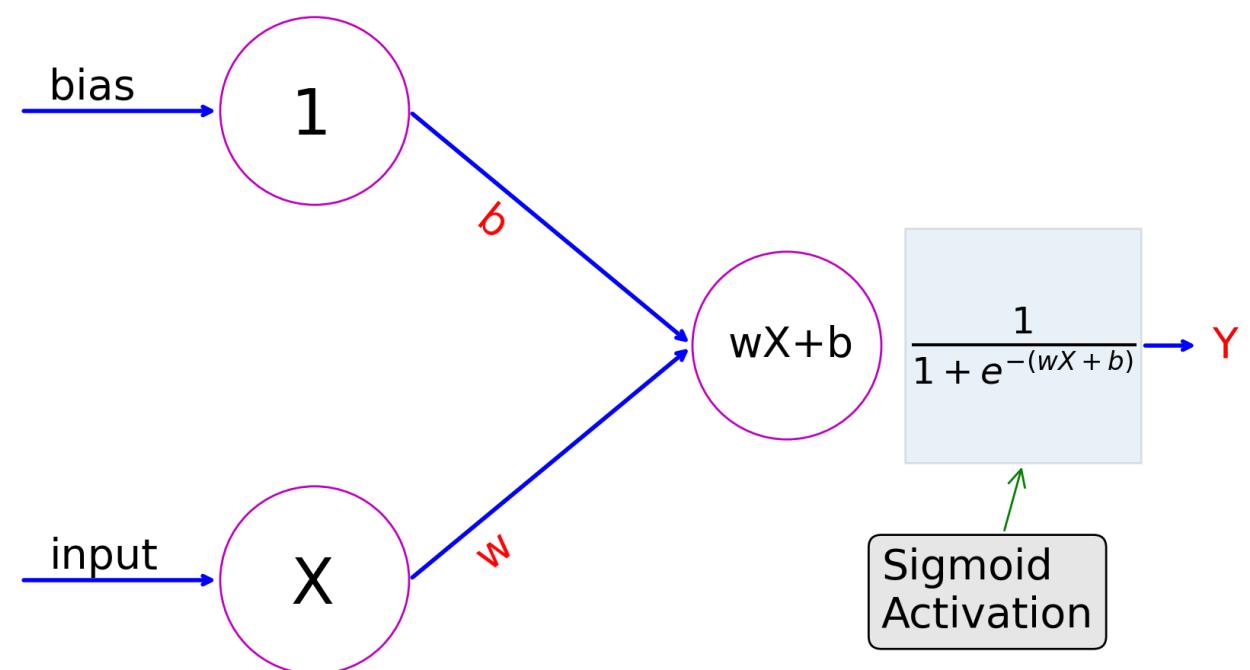
A neural network learns from ten pairs of values: $X=-40, Y=-40; X=0, Y=32; \dots; X=100, Y=212$. The model first assigns values to w and b in $Y=wX+b$; it then uses gradient descent to adjust the values of w and b ; finally it figures out a linear relation between X and Y that corresponds to the relation between Celsius and Fahrenheit $Y=1.8*X+32$. This animation shows the value of w and b in each step:



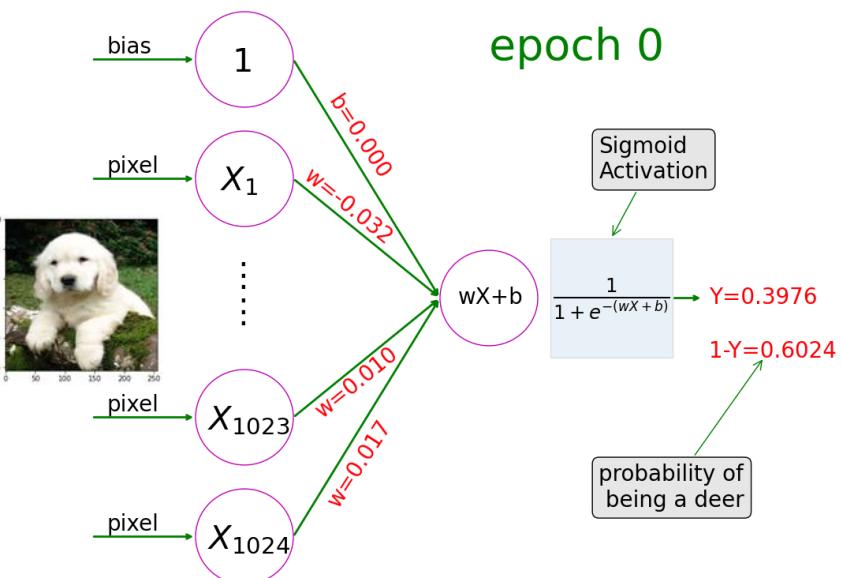
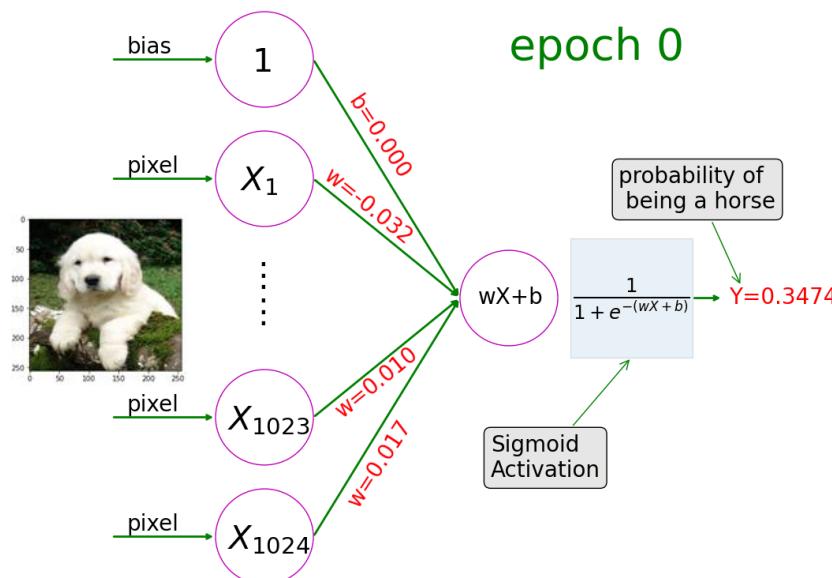
Activation Functions



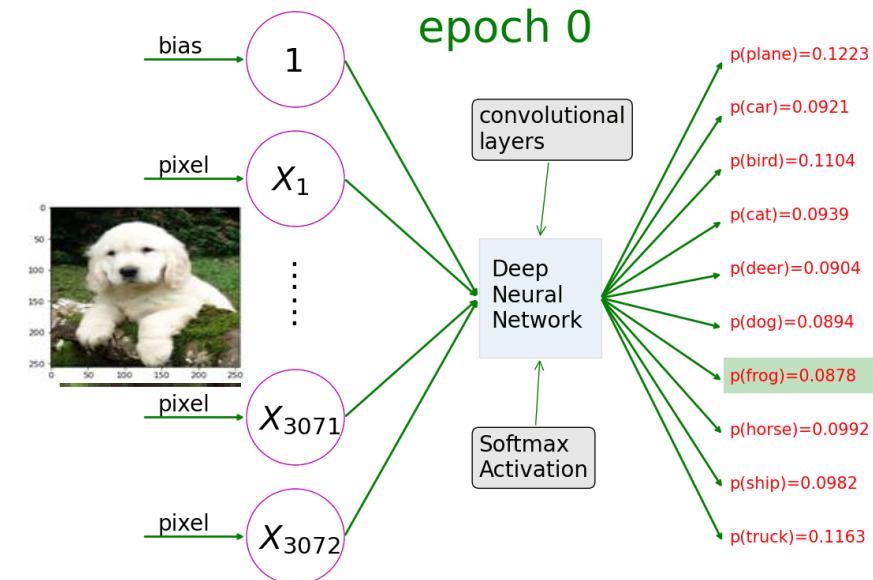
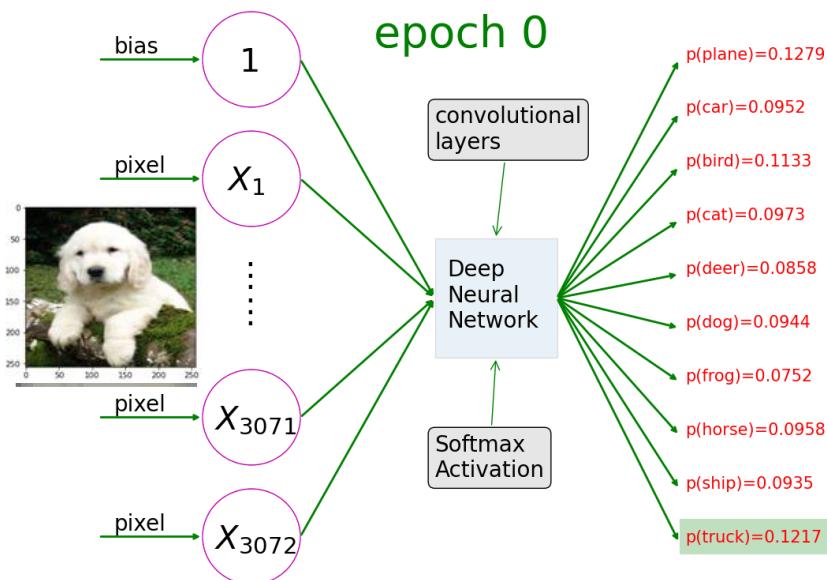




Binary Classifications



Multi-Class Image Classifications



Convolutional Neural Networks

0	0	0	0	0
0	1	1	1	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Apply a 3 by 3 filter on a 3 by 3 image
with zero-padding
stride=1

0	0	0
1	1	1
0	0	0

2.0	3.0	2.0
0.0	0.0	0.0
0.0	0.0	0.0

1	0	0	1	1	0
1	0	1	0	0	1
0	0	0	1	1	0
0	1	1	0	1	0
1	0	1	1	0	0
0	1	0	0	1	0

Apply a 2 by 2 filter on a 6 by 6 image
without zero-padding
stride=2

1	0
0	1

1.0	0.0	2.0
1.0	0.0	1.0
2.0	1.0	0.0